

Todo list

■ Important: Lay abstract.	iii
■ Important: Replace reading notes.	xv
■ OG ISO/IEC, 2005	2
■ OG IEEE, 1996	3
■ OG 2005	3
■ See #22	4
■ See #155	7
■ See #176	7
■ OG Myers 1976	7
■ OG ISO/IEC 2014	8
■ OG Black, 2009	8
■ See #21	8
■ See #44, #119, and #39	10
■ See #119	10
■ See #52	10
■ Present tense in footnote?	11
■ See #155	11
■ See #155	11
■ See #140	13
■ See #140 and #151	13
■ See #137 and #138	13
■ <i>Later:</i> Ensure this is up to date	14
■ Q #1: Does this make sense in isolation or do I need more detail?	14
■ See #176	14
■ Q #2: Better name for this?	17
■ Present tense?	17
■ Present tense?	19
■ See #89	19
■ See #55	19
■ Q #3: Should glossary headers be capitalized?	20
■ Q #4: Are these references helpful or distracting?	20
■ Q #5: Is my justification for inferring “Type” sufficient?	22
■ OG ISO/IEC, 2014	23
■ OG Miller et al., 1994	24
■ OG Hetzel88	24
■ Q #6: Is this sufficient to explain ChatGPT’s role in our methodology? .	29

■ Q #7: Is this sentence needed?	32
■ Q #8: Do I need an example of this?	35
■ Q #9: Since this doesn't come up in our thesis, do I even need an example?	35
■ See #83	37
■ See #83 and #176	39
■ See #137 and #138	40
■ Later: Ensure this is up to date	40
■ Later: Ensure this is up to date	41
■ See #184	41
■ Q #10: Is this too strong of a synonym for "ground truth"?	47
■ OG Hetzel88	52
■ Q #11: I'm counting this as an ambiguous parent relation occurring in ISTQB; is this sufficient?	52
■ Are these separate approaches?	53
■ See #14	53
■ OG Hass, 2008	54
■ OG Beizer	55
■ OG 2015	55
■ OG ISO1984	55
■ OG 2015	55
■ OG Reid, 1996	56
■ Q #12: Does this merit counting this as an Ambiguity as well as a Contradiction?	57
■ Is this a def flaw?	57
■ OG 2015	57
■ OG [14]	57
■ Q #13: I ignore "syntax errors, runtime errors, and logical errors" (Washizaki, 2024, p. 16-13, cf. p. 18-13) since they seem to be in dif- ferent domains. Does that make sense? How should I document this?	57
■ OG ISO/IEC, 2005	58
■ Q #14: Should we classify this as a Synonym Flaw at all?	59
■ Q #15: Shouldn't this be "is", referring to "test item"?	60
■ See #21	61
■ OG Beizer	61
■ See #64	62
■ OG 2013	62
■ See #14	68
■ OG Hass, 2008	68
■ FIND SOURCES	69
■ OG Hetzel88	69
■ more examples?	69
■ Better way to handle/display this?	70
■ Are these separate approaches?	71
■ OG 2015	72
■ OG Reid 1996	72

■ Q #16: I'm counting this as an ambiguous parent relation occurring in ISTQB; is this sufficient?	76
■ Q #17: Should we classify this as a Synonym Flaw at all?	76
■ Q #18: Shouldn't this be "is", referring to "test item"?	77
■ OG Beizer	78
■ OG 2015	78
■ OG ISO1984	78
■ OG 2015	78
■ OG 2015	78
■ OG [14]	78
■ Q #19: I ignore "syntax errors, runtime errors, and logical errors" (Washizaki, 2024, p. 16-13, cf. p. 18-13) since they seem to be in different domains. Does that make sense? How should I document this? .	78
■ OG Reid, 1996	79
■ Q #20: Does this merit counting this as an Ambiguity as well as a Contradiction?	80
■ Is this a def flaw?	82
■ OG ISO/IEC, 2005	82
■ See #182	83
■ van Vliet (2000, p. 399) may list these as synonyms; investigate	85
■ OG PMBOK 5th ed.	86
■ find more academic sources	87
■ See #59	91
■ OG Reid, 1996	91
■ Same label to different phantomsections; is that OK?	93
■ find more academic sources	94
■ See #35	96
■ OG IEEE 2013	98
■ investigate OG sources	102
■ Should I include the definition of Constraints ?	102
■ cite Dr. Smith	102
■ add refs to 'underlying Theory' comment and 'not all outputs be IMs' comment	102
■ add constraints	102
■ A justification for why we decided to do this should be added	104
■ add acronym?	104
■ is this punctuation right?	104
■ See #54	106
■ OG ISO 26262	106
■ see ISO 29119-11	107
■ Investigate	108
■ OG [11, 6]	108
■ OG Halfond and Orso, 2007	108
■ OG Artzi et al., 2008	108
■ Investigate!	109
■ Add paragraph/section number?	110

■ Add example	110
■ Add source(s)?	110
■ Important: “Important” notes.	112
■ Generic inlined notes.	112
■ <i>Later:</i> TODO notes for later! For finishing touches, etc.	112
■ <i>Easy:</i> Easier notes.	112
■ <i>Needs time:</i> Tedious notes.	112
■ Q #21: Questions I might have?	112
■ Investigate further	125
■ See #41 and #44	127
■ See #39, #44, and #28	127
■ OG PMBOK	127
■ See #63	127
■ OG Alalfi et al., 2010	127
■ OG Artzi et al., 2008	127
■ OG [19]	128
■ Q #22: Is this orthogonal? Investigate legacy testing	129
■ Q #23: Is this orthogonal? Investigate role-based testing	129
■ Q #24: Is this orthogonal? Investigate scenario-based testing	129
■ Q #25: Is this orthogonal? Investigate scenario-based testing	129
■ OG ISO/IEC, 2009	130
■ OG ISO/IEC, 2009	130
■ OG Beizer	133
■ OG 2015	134
■ OG Reid 1996	134
■ FIND SOURCES	134
■ OG Hetzel88	134
■ more examples?	134

PUTTING SOFTWARE TESTING TERMINOLOGY TO THE TEST

PUTTING SOFTWARE TESTING TERMINOLOGY TO THE TEST

By SAMUEL CRAWFORD, B.Eng.

A THESIS
SUBMITTED TO THE DEPARTMENT OF COMPUTING AND SOFTWARE
AND THE SCHOOL OF GRADUATE STUDIES
OF MCMASTER UNIVERSITY
IN PARTIAL FULFILLMENT OF THE REQUIREMENTS
FOR THE DEGREE OF
MASTER OF APPLIED SCIENCE

Master of Applied Science (2025)
(Department of Computing and Software)

McMaster University
Hamilton, Ontario

TITLE: Putting Software Testing Terminology to the Test
AUTHOR: Samuel Crawford, B.Eng.
SUPERVISOR: Dr. Carette and Dr. Smith
PAGES: xvi, 137

Lay Abstract

Important: Lay abstract.

Abstract

Despite the prevalence and importance of software testing, it lacks a standardized and consistent taxonomy, instead relying on a large body of literature with many flaws—even within individual documents! This hinders precise communication, contributing to misunderstandings when planning and performing testing. In this thesis, we explore the current state of software testing terminology by:

1. identifying established standards and prominent testing resources,
2. capturing relevant testing terms from these sources, along with their definitions and relationships (both explicit and implicit), and
3. constructing graphs to visualize and analyze these data.

This process uncovers 558 test approaches and four in-scope methods for deriving test approaches, such as those related to 75 software qualities. We also build a tool for generating graphs that illustrate relations between test approaches and track flaws captured by this tool and manually through the research process. This reveals 275 flaws, including nine terms used as synonyms to two (or more) disjoint test approaches and 17 pairs of test approaches that may either be synonyms or have a parent-child relationship. This also highlights notable confusion surrounding functional, operational acceptance, recovery, and scalability testing. Our findings make clear the urgent need for improved testing terminology so that the discussion, analysis and implementation of various test approaches can be more coherent. We provide some preliminary advice on how to achieve this standardization.

Acknowledgements

ChatGPT was used to help generate supplementary Python code for constructing graphs and generating \LaTeX code, including regex. ChatGPT and GitHub Copilot were both used for assistance with \LaTeX formatting. ChatGPT and ProWritingAid were both used for proofreading. Jason Balaci’s McMaster thesis template provided many helper \LaTeX functions. A special “thank you” to Christopher William Schankula for help with \LaTeX , Reed Mullanix with formal logic, and various friends for discussing software testing with me and providing many of the approaches in Section 3.6 (ChatGPT also provided pointers to the potential existence of some of these approaches). Finally, Dr. Spencer Smith and Dr. Jacques Carette have been great supervisors and valuable sources of guidance and feedback.

Contents

Todo list	i
Lay Abstract	iii
Abstract	iv
Acknowledgements	v
Contents	vi
List of Figures	x
List of Tables	xi
List of Source Codes	xii
List of Abbreviations and Symbols	xiii
Reading Notes	xv
Declaration of Academic Achievement	xvi
1 Introduction	1
2 Terminology	7
2.1 Test Approaches	7
2.1.1 Approach Categories	8
2.1.2 Synonym Relations	10
2.1.3 Parent-Child Relations	10
2.2 Flaws	11
2.2.1 Flaw Manifestations	11
2.2.2 Flaw Domains	13
2.2.3 Additional Flaw Terminology	13
2.3 Explicitness	14
2.4 Credibility	15

3	Methodology	16
3.1	Sources	17
3.1.1	Established Standards	17
3.1.2	Terminology Collections	18
3.1.3	Textbooks	18
3.1.4	Papers and Other Documents	19
3.2	Identifying Relevant Terms	19
3.3	Recording Relevant Terms	20
3.4	Implicit Information	23
3.5	Undefined Terms	25
3.6	Stopping Criteria	26
4	Threats to Validity	28
4.1	Threats from Approach Categories	28
4.1.1	Overloaded Terms	28
4.1.2	Alternate Approach Categorizations	29
5	Tools	33
5.1	Approach Relation Visualization	33
5.2	Flaw Analysis	36
5.2.1	Automated Flaw Analysis	37
5.2.2	Augmented Flaw Analysis	39
5.3	LaTeX Commands	41
6	Observed Flaws	47
6.1	Flaws by Manifestation	51
6.1.1	Mistakes	51
6.1.2	Omissions	52
6.1.3	Contradictions	53
6.1.4	Ambiguities	57
6.1.5	Overlaps	59
6.1.6	Redundancies	60
6.2	Flaws by Domain	60
6.2.1	Approach Category Flaws	61
6.2.2	Synonym Relation Flaws	68
6.2.3	Parent-Child Relation Flaws	71
6.2.4	Definition Flaws	76
6.2.5	Label Flaws	81
6.2.6	Scope Flaws	83
6.2.7	Traceability Flaws	84
6.3	Functional Testing	85
6.3.1	Specification-based Testing	85
6.3.2	Correctness Testing	85
6.3.3	Conformance Testing	86
6.3.4	Functional Suitability Testing	86
6.3.5	Functionality Testing	86

6.4	Operational (Acceptance) Testing	87
6.5	Recovery Testing	87
6.6	Scalability Testing	88
6.7	Compatibility Testing	89
6.8	Inferred Flaws	90
6.8.1	Inferred Synonym Flaws	90
6.8.2	Inferred Parent Flaws	90
6.8.3	Inferred Category Flaws	91
6.8.4	Other Inferred Flaws	92
7	Recommendations	93
7.1	Operational (Acceptance) Testing	94
7.2	Recovery Testing	94
7.3	Scalability Testing	95
7.4	Performance(-related) Testing	96
7.5	Compatibility Testing	98
8	Development Process	100
8.1	Improvements to Manual Test Code	101
8.1.1	Testing with Mocks	101
8.2	The Use of Assertions in Code	102
8.3	Generating Requirements	102
9	Research	104
9.1	Existing Taxonomies, Ontologies, and the State of Practice	104
9.2	Definitions	105
9.2.1	Documentation	105
9.3	General Testing Notes	106
9.3.1	Test Oracles	106
9.3.2	Generating Test Cases	107
9.4	Examples of Metamorphic Relations	109
9.5	Roadblocks to Testing	110
9.5.1	Roadblocks to Testing Scientific Software	110
10	Extras	111
10.1	Writing Directives	111
10.2	HREFs	111
10.3	Pseudocode Code Snippets	112
10.4	TODOs	112
	Bibliography	113
A	Detailed Scope Analysis	124
A.1	Hardware Testing	124
A.2	V&V of Other Artifacts	125
A.3	Static Testing	126
A.4	Vague Terminology	127

A.5	Language-specific Approaches	127
A.6	Orthogonally Derived Approaches	128
B	Sources by Tier	131
B.1	Established Standards	131
B.2	Terminology Collections	131
B.3	Textbooks	131
B.4	Papers and Other Documents	131
C	Full Lists of Flaws	132
C.1	Full Lists of Flaws by Manifestation	132
C.1.1	Full List of Mistakes	132
D	Code Snippets	135

List of Figures

3.1	Summary of how many sources comprise each source tier.	18
3.2	Procedure for recording test approaches in our glossary.	21
3.3	ISO/IEC and IEEE (2022, p. 1)’s glossary entry for “A/B testing”.	22
3.4	A/B testing’s inclusion in Firesmith (2015, p. 58)’s taxonomy.	22
3.5	Breakdown of how many test approaches are undefined.	26
5.1	Example generated visualizations of parent-child relations.	34
5.2	Synonym relations given explicitly by the literature.	36
5.3	The sets of authors of established standards (as defined in Section 3.1.1).	38
5.4	Example generated graphs containing flaws.	39
6.1	Identified flaws by the source tier responsible. Some bars are omitted as they correspond to comparisons we do not make; see Section 5.2.	48
7.1	Graphs of relations between terms related to recovery testing.	95
7.2	Graphs of relations between terms related to scalability testing.	97
7.3	Proposed relations between rationalized “performance-related testing” terms.	98

List of Tables

1.1	Selected entries from our test approach glossary with “Notes” column excluded for brevity.	6
2.1	Categories of test approaches given by ISO/IEC and IEEE.	9
2.2	Observed flaw manifestations.	12
2.3	Observed flaw domains.	13
3.1	Breakdown of implicit keywords used for recording and analyzing implicitness (defined in Section 2.3).	24
4.1	Categories of testing given by other sources.	30
4.2	Alternate categorizations found in the literature.	31
5.1	Example glossary entries demonstrating how we track parent-child relations (see Section 2.1.3).	34
5.2	\LaTeX macros for calculated values.	42
5.3	\LaTeX macros for referencing well-defined sections.	44
5.4	\LaTeX macros for reused text.	45
6.1	Breakdown of identified Flaws by Manifestation by Source Tier. . .	50
6.2	Breakdown of identified Flaws by Domain by Source Tier.	50
6.3	Different kinds of mistakes found in the literature.	51
6.4	Test approaches with more than one category.	64
6.5	Pairs of test approaches with a parent-child <i>and</i> synonym relation. .	73
6.6	Test approaches inferred to have more than one category.	91

List of Source Codes

10.1 Pseudocode: exWD	111
10.2 Pseudocode: exPHref	112
D.1 Tests for main with an invalid input file	135
D.2 Projectile’s choice for constraint violation behaviour in code	136
D.3 Projectile’s manually created input verification requirement	136
D.4 “MultiDefinitions” (MultiDefn) Definition	136
D.5 Pseudocode: Broken QuantityDict Chunk Retriever	136

List of Abbreviations and Symbols

c-use/C-use	Computation data Use
CT	Continuous Testing
DAC	Differential Assertion Checking
DbIPend	Double Pendulum
DOM	Document Object Model
du-path/DU-path	Definition-Use path
EMSEC	EManations SECurity
FIST	Fault Injection Security Tool
GlassBR	Glass BReaking
HREF	Hypertext REference
IEC	International Electrotechnical Commission
ISTQB	International Software Testing Qualifications Board
ML	Machine Learning
MR	Metamorphic Relation
OAT	Operational Acceptance/Orthogonal Array Testing
OT	Operational Testing
p-use/P-use	Predicate data Use
PAR	Product Anomaly Report
PDF	Portable Document Format
PIR	Product Incident Report
QAI	Quality Assurance Institute
RAC	Runtime Assertion Checking
RQ	Research Question
SglPend	Single Pendulum
SoS	System of Systems
SRS	Software Requirements Specification
SSP	Slope Stability analysis Program
SUT	System Under Test

SV	Software Verification
SWEBOK Guide	Guide to the SoftWare Engineering Body Of Knowledge
SWHS	Solar Water Heating System
TOAT	Taguchi's Orthogonal Array Testing
V&V	Verification and Validation

Reading Notes

Before reading this thesis, I encourage you to read through these notes, keeping them in mind while reading.

- The source code of this thesis is publicly available.
- This thesis template is primarily intended for usage by the computer science community¹. However, anyone is free to use it.
- I’ve tried my best to make this template conform to the thesis requirements as per those set forth in 2021 by McMaster University. However, you should double-check that your usage of this template is compliant with whatever the “current” rules are.

Important: Replace reading notes.

¹Hence why there are some \LaTeX macros for “code” snippets.

Declaration of Academic Achievement

This research and analysis was performed by Samuel Crawford under the guidance, supervision, and recommendations of Dr. Spencer Smith and Dr. Jacques Carette. The resulting contributions are three glossaries—one for each of test approaches, software qualities, and supplementary terms (see Section 3.3)—as well as the tools for data visualization and automated analysis outlined in Chapter 5. These are all available on an open-source repo for independent analysis and, ideally, extension as more test approaches are discovered and documented.

Chapter 1

Introduction

As with all fields of science and technology, software development should be approached systematically and rigorously. Peters and Pedrycz claim that “to be successful, development of software systems requires an engineering approach” that is “characterized by a practical, orderly, and measured development of software” (2000, p. 3). When a NATO study group decided to hold a conference to discuss “the problems of software” in 1968, they chose the phrase “software engineering” to “imply[] the need for software manufacture to be based on the types of theoretical foundations and practical disciplines, [sic] that are traditional in the established branches of engineering” (Naur and Randell, 1969, p. 13). “The term was not in general use at that time”, but conferences such as this “played a major role in gaining general acceptance ... for the term” (McClure, 2001). While one of the goals of the conference was to “discuss possible techniques, methods and developments which might lead to the[] solution” to these problems (Naur and Randell, 1969, p. 14), the format of the conference itself was difficult to document. Two competing classifications of the report emerged: “one following from the normal sequence of steps in the development of a software product” and “the other related to aspects like communication, documentation, management, [etc.]” (p. 10). Furthermore, “to retain the spirit and liveliness of the conference, ... points of major disagreement have been left wide open, and ... no attempt ... [was] made to arrive at a consensus or majority view” (p. 11)!

Perhaps unsurprisingly, there are still concepts in software engineering without consensus, and many of them can be found in the subdomain of software testing. Kaner et al. give the example of complete testing, which may require the tester to discover “every bug in the product”, exhaust the time allocated to the testing phase, or simply implement every test previously agreed upon (2011, p. 7). Having a clear definition of “complete testing” would reduce the chance for miscommunication and, ultimately, the tester getting “blamed for not doing ... [their] job” (p. 7). Because software testing uses “a substantial percentage of a software development budget (in the range of 30 to 50%)”, which is increasingly true “with the growing complexity of software systems” (Peters and Pedrycz, 2000, p. 438), this is crucial to the efficiency of software development. Even more foundationally, if software engineering holds code to high standards of clarity, consistency, and robustness, the same should apply to its supporting literature!

We noticed this lack of a standard language for software testing while working on our own software framework, Drasil (Carette et al., 2021), with the goal of “generating all of the software artifacts for (well understood) research software”. Currently, these include Software Requirements Specifications (SRSs), READMEs, Makefiles, and code in up to six languages, depending on the specific case study (Hunt et al., 2021). To improve the quality, functionality, and maintainability of Drasil, we want to add test cases to this list (Smith, 2024; Hunt et al., 2021). This process would be a part of our “continuous integration system, [sic] so that the generated code for the case studies is automatically tested with each build” and would be a big step forward, since we currently only “test that the generated code compiles” (Smith, 2024). However, before we can include test cases as a generated artifact, the underlying domain—software testing—needs to be “well understood”, which requires a “stable knowledge base” (Hunt et al., 2021).

Unfortunately, a search for a systematic, rigorous, and complete taxonomy for software testing revealed that the existing ones are inadequate and mostly focus on the high-level testing process rather than the testing approaches themselves:

- Tebes et al. (2020) focus on *parts* of the testing process (e.g., test goal, test plan, testing role, testable entity) and how they relate to one another,
- Souza et al. (2017) prioritize organizing test approaches over defining them,
- Firesmith (2015) similarly defines relations between test approaches but not the approaches themselves, and
- Unterkalmsteiner et al. (2014) focus on the “information linkage or transfer” (p. A:6) between requirements engineering and software testing and “do[] not aim at providing a systematic and exhaustive state-of-the-art survey of [either domain]” (p. A:2).

In addition to these taxonomies, many standards documents (see Section 3.1.1) and terminology collections (see Section 3.1.2) define testing terminology, albeit with their own issues.

For example, a common point of discussion in the field of software is the distinction between terms for when software does not work correctly. We find the following four to be most prevalent:

- **Error:** “a human action that produces an incorrect result” (ISO/IEC and IEEE, 2010, p. 128; van Vliet, 2000, p. 399).
- **Fault:** “an incorrect step, process, or data definition in a computer program” (ISO/IEC and IEEE, 2010, p. 140) inserted when a developer makes an error (pp. 128, 140; Washizaki, 2024, p. 12-3; van Vliet, 2000, pp. 399–400).
- **Failure:** the inability of a system “to perform a required function or ... within previously specified limits” (ISO/IEC and IEEE, 2019a, p. 7; 2010, p. 139; similar in van Vliet, 2000, p. 400) that is “externally visible” (ISO/IEC and IEEE, 2019a, p. 7; similar in van Vliet, 2000, p. 400) and caused by a fault (Washizaki, 2024, p. 12-3; van Vliet, 2000, p. 400).

OG ISO/IEC,
2005

- **Defect:** “an imperfection or deficiency in a project component where that component does not meet its requirements or specifications and needs to be either repaired or replaced” (ISO/IEC and IEEE, 2010, p. 96).

OG IEEE, 1996

OG 2005

This distinction is sometimes important, but not always (Bourque and Fairley, 2014, p. 4-3). The term “fault” is “overloaded with too many meanings, as engineers and others use the word to refer to all different types of anomalies” (Washizaki, 2024, p. 12-3), and “defect” may be used as a “generic term that can refer to either a fault (cause) or a failure (effect)” (ISO/IEC and IEEE, 2017, p. 124; 2010, p. 96). Software testers may even choose to ignore these nuances completely! Patton (2006, pp. 13–14) “just call[s] it what it is and get[s] on with it”, abandoning these four terms, “problem”, “incident”, “anomaly”, “variance”, “inconsistency”, “feature” (!), and “a list of unmentionable terms” in favour of “bug”; after all, “there’s no reason to dice words”!

These decisions are not inherently wrong, since they may be useful in certain contexts or for certain teams (cf. Section 2.1.2). Problems start to arise when teams need to make these decisions in the first place. Patton (2006, p. 14) notes that “a well-known computer company spent weeks in discussion with its engineers before deciding to rename Product Anomaly Reports (PARs) to Product Incident Reports (PIRs)”, a process that required “countless dollars” and updating “all the paperwork, software, forms, and so on”. While consistency and clear terminology may have been valuable to the company, “it’s unknown if [this decision] made any difference to the programmer’s or tester’s productivity” (p. 14). A potential way to avoid similar resource sinks would be to prescribe a standard terminology. Perhaps multiple sets of terms could be designed with varying levels of specificity so a company would only have to determine which one best suits their needs.

But why are minor differences between terms like these even important? The previously defined terms “error”, “fault”, “failure”, and “defect” are used to describe many test approaches, including:

- | | |
|------------------------------|-------------------------------|
| 1. Defect-based testing | 8. Fault injection testing |
| 2. Error forcing | 9. Fault seeding |
| 3. Error guessing | 10. Fault sensitivity testing |
| 4. Error tolerance testing | 11. Fault tolerance testing |
| 5. Error-based testing | 12. Fault tree analysis |
| 6. Error-oriented testing | 13. Fault-based testing |
| 7. Failure tolerance testing | |

When considering which approaches to use or when actually using them, the meanings of these four terms inform what their related approaches accomplish and how to they are performed. For example, the tester needs to know what a “fault” is to perform fault injection testing; otherwise, what would they inject? Information such as this is critical to the testing team, and should therefore be standardized.

These kinds of inconsistencies can lead to miscommunications—such as that previously mentioned by Kaner et al. (2011, p. 7)—and are prominent in the literature. ISO/IEC and IEEE categorize experience-based testing as both a test design technique and a test practice on the same page—twice (2022, Fig. 2, p. 34)! The structure of tours can be defined as either quite general (ISO/IEC and IEEE, 2022, p. 34) or “organized around a special focus” (Hamburg and Mogyorodi, 2024). Load testing is performed with loads “between anticipated conditions of low, typical, and peak usage” (ISO/IEC and IEEE, 2022, p. 5) or loads that are as large as possible (Patton, 2006, p. 86). Alpha testing is performed by “users within the organization developing the software” (ISO/IEC and IEEE, 2017, p. 17), “a small, selected group of potential users” (Washizaki, 2024, p. 5-8), or “roles outside the development organization” conducted “in the developer’s test environment” (Hamburg and Mogyorodi, 2024). It is clear that there is a notable gap in the literature, one which we attempt to describe. While the creation of a complete taxonomy is unreasonable, especially considering the pace at which the field of software changes, we can make progress towards this goal that others can extend and update as new test approaches emerge. The main way we accomplish this is by identifying “flaws” or “inconsistencies” in the literature, or areas where there is room for improvement. We track these flaws according to both *what* information is wrong and *how* (described in more detail in Section 2.2), which allows us to analyze them more thoroughly and reproducibly.

Based on this observed gap in software testing terminology and our original motivation for this research, we only consider the component of Verification and Validation (V&V) that tests code itself. However, some test approaches are only used to testing *other* artifacts, while others can be used for both! In these cases, we only consider the subsections that focus on code. For example, reliability testing and maintainability testing can start *without* code by “measur[ing] structural attributes of representations of the software” (Fenton and Pfleeger, 1997, p. 18), but only reliability and maintainability testing performed on code *itself* is in scope of this research. This is a high-level overview of what is in scope; see Appendix A for more detailed discussion on what we include and exclude.

This document describes our process, as well as our results, in more detail. We start by documenting the 558 test approaches mentioned by 66 sources (described in Section 3.1), recording their names, categories¹ (see Section 2.1.1), definitions, synonyms (see Section 2.1.2), parents (see Section 2.1.3), and flaws (see Section 2.2) (in separate documents) as applicable. We also record any other relevant notes, such as prerequisites, uncertainties, other sources. We follow the procedure laid out in Chapter 3 and use these Research Questions (RQs) as a guide:

1. What testing approaches do the literature describe?
2. Are these descriptions consistent?
3. Can we systematically resolve any of these inconsistencies?

¹A single test approach with more than one category (for example, A/B Testing in Table 1.1) often indicates an underlying flaw; see Section 6.2.1.

See #22

An excerpt of this recorded information (excluding other notes for brevity), is given in Table 1.1. We then create tools to support our analysis of our findings (Chapter 5). Despite the amount of well understood and organized knowledge, the literature is still quite flawed (Chapter 6). This reinforces the need for a proper taxonomy! We then provide some potential solutions covering some of these flaws (Chapter 7).

Table 1.1: Selected entries from our test approach glossary with “Notes” column excluded for brevity.

Name	Approach Category	Definition	Parent(s)	Synonym(s)
A/B Testing	Practice (ISO/IEC and IEEE, 2022, p. 22), Type (inferred from usability testing)	Testing “that allows testers to determine which of two systems or components performs better” (ISO/IEC and IEEE, 2022, pp. 1, 36)	Statistical Testing (ISO/IEC and IEEE, 2022, pp. 1, 36), Usability Testing (Firesmith, 2015, p. 58)	Split-Run Testing (ISO/IEC and IEEE, 2022, pp. 1, 36)
Big-Bang Testing	Technique (Sharma et al., 2021, pp. 601, 603, 605–606), Level (inferred from integration testing)	“Testing in which ... [components of a system] are combined all at once into an overall system, rather than in stages” (ISO/IEC and IEEE, 2017, p. 45)	Integration Testing (ISO/IEC and IEEE, 2017, p. 45; Washizaki, 2024, p. 5-7; Sharma et al., 2021, p. 603; Kam, 2008, p. 42; Peters and Pedrycz, 2000, p. 488, Tab. 12.8)	Sometimes spelled without a hyphen (Peters and Pedrycz, 2000, p. 489)
Classification Tree Method	Technique (ISO/IEC and IEEE, 2022, p. 22; 2021, pp. 2, 12, Fig. 2; Hamburg and Mogyorodi, 2024; Firesmith, 2015, p. 47)	Testing “based on exercising classes in a classification tree” (ISO/IEC and IEEE, 2021, p. 2)	Specification-based Testing (ISO/IEC and IEEE, 2022, p. 22; 2021, pp. 2, 12, Fig. 2; Hamburg and Mogyorodi, 2024; Firesmith, 2015, p. 47), Model-based Testing (ISO/IEC and IEEE, 2022, p. 13; 2021, pp. 6, 12)	Classification Tree Technique (Hamburg and Mogyorodi, 2024), Classification Tree Testing (Firesmith, 2015, p. 47)

Chapter 2

Terminology

Our research aims to describe the current state of software testing literature, including its flaws. Since we critique the lack of clarity, consistency, and robustness in the literature, we need to hold ourselves to a high standard in these areas when defining and using terms. For example, since we focus on how the literature describes “test approaches”, we first define this term (Section 2.1). Likewise, before we can constructively describe the flaws in the literature, we need to define what we mean by “flaw” (Section 2.2). To further prevent bias, we only name and describe classifications and relations that the literature already implies instead of inventing our own; for example, test approaches can have categories (Section 2.1.1), synonyms (Section 2.1.2), and parent-child relations (Section 2.1.3). We also observe flaws having both manifestations (Section 2.2.1) and domains (Section 2.2.2), even though these terms are not used by the literature. All of these classifications and relations follow logically from the literature and as such are technically “results” of our research, but we define them here for clarity since we use them throughout this thesis.

See #155

Since the literature is flawed, we need to be careful with what information we take at face value. We do this by tracking the nuance, or “explicitness”, of information (Section 2.3) found in sources and the “credibility” of these sources themselves (Section 2.4). We make these distinctions to reduce how much our preconceptions affect our analysis (or at least make it more obvious to future researchers).

See #176

2.1 Test Approaches

Software testing is “the process of executing a program with the intent of finding errors” (Peters and Pedrycz, 2000, p. 438). For each test, the main steps are to:

OG Myers 1976

1. identify the goal(s) of the test,
2. decide on an approach,
3. develop the tests,
4. determine the expected results,

5. run the tests, and
6. compare the expected results to the actual results (p. 443).

The end goal is to evaluate “some aspect of the system or component” based on the results of step 6 (ISO/IEC and IEEE, 2022, p. 10; 2021, p. 6; 2017, p. 465). When this evaluation reveals errors, “the faults causing them are what can and must be removed” (Washizaki, 2024, p. 5-3).

Of course, the approach chosen in step 2 influences what kinds of test cases should be developed and executed in later steps, so it is important that test approaches are defined correctly, consistently, and unambiguously. A “test approach” is a “high-level test implementation choice” (ISO/IEC and IEEE, 2022, p. 10) used to “pick the particular test case values” (2017, p. 465) used in step 5. The only approach that can “fully” test a system (exhaustive testing) is infeasible in most non-trivial situations (2022, p. 4; Washizaki, 2024, p. 5-5; Peters and Pedrycz, 2000, pp. 439, 461; van Vliet, 2000, p. 421), so multiple approaches are needed (ISO/IEC and IEEE, 2022, p. 18) to “suitably cover any system” (p. 33). This is why this process should be repeated and there are so many test approaches described in the literature.

2.1.1 Approach Categories

Since there are so many test approaches, it is helpful to categorize them. The literature provides many ways to do so, but perhaps the most widely used is the one given by ISO/IEC and IEEE (2022). This schema divides test approaches into levels, types, techniques, and practices (2022, Fig. 2; see Table 2.1). These categories seem to be pervasive throughout the literature, particularly “level” and “type”. For example, six non-IEEE sources also give unit testing, integration testing, system testing, and acceptance testing as examples of test levels (Washizaki, 2024, pp. 5-6 to 5-7; Hamburg and Mogyorodi, 2024; Perry, 2006, pp. 807–808; Peters and Pedrycz, 2000, pp. 443–445; Kuļešovs et al., 2013, p. 218; Gerrard, 2000a, pp. 9, 13). These categories seem to be orthogonal based on their definitions and usage. For example, “a test type can be performed at a single test level or across several test levels” (ISO/IEC and IEEE, 2022, p. 15; 2021, p. 7). We may assess this assumption more rigorously in the future, but for now, it implies that a specific test approach can be derived by combining multiple test approaches from different categories; see Appendix A.6 for more detailed discussion. Because of their widespread use and their usefulness in dividing the domain of software testing into more manageable subsets, we use these categories throughout this thesis. They can be loosely described by what they specify as follows:

- **Level:** What code is tested
- **Practice:** How the test is structured and executed
- **Technique:** How inputs and/or outputs are derived
- **Type:** Which software quality is evaluated

Table 2.1: Categories of test approaches given by ISO/IEC and IEEE.

Term	Definition	Examples
Test Level ^a	A stage of testing “typically associated with the achievement of particular objectives and used to treat particular risks”, each performed in sequence (ISO/IEC and IEEE, 2022, p. 12; 2021, p. 6) with their “own documentation and resources” (2017, p. 469)	unit/component testing, integration testing, system testing, acceptance testing (2022, p. 12; 2021, p. 6; 2017, p. 467)
Test Practice	A “conceptual framework that can be applied to ... [a] test process to facilitate testing” (2022, p. 14; 2017, p. 471; OG IEEE 2013)	scripted testing, exploratory testing, automated testing (2022, p. 20)
Test Technique ^b	A “procedure used to create or select a test model, identify test coverage items, and derive corresponding test cases” (2022, p. 11; similar in 2017, p. 467) that “generate evidence that test item requirements have been met or that defects are present in a test item” (2021, p. vii)	equivalence partitioning, boundary value analysis, branch testing (2022, p. 11)
Test Type	“Testing that is focused on specific quality characteristics” (2022, p. 15; 2021, p. 7; 2017, p. 473; OG IEEE 2013)	security testing, usability testing, performance testing (2022, p. 15; 2017, p. 473)

^a Also called “test phase” (see Synonym Flaw 2) or “test stage” (see Synonym Flaw 4).^b Also called “test design technique” (ISO/IEC and IEEE, 2022, p. 11; Hamburg and Mogyorodi, 2024).

While ISO/IEC and IEEE (2022)’s schema includes “static testing” as a test approach category, we omit it as it seems non-orthogonal to the others and thus less helpful for grouping test approaches. Other categorization schemas (see Section 4.1.2) may consider static testing orthogonal, and some may consider it out-of-scope entirely (see Appendix A.3)! We also introduce a “supplemental” category of “artifact” since some terms can refer to the application of a test approach *as well as* the resulting document(s). Therefore, we do *not* consider approaches categorized as an artifact *and* another category as flaws in Section 6.2.1. Finally, we also record the test category of “process” for completeness, although this seems to be a higher level classification and these approaches will likely be excluded during later analysis.

See #44, #119,
and #39

See #119

See #52

2.1.2 Synonym Relations

The same approach often has many names. For example, “specification-based testing” is also called “black-box testing” (ISO/IEC and IEEE, 2022, p. 9; 2021, p. 8; 2017, p. 431; Washizaki, 2024, p. 5-10; Hamburg and Mogyorodi, 2024; Firesmith, 2015, p. 46¹; Sakamoto et al., 2013, p. 344; van Vliet, 2000, p. 399). Throughout our work, we use the terms “specification-based testing” and “structure-based testing” to articulate the source of the information for designing test cases, but a team or project also using grey-box testing may prefer the terms “black-box” and “white-box testing” for consistency.

We can formally define the synonym relation S on the set T of terms used by the literature to describe test approaches based on how synonyms are used in natural language. S is symmetric and transitive, and although pairs of synonyms in natural language are implied to be distinct, a relation that is symmetric and transitive is provably reflexive; this implies that all terms are trivially synonyms of themselves. Since S is symmetric, transitive, *and* reflexive, it is an equivalence relation, reflecting the role of synonyms in natural language where they can be used interchangeably. While synonyms may emphasize different aspects or express mild variations, their core meaning is nevertheless the same.

2.1.3 Parent-Child Relations

Many test approaches are multi-faceted and can be “specialized” into others; for example, load testing and stress testing are some of the subtypes of performance-related testing (described in more detail in Section 7.4). We refer to these “specializations” as “children” or “subapproaches” of their multi-faceted “parent(s)”. This nomenclature also extends to approach categories (such as “subtype”; see Section 2.1.1 and Table 2.1) and software qualities (“subquality”; see Section 3.3).

We can formally define the parent-child relation P on the set T of terms used by the literature to describe test approaches based on directed relations between approach pairs. This relation should be irreflexive, asymmetric, and transitive, making it a strict partial order. A given child approach c may have more than one

¹Firesmith (2015) excludes the hyphen, calling it “black box testing”.

parent approach p . This relation often manifests when a “well-understood” test approach p is decomposed into smaller, independently performable approaches c_1, \dots, c_n , each with its own focus or nuance. This is frequently the case for hierarchies of approaches given in the literature (ISO/IEC and IEEE, 2022, Fig. 2; 2021, Fig. 2; Firesmith, 2015).

Another possible way for a parent-child relation to occur is when the completion of p indicates that “sufficient testing has been done” in regards to c (van Vliet, 2000, p. 402). While this only “compares the thoroughness of test techniques, not their ability to detect faults” (p. 434), it is sufficient to justify a parent-child relation between the two approaches. These relations may also be represented as hierarchies (ISO/IEC and IEEE, 2021, Fig. F.1; van Vliet, 2000, Fig. 13.17).

2.2 Flaws

Ideally, software testing literature would describe test approaches correctly, completely, consistently, and modularly, but this is not the case in reality. We refer to any instance where one of these ideals is violated as a “flaw”². We classify flaws by both their “manifestations” (*how* information is wrong; see Section 2.2.1) and their “domains” (*what* information is wrong; see Section 2.2.2). These are orthogonal classifications, since each flaw *manifests* in a particular *domain*, which we track by assigning each flaw one “key” for each classification. We list these keys in Tables 2.2 and 2.3, respectively, and use them to count and analyze the flaws we uncover; for example, we present the number of flaws according to these two views in Tables 6.1 and 6.2, respectively. We also introduce some terms to be used when discussing flaws based on how many sources contribute to them (Section 2.2.3).

2.2.1 Flaw Manifestations

Perhaps the most obvious example of something being “wrong” with the literature is that a piece of information it presents is incorrect—“wrong” in the literal sense. However, if our standards for correctness require clarity, consistency, and robustness, then there are many ways for a flaw to manifest. This is one view we take when observing, recording, and analyzing flaws: *how* information is “wrong”. We observe the “manifestations” described in Table 2.2 throughout the literature, and give each a unique key for later analysis and discussion. We list them in descending order of severity, although this is partially subjective. While some may disagree with our ranking, it is clear that information being incorrect is worse than it being repeated. Our ordering has the benefit of serving as a “flowchart” for classifying

²When picking a word to describe these violations, we wanted to avoid words that are “overloaded with too many meanings” like “error” and “fault” (Washizaki, 2024, p. 12-3; see Chapter 1 for more detailed discussion). A small literature review revealed that established standards (see Section 3.1.1) only use the term “flaw” to refer to requirements (ISO/IEC and IEEE, 2022, p. 38), design (p. 43), “system security procedures ... and internal controls” (IEEE, 2012, p. 194), or code itself (p. 92). Washizaki (2024, p. 7-9) also uses the word “flaw” in a way that aligns with our goal of analyzing what is wrong with software engineering’s testing literature.

Table 2.2: Observed flaw manifestations.

Manifestation	Description	Key
Mistake	Information is incorrect	WRONG ^a
Omission	Information that should be included is not	MISS ^b
Contradiction	Information from multiple places conflicts	CONTRA
Ambiguity	Information is unclear	AMBI
Overlap	Information is nonatomic or used in multiple contexts	OVER
Redundancy	Information is redundant	REDUN

^a We use **WRONG** here to avoid clashing with **MISS**.

^b We use **MISS** here to be more meaningful in isolation, as it implies the synonym of “missing”; **OMISS** would be unintuitive and **OMIT** would be inconsistent with the keys being adjective-based.

flaws. For example, if a piece of information is not intrinsically incorrect, then there are five remaining manifestation types for the flaw. Note that some flaws involve information from multiple sources (contradictions and overlaps in particular). We do not categorize these flaws as “mistakes” if finding the ground truth requires analysis that has not been performed yet.

Table 2.3: Observed flaw domains.

Domain	Description	Key
Categories	Approach categories, defined in Section 2.1.1	CATS
Synonyms	Synonym relations, defined in Section 2.1.2	SYNS
Parents	Parent-child relations, defined in Section 2.1.3	PARS
Definitions	Definitions given to terms	DEFS
Labels	Labels or names given to terms	LABELS
Scope	Scope of the information	SCOPE
Traceability	Records of the source(s) of information	TRACE

2.2.2 Flaw Domains

Another way to categorize flaws is by *what* information is wrong, which we call the flaw’s “domain”. We describe those we observe in Table 2.3, and tracking these uncovers which knowledge domains are less standardized (and should therefore be approached with more rigour) than others. We explicitly define some of these domains in previous sections and thus present them in that same order. Despite their nuance, the remaining domains are relatively straightforward, so we define them briefly as follows instead of defining them more rigorously in their own sections.

Terms can be thought of as definition-label pairs, but there is a meaningful distinction between definition flaws and label flaws. Definition flaws are quite self-explanatory, but label flaws are harder to detect, despite occurring independently. Examples of label flaws include terms that share the same acronym or contain typos or redundant information. Sometimes, an author may use one term when they mean another. One could argue that their “internal” definition of the term is the cause of this mistake, but we consider this a label flaw where the wrong label is used as we would change the *label* to fix it. Additionally, some information is presented with an incorrect scope and sometimes should not have been included at all! Finally, some traceability information is flawed, such as how one document cites another or even what information is included *within* a document.

2.2.3 Additional Flaw Terminology

Some flaws involve information from more than one source, but referring to this as a “flaw between two sources” is awkward. We instead refer to this kind of flaw more intuitively as an “inconsistency” between the sources. This clearly indicates that there is disagreement between the sources, but also does not imply that either one is correct—the inconsistency could be with some ground truth if *neither* source is correct!

Other flaws only involve one source, but we make a distinction between “self-contained” flaws and “internal” flaws. Self-contained flaws are those that manifest by comparing a document to an assertion of ground truth. These may appear once in a document or consistently throughout it. Sometimes, these do not require an

explicit comparison to ground truth; these often include omissions as the lack of information is contained within a single source and does not need to be cross-checked against an assertion of ground truth. On the other hand, internal flaws arise when a document disagrees with itself by containing two conflicting pieces of information; this includes many contradictions and overlaps. Internal flaws can even occur on the same page, such as when a source gives the same acronym to two distinct terms (see Overlap 8 and Overlap 9)!

2.3 Explicitness

When information is written in natural language, a considerable degree of nuance can get lost when interpreting or using it. For example, a source may provide data from which the reader can logically draw a conclusion, but may not state or prove this conclusion explicitly. Most information provided by sources we investigate is given explicitly; all sources cited throughout this thesis support their respective claims explicitly unless specified otherwise, usually via one of the keywords given below. In the cases where information is *not* explicit, we capture it when citing sources using (at least) one of the following keywords: “implied”, “can be”, “sometimes”, “should be”, “ideally”, “usually”, “most”, “likely”, “often”, “if”, and “although”. These keywords indicate the presence of some nuance, which we call “explicitness”, or how explicit a piece of information is (or is *not*). While these keywords often appear directly within the literature, we also use them to track explicitness in our test approach glossary to provide a more complete summary of the state of software testing literature without getting distracted by less relevant details. We outline this process in more detail in Section 3.4, but it is important to note that we use “implicit” (as well as “implied by” when describing sources of information) to refer to any instance of “not explicit” information for clarity. Any kind of information can be implicit, including the names, definitions, categories (see Section 2.1.1), synonyms (see Section 2.1.2), and parents (see Section 2.1.3) of identified test approaches.

As our research focuses on the flaws present in the literature, the explicitness of information affects how seriously we take it. We call flaws based on explicit information “objective”, since they are self-evident in the literature. On the other hand, we call flaws based on implicit information “subjective”, since some level of judgement is required to assess whether these flaws are actually problematic. By looking for the indicators of uncertainty mentioned above, we can automatically detect subjective flaws when generating graphs and performing analysis (see Sections 5.1 and 5.2, respectively).

Throughout our research, we also infer some information through “surface-level” analysis that follows straightforwardly but is not stated, explicitly or otherwise, by a source. Although these data originate from our judgement, we document them for completeness, using the phrase “inferred from” when relevant. All data in our glossaries without a citation are inferred, such as algebraic testing (Peters and Pedrycz, 2000, Fig. 12.2) being a child of mathematical-based testing. Additionally, some inferences are based on information given by a source,

Later: Ensure this is up to date

Q #1: Does this make sense in isolation or do I need more detail?

See #176

which we cite alongside these inferences. For example, Gerrard describes large scale integration testing and legacy system integration testing in (2000b, p. 30) and (2000a, Tab. 2; 2000b, Tab. 1), respectively. While he never explicitly says so, we infer that these approaches are children of integration testing and system integration testing, respectively. Similarly, some test approaches appear to be combinations of other (seemingly orthogonal) approaches (described in more detail in Appendix A.6), from which we can extrapolate other test approaches. For example, Moghadam (2019) uses the phrase “machine learning-assisted performance testing”; since performance testing is a known test approach, we infer the existence of the test approach “machine learning-assisted testing” and include it in our test approach glossary as such. We also infer that child approaches inherit their parents’ categories (see Section 2.1.1).

2.4 Credibility

In the same way we distinguish between the explicitness of information from different sources, we also wish to distinguish between the “explicitness” of the sources themselves! Of course, we do not want to overload terms, so we define a source as more “credible” if it:

- has gone through a peer-review process,
- is written by numerous, well-respected authors,
- cites a (comparatively) large number of sources, and/or
- is accepted and used in the field of software.

Sources may meet only some of these criteria, so we use our judgement (along with the format of the sources themselves) when comparing them (cf. Section 3.1).

Chapter 3

Methodology

Since we wish to analyze software testing literature, we collect data (both correct and incorrect) from a wide variety of relevant documents, focusing on test approaches and supporting information. This results in a big messy glossary of software approaches, some glossaries of supplementary terms, and a list of flaws. To ensure this data can be analyzed and expanded thoroughly and consistently, we need a process that can be repeated for future developments in the field of software testing, or by independent researchers seeking to verify our work. Our methodology for documenting the current (messy) state of software testing terminology consists of:

1. Identifying authoritative sources on software testing and “snowballing” from them (Section 3.1)
2. Identifying all test approaches (Section 2.1) and related testing terms that are used repeatedly and/or have complex definitions (Section 3.2)
3. Recording all relevant data (Section 3.3), including implicit data (Section 3.4), for each term identified in step 2; for test approaches, these are comprised of:
 - (a) Names
 - (b) Categories¹ (Section 2.1.1)
 - (c) Definitions
 - (d) Synonyms (Section 2.1.2)
 - (e) Parents (Section 2.1.3)
 - (f) Flaws (Section 2.2) (in separate documents)
 - (g) Other relevant notes (prerequisites, uncertainties, other sources, etc.)
4. Repeating steps 1 to 3 for any missing or unclear terms (Section 3.5) until some stopping criteria (Section 3.6)

¹A single test approach with more than one category (for example, A/B Testing in Table 1.1) often indicates an underlying flaw; see Section 6.2.1.

3.1 Sources

As there is no single authoritative source on software testing terminology, we need to look at many sources to observe how this terminology is used in practice. Since we are particularly interested in software engineering, we start from the vocabulary document for systems and software engineering (ISO/IEC and IEEE, 2017) and two versions of the Guide to the SoftWare Engineering Body Of Knowledge (SWEBOK Guide)—the newest one (Bourque and Fairley, 2014) and one submitted for public review² (Washizaki, 2024). To gather further sources, we then use a version of “snowball sampling”, which “is commonly used to locate hidden populations ... [via] referrals from initially sampled respondents to other persons” (Johnson, 2014). We apply this concept to “referrals” between sources. For example, Hamburg and Mogyorodi (2024) cite Gerrard and Thompson (2002) as the original source for their definition of “scalability” (see Section 7.3); we verified this by looking at this original source.

For ease of discussion and analysis, we group the complete set of sources into “tiers” based on their format, method of publication, and our heuristic of credibility (defined in Section 2.4). In order of descending credibility, we define the following tiers:

1. established standards (Section 3.1.1),
2. terminology collections (Section 3.1.2),
3. textbooks (Section 3.1.3), and
4. papers and other documents (Section 3.1.4).

We provide a summary of how many sources comprise each tier in Figure 3.1 and list all sources in each tier in Appendix B. The “papers” tier is quite large since we often “snowball” on terminology itself when a term requires more investigation (e.g., its definition is missing or unclear). This includes performing a miniature literature review on this subset to “fill in” missing information (see Section 3.5) and potentially fully investigating these additional sources, as opposed to just the original subset of interest, based on their credibility and how much extra information they provide. We use standards the second most frequently due to their high credibility and broad scope; for example, the glossary portion of ISO/IEC and IEEE (2017) has 514 pages! Using these standards allows us to record many test approaches in a similar context from a source that is widely used and well-respected.

3.1.1 Established Standards

These are documents written for the field of software engineering by reputable standards bodies, namely ISO, the International Electrotechnical Commission (IEC),

²(Washizaki, 2024) has been published since we investigated these sources; if time permits, we will revisit this published version.



Figure 3.1: Summary of how many sources comprise each source tier.

and IEEE. Their purpose is to “encourage the use of systems and software engineering standards” and “collect and standardize terminology” by “provid[ing] definitions that are rigorous, uncomplicated, and understandable by all concerned” (ISO/IEC and IEEE, 2017, p. viii) “that can be used by any organization when performing any form of software testing” (2022, p. vii; similar in 2016, p. ix). For these reasons, they are the most credible sources. However, this does not imply perfection, as we identify 51 flaws within these standards (see Tables 6.1 and 6.2)! Only standards for software development and testing are in scope for this research (see Chapter 1).

3.1.2 Terminology Collections

These are collections of software testing terminology built up from multiple sources, such as the established standards outlined in Section 3.1.1. For example, the SWE-BOK Guide is “proposed as a suitable foundation for government licensing, for the regulation of software engineers, and for the development of university curricula in software engineering” (Kaner et al., 2011, p. xix). Even though it is “published by the IEEE Computer Society”, it “reflects the current state of generally accepted, consensus-driven knowledge derived from the interaction between software engineering theory and practice” (Washizaki, 2025). Due to this combination of IEEE standards and state-of-the-practice observations, we designate it as a collection of terminology as opposed to an established standard. Collections such as this are often written by a large organization, such as the International Software Testing Qualifications Board (ISTQB), but not always. Firesmith (2015)’s taxonomy presents relations between many test approaches and Doğan et al. (2014)’s literature review cites many sources from which we can “snowball” if desired (see Section 3.1), so we include them in this tier as well.

3.1.3 Textbooks

We consider textbooks to be more credible than papers (see Section 3.1.4) because they are widely used as resources for teaching software engineering and industry frequently uses them as guides. Although textbooks have smaller sets of authors, they follow a formal review process before publication. Textbooks used at McMas-

ter University (Patton, 2006; Peters and Pedrycz, 2000; van Vliet, 2000) served as the original (albeit ad hoc and arbitrary) starting point of this research, and we investigate other books as they arise. For example, Hamburg and Mogyorodi (2024) cite Gerrard and Thompson (2002) as the original source for their definition of “scalability” (see Section 7.3); we verified this by looking at this original source.

Present tense?

3.1.4 Papers and Other Documents

The remaining documents all have much smaller sets of authors and are much less widespread than those in higher source tiers. While most of these are journal articles and conference papers, we also include the following document types. Some of these are not peer-reviewed works but are still useful for observing how terms are used in practice:

See #89

- Report (Kam, 2008; Gerrard, 2000a;b)
- Thesis (Bas, 2024)
- Website (LambdaTest, 2024; Pandey, 2023)
- Booklet (Knüvener Mackert GmbH, 2022)
- ChatGPT (GPT-4o) (2024) (with its claims supported by Rus et al. (2008))

3.2 Identifying Relevant Terms

Before we can consistently track software testing terminology used in the literature, we must first determine what to record. We use heuristics to guide this process to increase confidence that we identify all relevant terms, paying special attention to the following when investigating a new source:

- glossaries, taxonomies, hierarchies, and lists of terms,
- testing-related terms (e.g., terms containing “test(ing)”, “review(s)”, “audit(s)”, “attack(s)”, “validation”, or “verification”),
- terms that had emerged as part of already-discovered test approaches, *especially* those that were ambiguous or prompted further discussion (e.g., terms containing “performance”, “recovery”, “component”, “bottom-up”, “boundary”, or “configuration”), and
- terms that imply test approaches³, including:
 - coverage metrics that may imply related test techniques,

See #55

³Since these methods for deriving test approaches only arose as research progressed, some examples would have been missed during the first pass(es) of resources investigated earlier in the process. While reiterating over them would be ideal, this may not be possible due to time constraints.

- software qualities that may imply related test types, and
- software requirements that may imply related test approaches.

We apply these heuristics to the entirety of most investigated sources, especially established standards (see Section 3.1.1). However, we only partially investigate some sources, such as those chosen for a specific area of interest or based on a test approach that was determined to be out-of-scope. These include the following sources as described in Section 3.5: ISO (2022; 2015); Dominguez-Pumar et al. (2020); Pierre et al. (2017); Trudnowski et al. (2017); Yu et al. (2011); Tsui (2007); Goralski (1999).

3.3 Recording Relevant Terms

Once we have identified which terms from the literature are relevant, we can then track them consistently. We do this by building glossaries, mainly focusing on the one most central to our research—our test approach glossary—where we record all test approaches we identify in the literature. We give each test approach its own row where we record its name and any given definitions, categories (see Section 2.1.1), synonyms (see Section 2.1.2), and parents (see Section 2.1.3) (along with any other notes, such as questions, prerequisites, and other resources to investigate) following the procedure in Figure 3.2. Note that only the name and category fields are required; all other fields may be left blank, although a lack of definition indicates that the approach should be investigated further to see if its inclusion is meaningful (see Section 3.5). When recording flawed data, we record it in our glossary as dubious information (see Section 3.4) and/or in a separate document if it is clearly flawed (see Section 2.2).

For example, ISO/IEC and IEEE (2022, p. 1, 36) define “A/B testing” as shown in Figure 3.3. Since we do not yet have a row for “A/B testing” in our test approach glossary, we apply our procedure as follows:

Q #3: Should glossary headers be capitalized?

1. Create a new row with the name “A/B testing” and the category “Approach” (see Section 2.1.1).
2. Record the synonym “Split-Run Testing” (see Section 2.1.2).

Q #4: Are these references helpful or distracting?

3. Record the parent “Statistical Testing” (see Section 2.1.3).
4. Record the definition “Testing ‘that allows testers to determine which of two systems or components performs better’ ”; note that we abstract away information that we have previously captured (i.e., its synonym and parent).



Figure 3.2: Procedure for recording test approaches in our glossary.

3.1

A/B testing

split-run testing

statistical *testing* (3.131) approach that allows testers to determine which of two systems or components performs better

Figure 3.3: ISO/IEC and IEEE (2022, p. 1)’s glossary entry for “A/B testing”.

Note that we accompany all of these entries (excluding its name) with the citation “(ISO/IEC and IEEE, 2022, pp. 1, 36)”. This document provides more information which we capture as follows:

1. Record the note “It ‘can be time-consuming, although tools can be used to support it’, ‘is a means of solving the test oracle problem by using the existing system as a partial oracle’, and is ‘not a test case generation technique as test inputs are not generated’ ” (p. 36).
2. Replace the category of “Approach” with the more specific “Practice” (Fig. 2) (see Section 2.1.1); note that this is consistent with the exclusion of “Technique” as a possible category for this approach from page 36.

As we investigate other sources, we learn more about this approach. Firesmith (2015, p. 58) includes it in his taxonomy as shown in Figure 3.4. We add to our entry for “A/B Testing” as follows:

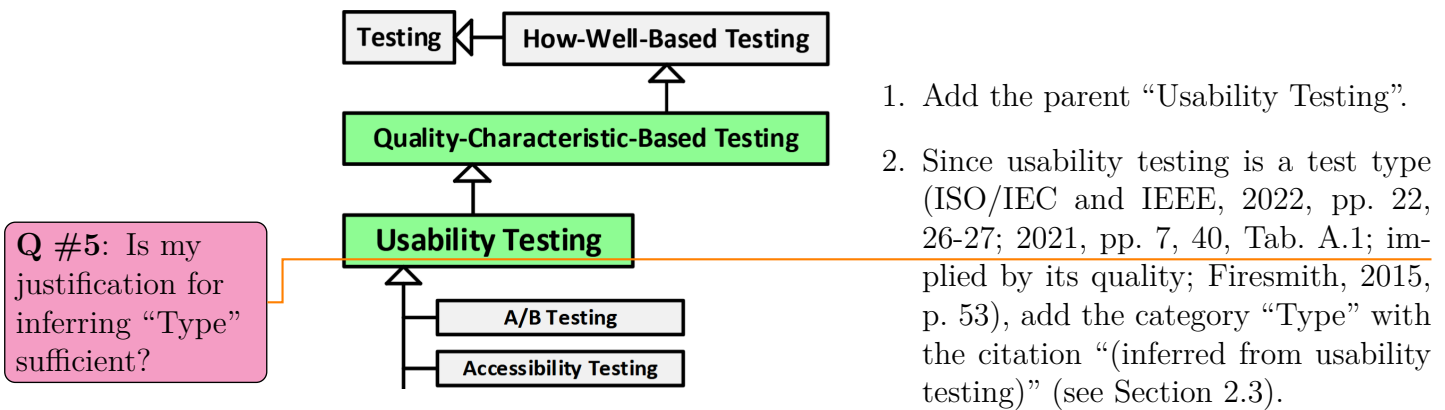


Figure 3.4: A/B testing’s inclusion in Firesmith (2015, p. 58)’s taxonomy.

This second change violates our assumption that categories are orthogonal (see Section 2.1.1), so we consider this to be an inferred flaw (defined in Section 2.3) that we automatically detect and document (see Section 5.2.1 and Table 6.6, respectively). This results in the corresponding row in Table 1.1 although we exclude the “Notes” column for brevity.

We use this same procedure to track software qualities and supplementary terminology that is either shared by multiple approaches or too complicated to explain inline. We create a separate glossary for both qualities and supplementary terms, each with a similar format to our test approach glossary. Since these terms do not have categories, the process of recording them is much simpler,

only requiring us to record the name, definition, and synonym(s) of these terms, along with any additional notes. The only new information we capture is the “precedence” for a software quality to have an associated test type, since each test type measures a particular software quality (see Table 2.1). These precedences are instances where a given software quality is related to, is covered by, or is a child, parent, or prerequisite of another quality with an associated test type, as given by the literature.

Tracking information about software qualities helps us investigate the literature more thoroughly, since these data may become relevant based on information from other yet-uninvestigated sources. When the literature mentions (or implies) a test approach that corresponds to a software quality we have recorded, we first follow the procedure given in Figure 3.2 with the information provided in the source that mentions it. We then repeat this process with the relevant data from our quality glossary, then remove its entry (since it has now been “upgraded” to a test type).

3.4 Implicit Information

As described in Section 2.3, the use of natural language introduces significant nuance that we need to document. Keywords such as “implied”, “can be”, “sometimes”, “should be”, “ideally”, “usually”, “most”, “likely”, “often”, “if”, and “although” indicate that information from the literature is *not* explicit. We find the following non-mutually exclusive cases of “implicit” information from the literature:

1. The information follows logically.

Firesmith (2015, pp. 53–58) lists a set of test approaches that are “based on the[ir] associated quality characteristic[s] and ... associated quality attributes” (p. 53). This matches our definition of “test type”, but this term is used more loosely in this document to refer to different kinds of testing that we would call “test approaches” (see Section 2.1.1). We therefore consider Firesmith’s categorizations of “test type” to be implicit (such as in Tables 6.4 and 6.6).

2. The information is not universal.

ISO/IEC and IEEE (2017, p. 372, emphasis added) define “regression testing” as “testing required to determine that a change to a system component has not adversely affected *functionality, reliability or performance* and has not introduced additional defects”. While reliability testing, for example, is not *always* a subset of regression testing (since it may be performed in other ways), it *can be* accomplished by regression testing. This means that these parent-child relations (defined in Section 2.1.3) between these pairs of approaches only exist *sometimes*, but this qualifier is given implicitly.

3. The information is conditional.

As a more specific case of information not being universal, sometimes prerequisites must be satisfied for information to apply. For example, branch

OG ISO/IEC,
2014

Table 3.1: Breakdown of implicit keywords used for recording and analyzing implicitness (defined in Section 2.3).

Keyword	Follows Logically	Not Universal	Conditional	Dubious
“implied”	X			
“can be”		X	X	
“sometimes”		X	X	
“should be”		X		X
“ideally”		X		X
“usually”		X		
“most”		X		
“likely”	X	X		X
“often”		X		
“if”		X	X	X
“although”				X
“ (Testing)”	X			

condition combination testing is equivalent to (and is therefore implied to be a synonym of) exhaustive testing *if* “each subcondition is viewed as a single input” (Peters and Pedrycz, 2000, p. 464). Likewise, statement testing can be used for (and is therefore implied to be a child of) unit testing *if* there are “less than 5000 lines of code” (p. 481).

OG Miller et al.,
1994

4. The information is dubious.

This happens when there is reason to doubt the information provided. If a source claims one thing that is not true, related claims lose credibility. For example, the incorrect claim that “white-box testing”, “grey-box testing”, and “black-box testing” are synonyms for “module testing”, “integration testing”, and “system testing”, respectively, (see Mistake 22) casts doubt on the claim that “red-box testing” is a synonym for “acceptance testing” (Sneed and Göschl, 2000, p. 18) (see Mistake 23).

OG Hetzel88

When we encounter information that meets one of these criteria, we use an appropriate keyword to capture this nuance in our test approach glossary (see Table 3.1). This also helps us identify implicit information when performing later analysis. Despite “implicit” only describing the first of these cases, we use it (as well as “implied by” when describing sources of information) as a shorthand for all “not explicit” manifestations throughout this thesis for clarity.

Regarding the last entry in Table 3.1, if a test approach in our test approach glossary has a name ending in “ (Testing)” (space included), then the word “Testing” might not be part of its name *or* it might not be a test approach at all! For example, the term “legacy system integration” is used in Gerrard (2000a,

pp. 12–13, Tab. 2; 2000b, Tab. 1), but the more accurate “legacy system integration testing” is used in (2000b, pp. 30–31). In other cases where a term is *not* explicitly labelled as “testing”, we add the suffix “(Testing)” (when it makes sense to do so) and consider the test approach to be implied.

3.5 Undefined Terms

The literature mentions many software testing terms without defining them. While this includes test approaches, software qualities, and more general software terms, we focus on the former as the main focus of our research. In particular, ISO/IEC and IEEE (2022) and Firesmith (2015) name many undefined test approaches. Once we exhaust the standards in Section 3.1.1, we perform miniature literature reviews on these subsets to “fill in” the missing definitions (along with any relations), essentially “snowballing” on these terms. This process uncovers even more approaches, although some are out of scope, such as EManations SECurity (EMSEC) testing, aspects of Orthogonal Array Testing (OAT) and loop testing (see Appendix A.1), and HTML testing (see Appendix A.5). We investigate the following terms (and their respective related terms) in the sources given:

- **Assertion Checking:** Lahiri et al. (2013); Chalin et al. (2006); Berdine et al. (2006)
- **Loop Testing**⁴: Dhok and Ramanathan (2016); Godefroid and Luchaup (2011); Preuße et al. (2012); Forsyth et al. (2004)
- **EMSEC Testing:** Zhou et al. (2012); ISO (2021)
- **Asynchronous Testing:** Jard et al. (1999)
- **Performance(-related) Testing:** Moghadam (2019)
- **Web Application Testing:** Doğan et al. (2014); Kam (2008)
 - **HTML Testing:** Choudhary et al. (2010); Sneed and Göschl (2000); Gerrard (2000b)
 - **Document Object Model (DOM) Testing:** Bajammal and Mesbah (2018)
- **Sandwich Testing:** Sharma et al. (2021); Sangwan and LaPlante (2006)
- **Orthogonal Array Testing**⁵: Mandl (1985); Valcheva (2013)
- **Backup Testing**⁶: Bas (2024)

⁴ISO (2022) and ISO (2015) were used as reference for terms but not fully investigated, Pierre et al. (2017) and Trudnowski et al. (2017) were added as potentially in scope, and Dominguez-Pumar et al. (2020) and Goralski (1999) were added as out-of-scope examples.

⁵Yu et al. (2011) and Tsui (2007) were added as out-of-scope examples.

⁶See Section 6.5.



Figure 3.5: Breakdown of how many test approaches are undefined.

- **SoS (Integration) Testing:** ISO/IEC and IEEE (2019b)

Applying our procedure shown in Figure 3.2 to these sources uncovers 87 new approaches and 71 new definitions. These definitions are either for existing undefined approaches or new uncovered approaches; while not every new approach is presented alongside a definition, if we assume that each of these definitions is for a new approach, we can deduce that about 82% of added test approaches are defined. This indicates that this procedure leads to a higher proportion of defined terms (63% vs. 66%), as shown in Figure 3.5, which helps verify that our procedure constructively uncovers *and* defines new terminology. Further iterating on it would continue this trend, creating something approaching a complete taxonomy.

3.6 Stopping Criteria

Unfortunately, we cannot continue looking for new approaches indefinitely. We therefore need a “stopping criteria” to let us know when we are “finished” looking for test approaches in the literature. Our original plan was to iterate steps 2 and 3 until there are diminishing returns, implying an approach of a complete taxonomy! However, due to time constraints, this is infeasible, so our stopping point is artificially imposed. With more time, we would find definitions for all terms we uncover as described in Section 3.5.

In addition to these undefined terms, some terms do not appear in the literature at all! While most test approaches arise as a result of our snowballing approach, we each have preexisting knowledge of what test approaches exist (a form of experience-based testing, if you will). As an example, we are surprised that property-based testing is not mentioned in any sources investigated, even considering it as a potential stopping point during our research. Test approaches such as these that arise independently of snowballing may serve as starting points for continued research if we do not find them in the literature using our iterative approach. The following terms come from previous knowledge, conversations with colleagues, research for other projects, or ad hoc cursory research to see what other

test approaches exist:

- | | |
|---------------------------------|-----------------------------------|
| 1. Chaos engineering | 9. Lunchtime attacks ^a |
| 2. Chosen-ciphertext attacks | 10. Parallel testing |
| 3. Concolic testing | 11. Property-based testing |
| 4. Concurrent testing | 12. Pseudo-random bit testing |
| 5. Destructive testing | 13. Rubber duck testing |
| 6. Dogfooding | 14. Scream testing |
| 7. Implementation-based testing | 15. Shadow testing |
| 8. Interaction-based testing | 16. Situational testing |

^aIn previous meetings, Dr. Smith mentioned that with the number of test approaches that suggest that people just like to label everything as “testing”, he would not be surprised if something like “Monday morning testing” existed. While independently researching chosen-ciphertext attacks out of curiosity, this prediction of a time-based test approach came true with “lunchtime attacks”.

Chapter 4

Threats to Validity

While we do our best to approach the software testing literature with the rigour it deserves, we make some decisions with which others may disagree. The benefit of our work being open-source means that others can make their own decisions, modify our data/code appropriately, and observe how this changes the results. Nevertheless, some potential threats to validity arise from the test approach categorization we use (Section 4.1).

4.1 Threats from Approach Categories

In Section 2.1.1, we outline our criteria for categorizing test approaches based on documents by ISO/IEC and IEEE (such as 2022, Fig. 2) and summarize it in Table 2.1. However, the terms we use are sometimes used in a general sense (as opposed to our specific definitions) which makes it hard to evaluate how consistently the literature uses them (Section 4.1.1). Furthermore, the literature provides other criteria for categorizing test approaches (Section 4.1.2); using any of these schemas would affect the category data we collect, if any, and how we use them to identify the flaws given in Section 6.2.1.

4.1.1 Overloaded Terms

A side effect of using the terms “level”, “type”, “technique”, and “practice” is that they—perhaps excluding “level”—can be used interchangeably or as synonyms for “approach”. Because natural language can be ambiguous, we need to exercise judgement when determining if these terms are being used in a general or technical sense. For example, Kam (2008, p. 45, emphasis added) defines interface testing as “an integration *test type* that is concerned with testing ... interfaces”, but since he does not define “test type”, this may not have special significance. We consider these kinds of “categorizations” to be inferred (see Section 2.3) and mark them with a question mark (?) during data collection. When we evaluate test approaches categorized more than once in Section 6.2.1, we list well-defined categorizations in Table 6.4 and inferred ones in Table 6.6.

We are not immune to using these terms sloppily either! For example, Patton

(2006, p. 88) says that if a specific defect is found, it is wise to look for other defects in the same location and for similar defects in other locations, but does not provide a name for this approach. After researching in vain, we ask ChatGPT¹ to name the “*type* of software testing that focuses on looking for bugs where others have already been found” (ChatGPT (GPT-4o), 2024, emphasis added), using the word “type” in a general sense, akin to “kind” or “subset”. Interestingly, ChatGPT “corrects” our imprecise term in its response, using the more correct term “approach” (although it may have been biased by our previous usage of these terms)!

Q #6: Is this sufficient to explain ChatGPT’s role in our methodology?

4.1.2 Alternate Approach Categorizations

While our categorization in Table 2.1 is used fairly consistently, other sources (such as Washizaki, 2024; Barbosa et al., 2006) propose similar yet distinct categories that clash or overlap with them. We give these alternate categorizations in Table 4.1, which could simply map to their “IEEE Equivalents” or could provide new perspectives and be useful in some contexts, either in place of or in tandem with ours.

Similarly, we find many other criteria for categorizing test approaches in the literature. In general, these are defined less systematically but are more fine-grained, seeming to “specialize” our categories from Table 2.1. The existence of these categorizations is not inherently wrong, as they may be useful for specific teams or in certain contexts. For example, functional testing and structural testing “use different sources of information and have been shown to highlight different problems”, and deterministic testing and random testing have “conditions that make one approach more effective than the other” (Washizaki, 2024, p. 5-16). Unfortunately, even these alternate categories are not used consistently (see Category Flaw 3)!

We include the most prominent of these alternate categorizations in Table 4.2 for completeness. In each row, the source given lists the example approaches and categorizes them as the given “Parent IEEE Category” unless stated otherwise (in some cases, the source gives additional approaches that we omit for brevity). For example, in the first row, ISO/IEC and IEEE (2022, pp. 22, 35) categorize both manual testing and automated testing as test practices. Note that since “approach” is a catch-all category (see Section 2.1), it does not require an explicit categorization, and that these categorizations may be flawed, as stated in the provided footnotes.

¹We do *not* take ChatGPT’s output to be true at face value; this approach seems to be called “defect-based testing” based on the principle of “defect clustering” (ChatGPT (GPT-4o), 2024), which Rus et al. (2008) support.

Table 4.1: Categories of testing given by other sources.

Term	Definition	Examples	IEEE Equivalent
Level (objective-based) ^a	Test levels based on the purpose of testing (Washizaki, 2024, p. 5-6) that “determine how the test suite is identified ... regarding its consistency ... and its composition” (p. 5-2)	conformance testing, installation testing, regression testing, performance testing, security testing (Washizaki, 2024, pp. 5-7 to 5-9)	Type?
Phase	none given by Perry (2006) or Barbosa et al. (2006)	unit testing, integration testing, system testing, regression testing (Perry, 2006, p. 221; Barbosa et al., 2006, p. 3)	Level
Procedure	The basis for how testing is performed that guides the process; “categorized in [to] testing methods, testing guidances ^b and testing techniques” (Barbosa et al., 2006, p. 3)	none given generally by Barbosa et al. (2006); see “Technique”	Approach
Process	“A sequence of testing steps” (Barbosa et al., 2006, p. 2) “based on a development technology and ... paradigm, as well as on a testing procedure” (p. 3)	none given by Barbosa et al. (2006)	Practice
Stage	An alternative to the “traditional ... test stages” based on “clear technical groupings” (Gerrard, 2000a, p. 13)	desktop development testing, infrastructure testing, post-deployment monitoring (Gerrard, 2000a, p. 13)	Level
Technique	“Systematic procedures and approaches for generating or selecting the most suitable test suites” (Washizaki, 2024, p. 5-10)	specification-based testing, structure-based testing, fault-based testing ^c (Washizaki, 2024, pp. 5-10, 5-13 to 5-15)	Technique

^a See Synonym Flaw 4.^b Testing methods and guidances are omitted from this table since Barbosa et al. (2006) do not define or give examples of them.^c Synonyms for these examples are used by Souza et al. (2017, p. 3; OG Mathur, 2012) and Barbosa et al. (2006, p. 3).

Table 4.2: Alternate categorizations found in the literature.

Categorization Basis	Example Approaches	Parent ^a IEEE Category ^b	Source
Level of Automation	Manual Testing Automated Testing	Practice ^c	(ISO/IEC and IEEE, 2022, pp. 22, 35)
Execution of Code ^d	Static Testing Dynamic Testing	Approach	(ISO/IEC and IEEE, 2022, pp. 17–18)
Goal	Verification Testing Validation Testing	Approach	(ISO/IEC and IEEE, 2022, p. 16)
Source of Information for Designing Tests	Specification-based Testing Structure-based Testing Experience-based Testing	Technique ^e	(ISO/IEC and IEEE, 2021, pp. 4, 8)
Selection Process	Deterministic Testing Random Testing	Technique	(Washizaki, 2024, pp. 5-12, 5-16)
Stage of Lifecycle	Waterfall Testing Incremental Testing Continuous Testing (CT)	Practice (inferred)	(Firesmith, 2015, p. 29)

^a Defined in Section 2.1.3.

^b Defined in Section 2.1.1.

^c These approaches may instead be techniques (implied by ISO/IEC and IEEE, 2022, p. 35; see Table 6.4).

^d We also consider this categorization meaningful (see Appendix A.3).

^e Experience-based testing may instead be a practice (ISO/IEC and IEEE, 2022, pp. 22, 34; 2021, p. viii; see Table 6.4).

Another way to subdivide the IEEE categories is by grouping related test approaches into a “class” or “family” with “commonalities and well-identified variabilities that can be instantiated”, where “the commonalities are large and the variabilities smaller” (Carette, 2024). Examples of these are the classes of combinatorial (ISO/IEC and IEEE, 2021, p. 15) and data flow testing (p. 3) and the family of performance-related testing (Moghadam, 2019, p. 1187)²; we explore the notion of test approach families in more detail in Section 6.2.1. Note that “there is a lot of overlap between different classes of testing” (Firesmith, 2015, p. 8), meaning that “one category [of test techniques] might deal with combining two or more techniques” (Washizaki, 2024, p. 5-10). For example, “performance, load and stress testing might considerably overlap in many areas” (Moghadam, 2019, p. 1187). A side effect of this is that it is difficult to “untangle” these classes; for example, take the following sentence: “whitebox fuzzing extends dynamic test generation based on symbolic execution and constraint solving from unit testing to whole-application security testing” (Godefroid and Luchaup, 2011, p. 23)! This is, in part, why research on software testing terminology is so vital.

Q #7: Is this sentence needed?

²The original source describes “performance testing ... as a family of performance-related testing techniques”, but it makes more sense to consider “performance-related testing” as the “family” with “performance testing” being one of the variabilities (see Section 7.4).

Chapter 5

Tools

To better understand our findings, we build tools to more intuitively visualize relations between test approaches (Section 5.1) and automatically track their flaws (Section 5.2). Doing this manually would be error-prone due to the amount of data involved (for example, we identify 558 test approaches) and the number of situations where the underlying data would change, including more detailed analysis, error corrections, and the addition of data. These all require tedious updates to the corresponding graphs that may be overlooked or done incorrectly, so automating these processes allows for our results to be reproduced (by us or others) and account for new data. Besides being more systematic, this also allows us to observe the impacts of smaller changes, such as unexpected flaws that arise from a new relation between two approaches. It also helps verify the tools themselves; for example, tracking a flaw manually should affect relevant flaw counts, which we can double-check. We also define \LaTeX macros to help achieve our goals of maintainability, traceability, and reproducibility (Section 5.3).

5.1 Approach Relation Visualization

To better visualize how test approaches relate to each other, we develop a tool to automatically generate visualizations of these relations. Since we use a consistent format to track synonym and parent-child relations (see Sections 2.1.2 and 2.1.3, respectively) between approaches in our test approach glossary, we can parse them systematically. For example, if the entries in Table 5.1 appear in the glossary, then their parent-child relations are displayed as Figure 5.1a in the generated graph. We also capture relevant citation information in our glossary in the author-year citation format, “reusing” information from previous citations when applicable. For example, the first row of Table 5.1 contains the citation “(Author, 0000; 0001)”, which means that this information was present in two documents by Author: one written in the year 0000, and one in 0001. The following citation, “(0000)”, contains no author, which means it was written by the same one as the previous citation (Author). We process these citations according to this logic (see the relevant source code) so we can consistently track them throughout our analysis. We also colour each relation according to its source tier (see Section 3.1), including

Table 5.1: Example glossary entries demonstrating how we track parent-child relations (see Section 2.1.3).

Name ^a	Parent(s)
A	B (Author, 0000; 0001), C (0000)
B	C (implied by Author, 0000)
C	D (Author, 0002)
D (implied by Author, 0002)	

^a “Name” can refer to the name of a test approach, software quality, or other testing-related term, but we only generate graphs for test approaches.




Figure 5.1: Example generated visualizations of parent-child relations.

inferences (see Section 2.3) and proposals (see Chapter 7), but only display the relation with the most credible source tier (see Section 2.4). Although we omit this colouring from Figures 5.1 and 5.4 for brevity, Figure 5.2 shows how this affects our visualizations; each line representing a synonym relation is coloured based on the most credible source that describes that relation.

We graph all parent-child relations, since they are guaranteed to be visually meaningful. However, since each term is trivially a synonym of itself and there are many non-problematic synonyms that do not imply flaws (see Section 2.1.2), we only visualize the following synonym relations, which may indicate flaws:

1. Synonyms between approaches defined independently.

If two separate approaches have their own definitions, nuances, etc. but are also labelled as synonyms, this may indicate that the two terms are interchangeable and could be merged *or* that either their definitions or this synonym relation is incorrect.

2. Synonyms that violate transitivity.

If two distinct approaches share a synonym, that implies that they are synonyms themselves. If they are *not*, one or more relations may be incorrect or missing.

We deduce these conditions from the information we parse from our glossary. Applying this technique to our test approach glossary (and manually modifying it for legibility) results in Figure 5.2 (note that we exclude implicit synonyms from this visualization, but we include them in others). These relations are given as described by the literature and are therefore flawed. In particular, we later discuss synonyms that violate transitivity (as described earlier) in Section 6.2.2 and other kinds of flawed synonym relations in Sections 6.2.3, 6.3, 6.6, and 6.7.

Since we also track the “explicitness” of information (see Section 2.3), we can represent explicit *and* implicit relations without double counting them during the analysis in Section 5.2. If a relation is both explicit *and* implicit, we only include the implicit relation in the graph if its source is more “credible” (see Section 2.4).

Implicit approaches and relations are denoted by dashed lines, as shown in Figure 5.1a; explicit approaches are *always* denoted by solid lines, even if they are also implicit. We can also generate “explicit” versions of these graphs that exclude implicit approaches and relations; for example, Figure 5.1b is the explicit version of Figure 5.1a, and Figure 5.2 likewise only contains explicit approaches and synonym relations.

Since these graphs tend to be large, it is useful to focus on specific subsets of them. For each approach category (see Section 2.1.1), we generate a graph restricted to its approaches and the relations between them. We also generate a graph of all static approaches along with the relations between them *and* between a static approach and a dynamic approach. We generate this static-focused graph because static testing is sometimes considered to be a separate approach category (see Definition Flaw 4), but we include the relations with dynamic approaches since they are our primary focus (see Appendix A.3). We colour all dynamic approach nodes in these graphs grey to distinguish them.

We can also generate more specific subsets of these graphs from a given subset of approaches, such as those pertaining to recovery or scalability. We generate graphs of just these subsets of software testing to help visualize the relations between these approaches, but we can also update them based on our recommendations by specifying sets of approaches and relations to add or remove. When doing

Q #8: Do I need an example of this?

Q #9: Since this doesn't come up in our thesis, do I even need an example?

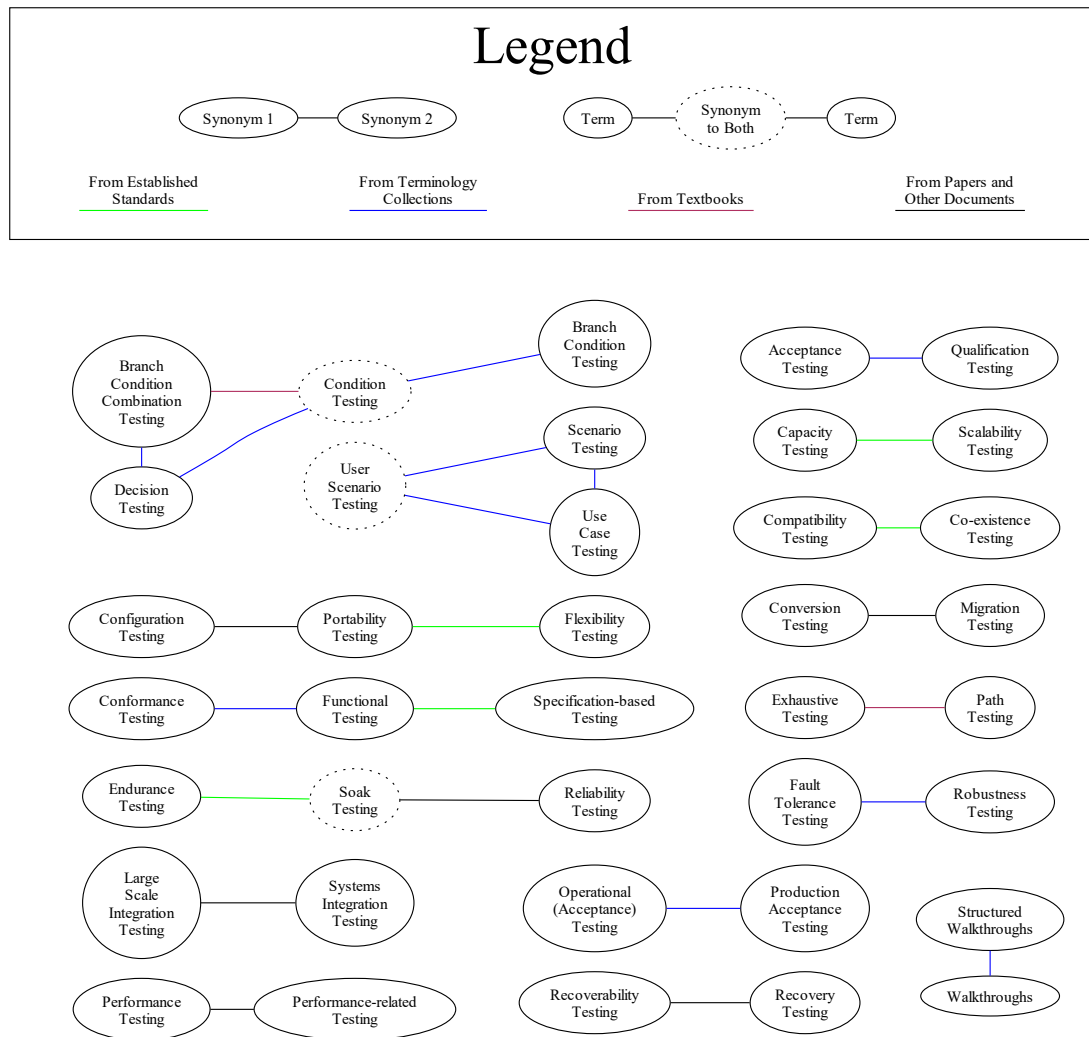


Figure 5.2: Synonym relations given explicitly by the literature.

so, we colour any added approaches or relations orange to distinguish them. See Chapter 7 for examples of this process’s output.

5.2 Flaw Analysis

In addition to analyzing specific flaws, it is also useful to examine them at a higher level. We automate subsets of this task where applicable (Section 5.2.1) and augment the remaining manual portion with automated tools (Section 5.2.2). This gives us an overview of:

- how many flaws (defined in Section 2.2) there are,
- how these flaws present themselves (see Section 2.2.1),
- in which knowledge domains these flaws occur (see Section 2.2.2),

- how explicit (see Section 2.3) these flaws are, and
- how responsible each source tier (see Section 3.1) is for these flaws.

To understand where flaws exist in the literature, we group them based on the source tier(s) responsible for them. We then count each flaw *once* per source tier if it appears within it *and/or* between it and a more credible tier¹ (see Sections 2.4 and 3.1). This avoids counting the same flaw more than once for a given source tier, which would give the number of *occurrences* of all flaws instead of the more useful number of flaws *themselves*. When taking a more detailed look at the *sources* of flaws (as opposed to just the responsible source *tiers*), as we do in Figure 6.1, we also count the following sources of flaws separately:

1. self-contained flaws (defined in Section 2.2.3),
2. internal flaws (defined in Section 2.2.3),
3. those between documents with the same set of authors, which includes
 - (a) the various combinations of ISO, the International Electrotechnical Commission (IEC), and IEEE shown in Figure 5.3 and
 - (b) the different versions of the Guides to the SoftWare Engineering Body Of Knowledge (SWEBOK Guides), which have different editors (Washizaki, 2024; Bourque and Fairley, 2014) but are written by the same organization: the IEEE Computer Society (Washizaki, 2025; see Section 3.1.2), and
4. those within a single source tier.

As before, we do not double count these sources of flaws, meaning that the maximum number of counted flaws possible within a *single* source tier in this more detailed view is four (one for each type). This only occurs if there is an example of each flaw source that is *not* ignored to avoid double counting; for example, while a single flaw within a single document would technically and trivially fulfill all four criteria, we would only count it once.

5.2.1 Automated Flaw Analysis

As outlined in Section 5.1, we automatically detect synonym relations from our test approach glossary that violate transitivity to generate our graphs. These relations are significant because they indicate potential flaws. We automatically detect and format these flaws to present them when discussing synonym relation flaws in Section 6.2.2. The next logical step is then to detect other classes of flaws.

Parent-child relations that violate irreflexivity as outlined in Section 2.1.3 (i.e., cases where a child is given as a parent of itself) are also trivial to automate by

¹If an inconsistency occurs between two source tiers and the more credible one is *incorrect*, we instead count it as an inconsistency between it and the asserted truth from the less credible source, as described in Section 5.2.2.

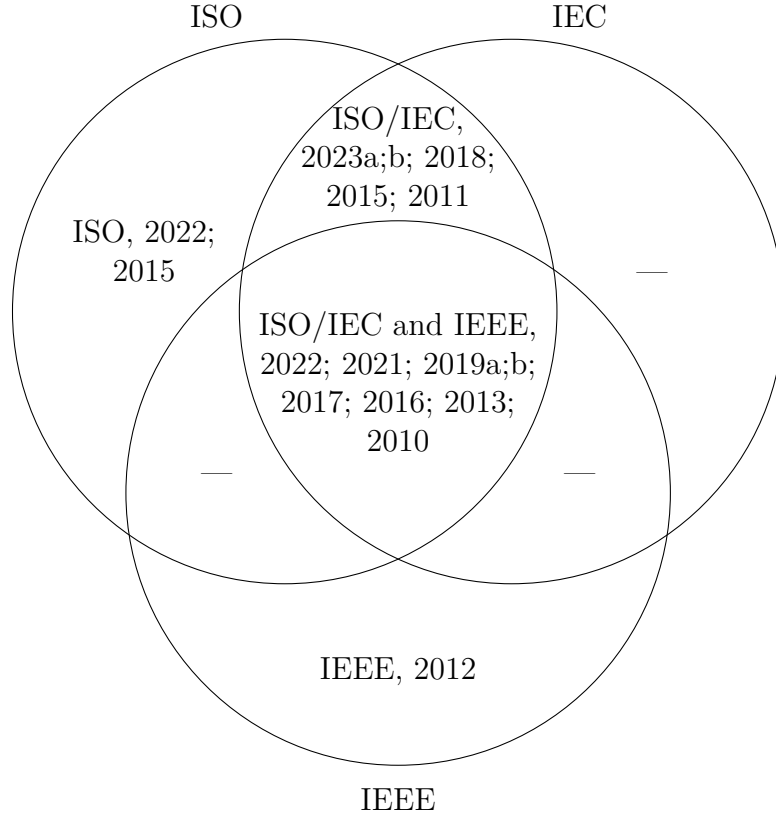


Figure 5.3: The sets of authors of established standards (as defined in Section 3.1.1).

looking for lines in the generated \LaTeX files that start with $I \rightarrow I$, where I is the label used for a test approach node in these graphs. This process results in output similar to Figure 5.4a. We use a similar process to detect pairs of approaches with a synonym relation *and* a parent-child relation; these are flaws since synonym relations are symmetric while parent-child relations are asymmetric as outlined in Sections 2.1.2 and 2.1.3, respectively. To find these flaws, we build a dictionary of each term’s synonyms to evaluate which synonym relations are notable enough to include in the graph, and then check these mappings to see if one appears as a parent of the other. For example, if J and K are synonyms, a generated \LaTeX file with a parent line starting with $J \rightarrow K$ or $K \rightarrow J$ would result in these approaches being graphed as shown in Figure 5.4b. We present these two classes of flaws when discussing parent-child relation flaws in Section 6.2.3.

While just counting the total number of flaws (found automatically *or* manually) is trivial, tracking the source(s) of these flaws is more useful, albeit more involved. Since we consistently track the appropriate citations for each piece of information we record (see Table 5.1 for an example of how we format these citations in our glossaries), we can use them to identify the offending source tier(s). This comes with the added benefit that we can format these citations to use with \LaTeX ’s citation commands in this thesis, including generating the comments used to analyze flaws as described in Section 5.2.2.




Figure 5.4: Example generated graphs containing flaws.

Alongside this citation information, we include keywords so we can assess how “explicit” a piece of information is (see Section 2.3). This is useful when counting flaws, since they can be both objective and subjective but should not be double counted as both! When presenting the numbers of flaws sorted by various criteria in Tables 6.1 and 6.2, each flaw is counted only for its most “explicit” manifestation (i.e., it will only increment a value in the “(Sub)jective” column if it is *not* also “(Obj)ective”), similarly to how we generate graphs (see Section 5.1).

5.2.2 Augmented Flaw Analysis

While we can detect some subsets of flaws automatically by analyzing our test approach glossary, most are too complex and need to be tracked manually. We record these more detailed flaws as \LaTeX enumeration items along with comments that we can parse automatically, allowing us to analyze them more broadly. We also add these comments to flaws we detect automatically before generating the corresponding \LaTeX file to ensure these flaws also get analyzed. These comments have the following format:

```
% Flaw count (MNFST, DMN): {A1} {A2} ... | {B1} ... | {C1} ...
```

MNFST and DMN are placeholders for the “keys” given in Tables 2.2 and 2.3, respectively, that we use to track a flaw’s manifestation(s) and domain(s) (defined in Section 2.2). For example, the comment line for an incorrect synonym relation would start with % Flaw count (WRONG, SYNS) and one for a redundant label would start with % Flaw count (REDUN, LABELS). We omit these keys from constructed examples of these comments without associated flaws throughout this chapter for brevity. Finally, A1, A2, B1, and C1 are each placeholders for a source involved in this example flaw; in general, there can be arbitrarily many. We represent each source by its \LaTeX key, and wrap each one in curly braces (with the exception of the ISTQB glossary due to its use of custom commands via $\text{\citealias{}}$) to mimic \LaTeX ’s citation commands for ease of parsing. We then separate each

“group” of sources with a pipe symbol (|) so we can compare each pair of groups; in general, a flaw can have any number of groups of sources.

As mentioned in Section 2.2.3, we make a distinction between “self-contained” flaws and “internal” flaws. If only one group of sources is present in a flaw’s comment, such as the first line below, we consider it to be a self-contained flaw. If a source appears in multiple groups in a flaw’s comment, we consider it to be an internal flaw. The second line is a standard example of this, while the third is more complex; in this case, source Y agrees with only one of the conflicting sources of information in X.

```
% Flaw count: {X}
% Flaw count: {X} | {X}
% Flaw count: {X} | {X} {Y}
```

We do not double count flaws that reappear when comparing between pairs of groups; this means the following line adds an inconsistency between X and Z *and* between Y and Z *without* double counting the former.

```
% Flaw count: {X} | {X} {Y} | {Z}
```

We compare the authors and years of each source involved with a given flaw to determine if it manifests within a single document and/or between documents with the same set of authors. Then, we group these sources into their tiers (see the relevant source code). We then distill these lists of sources down to sets of tiers and compare them against each other to determine how many times a given flaw manifests between source tiers. This determines which row(s) of Tables 6.1 and 6.2 and which bar(s) of Figure 6.1 that the given flaw should count toward. We describe this process in more detail in Section 5.2.2.

To give a more complete example, we track Synonym Flaw 2 with the following comment line:

```
% Flaw count (OVER, SYNS): {IEEE2017} {IEEE2013} | {IEEE2022}
↪ {IEEE2017} {Perry2006}
```

Since IEEE2022, IEEE2017, and IEEE2013 are all written by the same standards organizations (ISO/IEC and IEEE), we count this as an inconsistency between documents with the same set of authors in Figure 6.1, but only once to avoid double counting.

We can also specify the “explicitness” (see Section 2.3) of a flaw by inserting the phrase “implied by” after the sources of explicit information and before those of implicit information. This information is parsed following the same rules described in Section 5.2.1 for automatically detected flaws. Note that we only count subjective flaws if there is not an equivalent objective flaw, as we do when generating graphs (Section 5.1).

```
% Flaw count (CONTRA, DEFS): {IEEE2021} {IEEE2017} |
↪ {vanVliet2000} implied by {IEEE2021}
```

See #137 and
#138

Later: Ensure
this is up to date

Later: Ensure
this is up to date

For example, the above comment line from Definition Flaw 17 indicates that the flaws given below are present. The third flaw only affects Figure 6.1 due to its more nuanced breakdown of the sources of flaws. The rest increment their corresponding count in Figure 6.1 and Tables 6.1 and 6.2 by only one:

- an objective inconsistency between a textbook and a standard,
- a subjective flaw within a single document, and
- a subjective inconsistency between documents with the same set of authors (ISO/IEC and IEEE).

See #184

Occasionally, we assert that a source from a less credible tier is more correct than a source from a more credible tier. For example, the terms “acceleration tolerance testing” and “acoustic tolerance testing” seem to only refer to software testing in (Firesmith, 2015, p. 56); elsewhere, they seem to refer to testing the acoustic tolerance of rats (Holley et al., 1996) or the acceleration tolerance of astronauts (Morgun et al., 1999, p. 11), aviators (Howe and Johnson, 1995, pp. 27, 42), or catalysts (Liu et al., 2023, p. 1463), which don’t exactly seem relevant... This flaw is supported by additional papers found via a miniature literature review (described in Section 3.5) from a lower source tier than (Firesmith, 2015) (which is a terminology collection; see Section 3.1). However, this flaw is really based in (Firesmith, 2015) and not in these additional papers, but this would be counted as a flaw in these papers if they were included as detailed above. Therefore, we document these “assertion” sources separately to track them for traceability without incorrectly counting flaws, such as the following for this specific example (Scope Flaw 2):

```
% Flaw count (WRONG, LABELS): {Firesmith2015}
% Assertion: {LiuEtAl2023} {MorgunEtAl1999} {HolleyEtAl1996}
↪ {HoweAndJohnson1995}
```

5.3 L^AT_EX Commands

To improve maintainability, traceability, and reproducibility, we define helper commands (also called “macros”) for content that is prone to change or used in multiple places. For example, we use Python scripts to calculate values based on our glossaries and save them to files to be assigned to corresponding L^AT_EX macros. We use these throughout our documents instead of manually updating these constantly changing values, which is prone to error. Table 5.2 lists these macros, along with their current values and descriptions of what they represent. Our Python scripts convert numbers to their textual equivalents when necessary to follow IEEE guidelines.

Table 5.2: \LaTeX macros for calculated values.

Macro	What it Counts	Value
$\backslash\text{approachCount}\{\}$	Identified test approaches	558
$\backslash\text{qualityCount}\{\}$	Identified software qualities	75
$\backslash\text{srcCount}\{\}^a$	Sources used in glossaries	66
$\backslash\text{flawCount}\{\}^b$	Identified flaws	275
$\backslash\text{TotalBefore}\{\}^c$	Test approaches identified before process in Section 3.5	471
$\backslash\text{UndefBefore}\{\}^c$	Undefined test approaches identified before process in Section 3.5	174
$\backslash\text{TotalAfter}\{\}^c$	Test approaches identified after process in Section 3.5	558
$\backslash\text{UndefAfter}\{\}^c$	Undefined test approaches identified after process in Section 3.5	190
$\backslash\text{multiSynCount}\{\}$	Terms given as synonyms for multiple discrete terms	nine
$\backslash\text{parSynCount}\{\}$	Pairs of test approaches with a child-parent <i>and</i> synonym relation	17
$\backslash\text{selfParCount}\{\}$	Test approaches that are a parent of themselves	three

^a Calculated in \LaTeX from source tier lists; see Section 5.3.

^b Alias for $\backslash\text{totalFlawDmnBrkdwn}\{13\}$; see Section 5.3.

^c These macros are defined as counters to allow them to be used in calculations within \LaTeX (such as in Section 3.5 and Fig. 3.5).

Additionally, we count flaws based on their manifestation and domain, explicitness, and source tier (defined in Sections 2.2, 2.3, and 3.1, respectively). For each source tier, we create two files that each include both levels of explicitness: one for manifestations and one for domains. For example, flaws in standards are saved to `build/stdFlawMnfstBrkdwn.tex` by manifestation. We then assign these data to macros (such as $\backslash\text{stdFlawMnfstBrkdwn}\{\}$) to populate Tables 6.1 and 6.2. For example, we access the number of objective and subjective mistakes in standards by using $\backslash\text{stdFlawMnfstBrkdwn}\{1\}$ and $\backslash\text{stdFlawMnfstBrkdwn}\{2\}$, respectively. We follow a similar process for tracking the total numbers of flaws; this includes $\backslash\text{totalFlawMnfstBrkdwn}\{13\}$ and $\backslash\text{totalFlawDmnBrkdwn}\{13\}$ which are identical and track the total number of identified flaws.

Just as with calculated values, it is important that repeated text is updated consistently, which we accomplish by defining more macros. Some of these are generated by Python scripts in a similar fashion to calculated values, such as those for flaw manifestations and domains in Table 5.3 and the lists of sources in each source tier. The latter are built by extracting all sources cited in our

three glossaries, categorizing, sorting, and formatting them (including handling edge cases), and saving them to a file. These are then assigned to `\stdSources{}`, `\metaSources{}`, `\textSources{}`, and `\paperSources{}` and include:

1. the source tier’s name,
2. the list of sources in the tier, and
3. the number of sources in the tier.

These are accessed by passing in the corresponding number in the above enumeration (e.g., `\paperSources{2}`). We use the first value for the subheadings in Section 3.1, the first two for Appendix B and the third to build Figure 3.1 and calculate `\srcCount{}` (see Table 5.2).

However, we create most of these macros for reused text manually when we first notice the reuse, including the source tier macros in Table 5.3. Some of these macros account for context-specific formatting depending on how they are used, such as capitalization. These tend to be less well-defined, since they arise naturally from the writing process, so we omit these details from the manually created text macros in Table 5.4, which are grouped based on the type of information they contain.

Table 5.3: L^AT_EX macros for referencing well-defined sections.

	Flaw Manifestations ^a	Flaw Domains ^b	Source Tiers ^c
Macros (Values)	<code>\wrong{}</code> (Mistakes)	<code>\cats{}</code> (Categories)	<code>\stds{}</code> (Established Standards)
	<code>\miss{}</code> (Omissions)	<code>\syns{}</code> (Synonyms)	<code>\metas{}</code> (Terminology Collections)
	<code>\contra{}</code> (Contradictions)	<code>\pars{}</code> (Parents)	<code>\texts{}</code> (Textbooks)
	<code>\ambi{}</code> (Ambiguities)	<code>\defs{}</code> (Definitions)	<code>\papers{}</code> (Papers and Other Documents)
	<code>\over{}</code> ^d (Overlaps)	<code>\labels{}</code> (Labels)	<code>\papers*{}</code> (Papers and Others)
	<code>\redun{}</code> ^e (Redundancies)	<code>\scope{}</code> (Scope)	
		<code>\trace{}</code> (Traceability)	
Used In	Tables 2.2 and 6.1 (in both thesis and paper)	Tables 2.3 and 6.2 (in both thesis and paper)	Figures 3.1 and 6.1 and Tables 6.1 and 6.2

^a Defined in Section 2.2.1; we also define starred versions, such as `\wrong*{}` (Mistake), that use the singular noun.

^b Defined in Section 2.2.2; we also define starred versions, such as `\cats*{}` (Category), that use the singular noun.

^c Defined in Section 3.1.

^d We overwrite the primitive T_EX command `\over{}` since we do not otherwise use it.

^e This takes an optional parameter for a `note-tag` that corresponds to `\redunNote{}` (“Section omitted for brevity.”; see Table 5.4) when used in `tabulararrays` (i.e., `\redun[c]{}`) since we omit this section from our paper.

Table 5.4: L^AT_EX macros for reused text.

Type	Macro	Used in
Flaws ^a	<code>\bugPattonFlaw{}</code>	Chapter 1 and Synonym Flaw 6
	<code>\alphaFlaw{}</code>	Chapter 1 and Definition Flaw 10
	<code>\loadFlaw{}</code>	Chapter 1 and Definition Flaw 12
	<code>\expBasedCatMain{}</code>	Chapter 1 and Section 6.2.1
	<code>\tourFlaw{}</code>	Chapters 1 and 6 and Definition Flaw 9
	<code>\redBoxFlaw{}</code>	Sections 2.3 and 6.1.1 and Mistake 23
	<code>\perfAsFamily{}</code>	Sections 4.1.2 and 6.2.1
	<code>\tolTestFlaw{}</code>	Sections 5.2.2 and 6.1.1 and Scope Flaw 2
	<code>\accelTolTest{}</code>	<code>\tolTestFlaw{}</code> and Appendix A.1
	<code>\errorGuessFlaw{}</code>	Section 6.1.1 and Mistake 1
	<code>\parSheetTestFlaw{}</code>	Section 6.1.1 and Mistake 14
	<code>\perfSecParFlaw{}</code>	Parent Flaw 2 and paper version of Section 6.2.3
Footnotes	<code>\ftrnote{}</code>	Thesis (automated) and paper (manual) versions of Table 6.5
	<code>\specfn{}</code>	
	<code>\ucstn{}</code>	
	<code>\redunNote{}</code>	<code>\redun{}</code> and Table 2.2 ^b
	<code>\notDefDistinctIEEE{}</code>	Definition Flaw 4 and Section 9.1
	<code>\gerrardDistinctIEEE{}</code>	Table 4.2 and Definition Flaw 38
Links	<code>\distinctIEEE{}</code>	<code>\gerrardDistinctIEEE{}</code> , <code>\notDefDistinctIEEE{}</code> , and Sections 4.1.2 and 6.2.1
	<code>\ourApproachGlossary{}</code>	Table 1.1 and Sections 2.3, 3.3, 3.4, 5.1, 5.2.1, 5.2.2, and 7.1
	<code>\seeSrcCode{}</code>	Sections 5.1 and 5.2.1
RQs	<code>\recFigs{}</code>	Chapters 6 and 7
	<code>\rqatext{}</code>	Chapter 1 and seminar slides
	<code>\rqbtext{}</code>	
	<code>\rqctext{}</code>	

Continued on next page

Table 5.4: \LaTeX macros for reused text. (Continued)

Type	Macro	Used in
Misc.	<code>\supersAck{}</code>	Acknowledgements and seminar slides
	<code>\supers{}</code>	Declaration of Academic Achievement and <code>\supersAck{}</code>
	<code>\defRel{}</code>	Sections 2.1.2 and 2.1.3
	<code>\defLabelDistinct{}</code>	Sections 2.2.2 and 6.2.5
	<code>\oneSrcDistinct{}</code>	Sections 2.2.3 and 5.2.2
	<code>\approachFields{}</code>	Sections 2.3 and 3.3
	<code>\impKeywords{}</code>	Sections 2.3 and 3.4
	<code>\orthTestIntro{}</code>	Section 2.3 and Appendix A.6
	<code>\addTextEx{}</code>	Sections 3.1 and 3.1.3
	<code>\displayNL{}</code>	Sections 5.1 and 5.2.2

^a Defined in Section 2.2.^b See Table 5.3 for more details on how we use `\redunNote{}` alongside `\redun{}`.

Chapter 6

Observed Flaws

After gathering all these data¹, we find many flaws. Figure 6.1 shows the source tiers (see Section 3.1) responsible for these flaws, which reveals a lot about software testing literature:

1. Established standards (Section 3.1.1) aren't actually standardized, since:
 - (a) other documents disagree with them *very* frequently and
 - (b) they are the most internally inconsistent source tier!
2. Less standardized (or “credible”; see Section 2.4) documents, such as terminology collections and textbooks (Sections 3.1.2 and 3.1.3, respectively) are also not followed to the extent they should be.
3. Documents across the board have flaws within the same document, between documents with the same author(s), or even with reality!

Q #10: Is this too strong of a synonym for “ground truth”?

To better understand and analyze these flaws, we group them by their manifestations and their domains as defined in Section 2.2. We present the total number of flaws by manifestation and by domain in Tables 6.1 and 6.2, respectively, where a given row corresponds to the number of flaws either within that source tier and/or with a “more credible” one (i.e., a previous row in the table; see Sections 2.4 and 3.1). We also group these flaws by their explicitness (defined in Section 2.3) by counting (Obj)ective and (Sub)jective flaws separately (see Tables 6.1 and 6.2), since additional context may rectify them. Additionally, Tables 6.4 to 6.6 contain relations that are explicit, implicit, and both; implicit relations are marked by the phrases “implied by”.

Since we give each flaw a manifestation *and* a domain, the totals per source and grand totals in these tables are equal. As an example, the structure of tours can be defined as either quite general (ISO/IEC and IEEE, 2022, p. 34) or “organized around a special focus” (Hamburg and Mogyorodi, 2024). This is a contradictory definition, so it appears as both Contradiction 7 and Definition Flaw 9. Within these sections, we list “more significant” flaws first, followed by “less significant”

¹Available in `ApproachGlossary.csv`, `QualityGlossary.csv`, and `SuppGlossary.csv` at <https://github.com/samm82/TestingTesting>.




Figure 6.1: Identified flaws by the source tier responsible. Some bars are omitted as they correspond to comparisons we do not make; see Section 5.2.

ones we omit from the paper version of this thesis. We then sort these flaws by their source tier (see Section 3.1).

We only list flaws we automatically uncover based on their domain (see Section 5.2.1) in their corresponding domain section for clarity, although they still contribute to the counts in Table 6.1. Moreover, certain “subsets” of testing contain many interconnected flaws, which we present in subsections as a “third view”. This keeps related information together but causes a further apparent mismatch between the counts in Tables 6.1 and 6.2 and the number of flaws in Sections 6.1 and 6.2; we still include these flaws in the manifestation and domain counts (Section 5.2.2 outlines how this is done). These subsets include functional testing (Section 6.3), operational (acceptance) testing (Section 6.4), recovery testing (Section 6.5), scalability testing (Section 6.6), and compatibility testing (Section 6.7). Finally, we present inferred flaws, as described in Section 2.3, for completeness in

Section 6.8; these do not contribute to any counts. We colour inferred approaches and relations grey in our generated graphs, such as Figures 7.1 to 7.3.

Table 6.1: Breakdown of identified Flaws by Manifestation by Source Tier.

Source Tier	Mistakes		Omissions		Contradictions		Ambiguities		Overlaps		Redundancies		Total
	Obj	Sub	Obj	Sub	Obj	Sub	Obj	Sub	Obj	Sub	Obj	Sub	
Established Standards	6	1	2	0	18	10	5	0	9	0	0	0	51
Terminology Collections	13	0	1	0	33	17	15	3	5	1	2	0	90
Textbooks	6	2	2	0	36	5	4	0	1	0	0	0	56
Papers and Others	11	1	5	0	24	22	9	3	0	1	2	0	78
Total	36	4	10	0	111	54	33	6	15	2	4	0	275

Table 6.2: Breakdown of identified Flaws by Domain by Source Tier.

Source Tier	Categories		Synonyms		Parents		Definitions		Labels		Scope		Traceability		Total
	Obj	Sub	Obj	Sub	Obj	Sub	Obj	Sub	Obj	Sub	Obj	Sub	Obj	Sub	
Established Standards	10	7	2	2	7	0	15	2	6	0	0	0	0	0	51
Terminology Collections	7	14	10	3	10	2	20	0	11	2	5	0	6	0	90
Textbooks	0	0	15	0	9	6	18	1	7	0	0	0	0	0	56
Papers and Others	9	13	14	9	11	3	9	0	7	2	0	0	1	0	78
Total	26	34	41	14	37	11	62	3	31	4	5	0	7	0	275

Table 6.3: Different kinds of mistakes found in the literature.

Description	Count
Information is incorrect based on an assertion from another source	9
Information is not present where it is claimed to be	6
Information is provided with an incorrect scope	5
Information contains a minor mistake	3 ^a
Incorrect information makes other information incorrect	1

^a Comprises two typos and one duplication.

6.1 Flaws by Manifestation

The following sections list observed flaws grouped by *how* they manifest as presented in Section 2.2.1. These include mistakes (Section 6.1.1), omissions (Section 6.1.2), contradictions (Section 6.1.3), ambiguities (Section 6.1.4), overlaps (Section 6.1.5), and redundancies (Section 6.1.6).

6.1.1 Mistakes

There are many ways that information can be incorrect which we identify in Table 6.3. We provide an example below for those that are unintuitive; see Appendix C.1.1 for the full list of mistakes.

Information is incorrect based on an assertion from another source

Since errors are distinct from defects/faults (ISO/IEC and IEEE, 2010, pp. 128, 140; Washizaki, 2024, p. 12-3; van Vliet, 2000, pp. 399–400), error guessing should instead be called “defect guessing” if it is based on a “checklist of potential defects” (ISO/IEC and IEEE, 2021, p. 29) or “fault guessing” if it is a “fault-based technique” (Bourque and Fairley, 2014, p. 4-9) that “anticipate[s] the most plausible faults in each SUT” (Washizaki, 2024, p. 5-13). One (or both) of these proposed terms may be useful in tandem with “error guessing”, which would focus on errors as traditionally defined; this would be a subapproach of error-based testing (implied by van Vliet, 2000, p. 399).

Information is provided with an incorrect scope

“Par sheet testing” from (Hamburg and Mogyorodi, 2024) seems to refer to the specific example mentioned in Mistake 12 and does not seem more widely applicable, since a “PAR sheet” is “a list of all the symbols on each reel of a slot machine” (Bluejay, 2024). Conversely, the terms “acceleration tolerance testing” and “acoustic tolerance testing” seem to only refer to software testing in (Fire-smith, 2015, p. 56); elsewhere, they seem to refer to testing the acoustic tolerance of rats (Holley et al., 1996) or the acceleration tolerance of astronauts (Morgun

et al., 1999, p. 11), aviators (Howe and Johnson, 1995, pp. 27, 42), or catalysts (Liu et al., 2023, p. 1463), which don’t exactly seem relevant...

Incorrect information makes other information incorrect

The incorrect claim that “white-box testing”, “grey-box testing”, and “black-box testing” are synonyms for “module testing”, “integration testing”, and “system testing”, respectively, (see Mistake 22) casts doubt on the claim that “red-box testing” is a synonym for “acceptance testing” (Sneed and Göschl, 2000, p. 18) (see Mistake 23).

OG Hetzel88

6.1.2 Omissions

The following are cases where information (usually definitions; see Section 6.2.4) *should be* included but is not:

1. Integration testing, system testing, and system integration testing are all listed as “common test levels” (ISO/IEC and IEEE, 2022, p. 12; 2021, p. 6), but the latter two are not defined. This makes the relations between these three terms unclear; system integration testing is listed as a child of both integration testing (Hamburg and Mogyorodi, 2024) and system testing (Firesmith, 2015, p. 23).
2. Similarly, component testing, integration testing, and component integration testing are all mentioned by ISO/IEC and IEEE (2017), but the latter is only defined as “testing of groups of related components” (p. 82). As before, the relations between these three terms are unclear; component integration testing is only listed as a child of integration testing (Hamburg and Mogyorodi, 2024).
3. The acronym for System of Systems (SoS) (ISO/IEC and IEEE, 2019b) is used but not defined by Firesmith (2015, p. 23).
4. Van Vliet (2000, p. 425) defines many types of data flow coverage, including all-p-uses, all-p-uses/some-c-uses, and all-c-uses/some-p-uses, but excludes all-c-uses, which is implied by these definitions and defined elsewhere (ISO/IEC and IEEE, 2021, p. 27; 2017, p. 83; Peters and Pedrycz, 2000, p. 479).
5. Bas (2024, p. 16) lists “three [backup] location categories: local, offsite and cloud based [sic]” but does not define or discuss “offsite backups” (pp. 16-17).
6. Gerrard (2000a, Tab. 2) makes a distinction between “transaction verification” and “transaction testing” and uses the phrase “transaction flows” (Fig. 5) but doesn’t explain them.

Q #11: I’m counting this as an ambiguous parent relation occurring in ISTQB; is this sufficient?

7. Availability testing isn't assigned to a test priority (Gerrard, 2000a, Tab. 2), despite the claim that “the test types² have been allocated a slot against the four test priorities” (p. 13); usability testing and/or performance testing would have been good candidates.

6.1.3 Contradictions

The following are cases where multiple sources of information (sometimes within the same document!) disagree:

1. Regression testing and retesting are sometimes given as two distinct approaches (ISO/IEC and IEEE, 2022, p. 8; Firesmith, 2015, p. 34), but sometimes regression testing is defined as a form of “selective retesting” (ISO/IEC and IEEE, 2017, p. 372; Washizaki, 2024, pp. 5-8, 6-5, 7-5 to 7-6; Barbosa et al., 2006, p. 3). Moreover, the two possible variations of regression testing given by van Vliet (2000, p. 411) are “retest-all” and “selective retest”, which is possibly the source of the above misconception. This creates a cyclic relation between regression testing and selective retesting.
2. “Software testing” is often defined to exclude static testing (Firesmith, 2015, p. 13; Ammann and Offutt, 2017, p. 222; Peters and Pedrycz, 2000, p. 439), restricting “testing” to mean dynamic validation (Washizaki, 2024, p. 5-1) or verification “in which a system or component is executed” (ISO/IEC and IEEE, 2017, p. 427). However, “terminology is not uniform among different communities, and some use the term ‘testing’ to refer to static techniques³ as well” (Washizaki, 2024, p. 5-2). This is done by ISO/IEC and IEEE (2022, pp. 16–17) and Gerrard (2000a, pp. 8–9); the former even explicitly *exclude* static testing in another document (2017, p. 440)!
3. A component is an “entity with discrete structure ... within a system considered at a particular level of analysis” (ISO/IEC, 2023b) and “the terms module, component, and unit [sic] are often used interchangeably or defined to be subelements of one another in different ways depending upon the context” with no standardized relationship (ISO/IEC and IEEE, 2017, p. 82). For example, Hamburg and Mogyorodi (2024) define them as synonyms while Baresi and Pezzè (2006, p. 107) say “components differ from classical modules for being re-used in different contexts independently of their development”. Additionally, since components are structurally, functionally, or logically discrete (ISO/IEC and IEEE, 2017, p. 419) and “can be tested in isolation” (Hamburg and Mogyorodi, 2024), “unit/component/module testing” could refer to the testing of both a module *and* a specific function in a module, introducing a further level of ambiguity.

Are these separate approaches?

See #14

²“Each type of test addresses a different risk area” (Gerrard, 2000a, p. 12), which is distinct from the notion of “test type” described in Table 2.1.

³Not formally defined, but distinct from the notion of “test technique” described in Table 2.1.

4. Performance testing and security testing are given as subtypes of reliability testing by ISO/IEC (2023a), but these are all listed separately by Firesmith (2015, p. 53).
5. Similarly, random testing is a subtechnique of specification-based testing (ISO/IEC and IEEE, 2022, pp. 7, 22; 2021, pp. 5, 20, Fig. 2; Washizaki, 2024, p. 5-12; Hamburg and Mogyorodi, 2024) but is listed separately by Firesmith (2015, p. 46).
6. Path testing “aims to execute all entry-to-exit control flow paths in a SUT’s control flow graph” (Washizaki, 2024, p. 5-13; similar in Patton, 2006, p. 119), but ISO/IEC and IEEE (2017, p. 316) add that it can also be “designed to execute ... selected paths.”
7. The structure of tours can be defined as either quite general (ISO/IEC and IEEE, 2022, p. 34) or “organized around a special focus” (Hamburg and Mogyorodi, 2024).
8. Alpha testing is performed by “users within the organization developing the software” (ISO/IEC and IEEE, 2017, p. 17), “a small, selected group of potential users” (Washizaki, 2024, p. 5-8), or “roles outside the development organization” conducted “in the developer’s test environment” (Hamburg and Mogyorodi, 2024).
9. “Use case testing” is given as a synonym of “scenario testing” by Hamburg and Mogyorodi (2024) but listed separately by ISO/IEC and IEEE (2022, Fig. 2) and described as a “common form of scenario testing” in (2021, p. 20). This implies that use case testing may instead be a child of user scenario testing (see Table 6.5).
10. The terms “test level” and “test stage” are given as synonyms (Hamburg and Mogyorodi, 2024; implied by Gerrard, 2000a, p. 9), but Washizaki (2024, p. 5-6) says “[test] levels can be distinguished based on the object of testing, the *target*, or on the purpose or *objective*” and calls the former “test stages”, giving the term a child relation (see Section 2.1.3) to “test level” instead. However, the examples of “test stages” listed—unit testing, integration testing, system testing, and acceptance testing (Washizaki, 2024, pp. 5-6 to 5-7)—are commonly categorized as “test levels” (see Section 2.1.1).
11. “Operational acceptance testing” and “production acceptance testing” are given as synonyms by Hamburg and Mogyorodi (2024) but listed separately by Firesmith (2015, p. 30).
12. While Patton (2006, p. 120) implies that condition testing is a subtechnique of path testing, van Vliet (2000, Fig. 13.17) says that multiple condition coverage (which seems to be a synonym of condition coverage (p. 422)) does not subsume and is not subsumed by path coverage.

OG Hass, 2008

13. Load testing is performed with loads “between anticipated conditions of low, typical, and peak usage” (ISO/IEC and IEEE, 2022, p. 5) or loads that are as large as possible (Patton, 2006, p. 86).
14. State testing requires that “all states in the state model ... [are] ‘visited’ ” (ISO/IEC and IEEE, 2021, p. 19), but Patton (2006, pp. 82–83) lists this as only one of its possible criteria.
15. System testing is “conducted on a complete, integrated system” (ISO/IEC and IEEE, 2017, p. 456; similar in Peters and Pedrycz, 2000, Tab. 12.3; van Vliet, 2000, p. 439), but Patton (2006, p. 109) says it can also be done on “at least a major portion” of the product.
16. “Walkthroughs” and “structured walkthroughs” are given as synonyms by Hamburg and Mogyorodi (2024) but Peters and Pedrycz (2000, p. 484) imply that they are different, saying a more structured walkthrough may have specific roles.
17. While Patton (2006, p. 92, emphasis added) says that reviews are “*the* process[es] under which static white-box testing is performed”, van Vliet (2000, pp. 418–419) gives correctness proofs as another example.
- OG Beizer — 18. Kam (2008, p. 46) says “negative testing is related to the testers’ attitude rather than a specific test approach or test design technique”; while ISO/IEC and IEEE (2021) seem to support this idea of negative testing being at a “higher” level than other approaches, they also imply that it is a test technique (pp. 10, 14).
- OG 2015 — 19. While a computation data use is defined as the “use of the value of a variable in *any* type of statement” (ISO/IEC and IEEE, 2021, p. 2; 2017, p. 83, emphasis added), it is often qualified to *not* be a predicate data use (van Vliet, 2000, p. 424; implied by ISO/IEC and IEEE, 2021, p. 27).
- OG ISO1984 — 20. ISO/IEC and IEEE define an “extended entry (decision) table” both as a decision table where the “conditions consist of multiple values rather than simple Booleans” (2021, p. 18) and one where “the conditions and actions are generally described but are incomplete” (2017, p. 175).
- OG 2015 — 21. ISO/IEC and IEEE (2017, p. 83)’s definition of “all-c-uses testing”—testing that aims to execute all data “use[s] of the value of a variable in any type of statement”—is *much* more vague than the definition they give in (2021, p. 27; similar in van Vliet, 2000, p. 425; Peters and Pedrycz, 2000, p. 479)—testing that exercises “control flow sub-paths from each variable definition to each c-use of that definition (with no intervening definitions)”.
22. Since keyword-driven testing can be used for automated *or* manual testing (ISO/IEC and IEEE, 2016, pp. 4, 6), the claim that “test cases can be either manual test cases or keyword test cases” (p. 6) is incorrect.

23. ISO/IEC and IEEE (2022, p. 36) say “A/B testing is not a test case generation technique as test inputs are not generated”, where “test case generation technique” may be a synonym of “test design technique”. However, the inclusion of A/B testing under the heading “Test design and execution” in the same document implies that it may be considered a test technique.⁴
24. The claim that “test cases can be either manual test cases or keyword test cases” (ISO/IEC and IEEE, 2016, p. 6) implies that “keyword-driven testing” could be a synonym of “automated testing” instead of its child, which seems more reasonable (2016, p. 4; 2022, p. 35).
25. Different capitalizations of the abbreviations of “computation data use” and “predicate data use” are used: the lowercase “c-use” and “p-use” (ISO/IEC and IEEE, 2021, pp. 3, 27–29, 35–36, 114–155, 117–118, 129; 2017, p. 124; Peters and Pedrycz, 2000, p. 477, Tab. 12.6) and the uppercase “C-use” and “P-use” (van Vliet, 2000, pp. 424–425).
26. Similarly for “definition-use” (such as in “definition-use path”), both the lowercase “du” (ISO/IEC and IEEE, 2021, pp. 3, 27, 29, 35, 119–121, 129; Peters and Pedrycz, 2000, pp. 478–479) and the uppercase “DU” (van Vliet, 2000, p. 425) are used.
27. Van Vliet (2000, pp. 424–425) specifies that every successor of a data definition use needs to be executed as part of all-uses testing, but this condition is not included elsewhere (ISO/IEC and IEEE, 2021, pp. 28–29; 2017, p. 120; Peters and Pedrycz, 2000, pp. 478–479).
28. All-du-paths testing is usually defined as exercising all “loop-free control flow sub-paths from each variable definition to every use (both p-use and c-use) of that definition (with no intervening definitions)” (ISO/IEC and IEEE, 2021, p. 29; similar in 2017, p. 125; Washizaki, 2024, p. 5–13; Peters and Pedrycz, 2000, p. 479); however, paths containing simple cycles may also be required (van Vliet, 2000, p. 425).
29. Van Vliet (2000, pp. 432–433) says that all-p-uses testing is only stronger than all-edges (branch) testing if there are infeasible paths, but ISO/IEC and IEEE (2021, Fig. F.1) do not specify this caveat.
30. Similarly, van Vliet (2000, pp. 432–433) says that all-du-paths testing is only stronger than all-uses testing if there are infeasible paths, but Washizaki (2024, p. 5–13) does not specify this caveat.
31. Acceptance testing is “usually performed by the purchaser ... with the ... vendor” (ISO/IEC and IEEE, 2017, p. 5), “may or may not involve the developers of the system” (Bourque and Fairley, 2014, p. 4–6), and/or “is often performed under supervision of the user organization” (van Vliet, 2000, p. 439); these descriptions of who the testers are contradict each other *and*

⁴For simplicity, this implied categorization as “technique” is omitted from Table 6.4.

OG Reid, 1996

Q #12: Does this merit counting this as an Ambiguity as well as a Contradiction?

all introduce some uncertainty (“usually”, “may or may not”, and “often”, respectively).

32. Although ad hoc testing is classified as a “technique” (Washizaki, 2024, p. 5-14), it is one in which “no recognized test design technique is used” (Kam, 2008, p. 42).

6.1.4 Ambiguities

The following are cases where information (usually definitions or distinctions between terms; see Sections 6.2.4 and 6.2.5, respectively) is unclear:

Is this a def
flaw?

1. The distinctions between development testing (ISO/IEC and IEEE, 2017, p. 136), developmental testing (Firesmith, 2015, p. 30), and developer testing (p. 39; Gerrard, 2000a, p. 11) are unclear and seem miniscule.
2. Hamburg and Mogyorodi (2024) define “Machine Learning (ML) model testing” and “ML functional performance” in terms of “ML functional performance criteria”, which is defined in terms of “ML functional performance metrics”, which is defined as “a set of measures that relate to the functional correctness of an ML system”. The use of “performance” (or “correctness”) in these definitions is at best ambiguous and at worst incorrect.
3. “Installability testing” is given as a test type (ISO/IEC and IEEE, 2022, p. 22; 2021, p. 38; 2017, p. 228), while “installation testing” is given as a test level (van Vliet, 2000, p. 439; implied by Washizaki, 2024, p. 5-8). Since “installation testing” is not given as an example of a test level throughout the sources that describe them (see Section 2.1.1), it is likely that the term “installability testing” with all its related information should be used instead.
4. Hamburg and Mogyorodi (2024) claim that code inspections are related to peer reviews but Patton (2006, pp. 94–95) makes them quite distinct.

OG 2015

5. “Data definition” is defined as a “statement where a variable is assigned a value” (ISO/IEC and IEEE, 2021, p. 3; 2017, p. 115; similar in 2012, p. 27; van Vliet, 2000, p. 424), but for functional programming languages such as Haskell with immutable variables (Wikibooks Contributors, 2023), this could cause confusion and/or be imprecise.

OG [14]

6. While “error” is defined as “a human action that produces an incorrect result” (Washizaki, 2024, pp. 12-3; van Vliet, 2000, p. 399), Washizaki does not use this consistently, sometimes implying that errors can be intrinsic to software itself (2024, pp. 4-9, 6-5, 7-3, 12-4, 12-9, 12-13).

Q #13: I ignore “syntax errors, runtime errors, and logical errors” (Washizaki, 2024, p. 16-13, cf. p. 18-13) since they seem to be in different domains. Does that make sense?

7. Washizaki (2024, p. 1-1) defines “defect” as “an observable difference between what the software is intended to do and what it does”, but this seems to match the definition of “failure”: the inability of a system “to perform a required function or ... within previously specified limits” (ISO/IEC and

IEEE, 2019a, p. 7; 2010, p. 139; similar in van Vliet, 2000, p. 400) that is “externally visible” (ISO/IEC and IEEE, 2019a, p. 7; similar in van Vliet, 2000, p. 400).

8. Similarly, Washizaki (2024, p. 4-11) says “fault tolerance is a collection of techniques that increase software reliability by detecting errors and then recovering from them or containing their effects if recovery is not possible.” This should either be called “*error* tolerance” or be described as “detecting *faults* and then recovering from them”, since “error” and “fault” have distinct meanings (Washizaki, 2024, p. 5-3; van Vliet, 2000, pp. 399–400). The intent of this term-definition pair is unclear, as the strategies given—“backing up and retrying, using auxiliary code and voting algorithms, and replacing an erroneous value with a phony value that will have a benign effect” (Washizaki, 2024, p. 4-11)—could be used for errors or faults.
9. While ergonomics testing is out of scope (as it tests hardware, not software; see Appendix A.1), its definition of “testing to determine whether a component or system and its input devices are being used properly with correct posture” (Hamburg and Mogyorodi, 2024) seems to focus on how the system is *used* as opposed to the system *itself*.
10. Similarly, end-to-end testing is defined as testing “in which business processes are tested from start to finish under production-like circumstances” (Hamburg and Mogyorodi, 2024); it is unclear whether this tests the business processes themselves or the system’s role in performing them.
11. Hamburg and Mogyorodi (2024) describe the term “software in the loop” as a kind of testing, while the source they reference seems to describe “Software-in-the-Loop-Simulation” as a “simulation environment” that may support software integration testing (Knüvener Mackert GmbH, 2022, p. 153); is this a testing approach or a tool that supports testing?
12. While model testing is said to test the object under test, it seems to describe testing the models themselves (Firesmith, 2015, p. 20); using the models to test the object under test seems to be called “driver-based testing” (p. 33).
13. Similarly, it is ambiguous whether “tool/environment testing” refers to testing the tools/environment *themselves* or *using* them to test the object under test; the wording of its subtypes (Firesmith, 2015, p. 25) seems to imply the former.
14. Retesting and regression testing seem to be separated from the rest of the testing approaches (ISO/IEC and IEEE, 2022, p. 23), but it is not clearly detailed why; Barbosa et al. (2006, p. 3) consider regression testing to be a test level, but since it is not given as an example of a test level throughout the sources that describe them (see Section 2.1.1), it is likely that it is at best not universal and at worst incorrect.

OG ISO/IEC,
2005

15. “Conformance testing” is implied to be a synonym of “compliance testing” by Kam (2008, p. 43) which only makes sense because of the vague definition of the latter: “testing to determine the compliance of the component or system”.

6.1.5 Overlaps

The following are cases where information overlaps, such as nonatomic definitions and terms (see Sections 6.2.4 and 6.2.5, respectively):

1. ISO/IEC and IEEE (2022, p. 34) give the “landmark tour” as an example of “a tour used for exploratory testing”, but they also use the analogy of “a tour guide lead[ing] a tourist through the landmarks of a big city” to describe tours in general. Is the distinction between them the fact that landmark tours are pre-planned and follow a decided-upon sequence (p. 34)?
2. ISO/IEC and IEEE (2017, pp. 469, 470; 2013, p. 9) say that “test level” and “test phase” are synonyms, both meaning a “specific instantiation of [a] test sub-process”, but they have other definitions as well. “Test level” can also refer to the scope of a test process; for example, “across the whole organization” or only “to specific projects” (2022, p. 24) and “test phase” can also refer to the “period of time in the software life cycle” when testing occurs (2017, p. 470), usually after the implementation phase (pp. 420, 509; Perry, 2006, p. 56).
3. ISO/IEC and IEEE (2010, p. 128) define “error” as “a human action that produces an incorrect result”, but also as “an incorrect result” itself. Since faults are inserted when a developer makes an error (pp. 128, 140; Washizaki, 2024, p. 12-3; van Vliet, 2000, pp. 399–400), this means that they are “incorrect results”, making “error” and “fault” synonyms and the distinction between them less useful.
4. Additionally, “error” can also be defined as “the difference between a computed, observed, or measured value or condition and the true, specified, or theoretically correct value or condition” (ISO/IEC and IEEE, 2010, p. 128; similar in Washizaki, 2024, pp. 17-18 to 17-19, 18-7 to 18-8). While this is a widely used definition, particularly in mathematics, it makes some test approaches ambiguous. For example, back-to-back testing is “testing in which two or more variants of a program are executed with the same inputs, the outputs are compared, and errors are analyzed in case of discrepancies” (ISO/IEC and IEEE, 2010, p. 30; similar in Hamburg and Mogyorodi, 2024), which seems to refer to this definition of “error”.
5. The SWEBOK Guide V4 defines “privacy testing” as testing that “assess[es] the security and privacy of users’ personal data to prevent local attacks” (Washizaki, 2024, p. 5-10). This seems to overlap (both in scope and name) with the definition of “security testing” in (ISO/IEC and IEEE, 2022, p. 7):

Q #14: Should we classify this as a Synonym Flaw at all?

testing “conducted to evaluate the degree to which a test item, and associated data and information, [sic] are protected so that” only “authorized persons or systems” can use them as intended.

Q #15:
Shouldn't this be “is”, referring to “test item”?

6. “Orthogonal array testing” (Washizaki, 2024, pp. 5-1, 5-11; implied by Valcheva, 2013, pp. 467, 473; Yu et al., 2011, pp. 1573–1577, 1580) and “operational acceptance testing” (Firesmith, 2015, p. 30) have the same acronym (“OAT”).
7. ISO/IEC and IEEE (2017, p. 228) provide a definition for “inspections and audits”, despite also giving definitions for “inspection” (p. 227) and “audit” (p. 36); while the first term *could* be considered a superset of the latter two, this distinction doesn't seem useful.
8. “Customer acceptance testing” and “contract(ual) acceptance testing” have the same acronym (“CAT”) (Firesmith, 2015, p. 30).
9. The same is true for “hardware-in-the-loop testing” and “human-in-the-loop testing” (“HIL”) (Firesmith, 2015, p. 23), although Preuße et al. (2012, p. 2) use “HiL” for the former.

6.1.6 Redundancies

The following are cases of redundant information:

1. “Ethical hacking testing” is given as a synonym of penetration testing by Washizaki (2024, p. 13-4), which seems redundant; Gerrard (2000b, p. 28) uses the term “ethical hacking” which is clearer.
2. While correct, ISTQB's definition of “specification-based testing” is not helpful: “testing based on an analysis of the specification of the component or system” (Hamburg and Mogyorodi, 2024).
3. The phrase “continuous automated testing” (Gerrard, 2000a, p. 11) is redundant since Continuous Testing (CT) is already a subapproach of automated testing (ISO/IEC and IEEE, 2022, p. 35; Hamburg and Mogyorodi, 2024).

6.2 Flaws by Domain

The following sections list observed flaws grouped by *what* information is flawed as presented in Section 2.2.2. This includes test approach categories (Section 6.2.1), synonym relations (Section 6.2.2), parent-child relations (Section 6.2.3), definitions (Section 6.2.4), labels (Section 6.2.5), scope (Section 6.2.6), and traceability information (Section 6.2.7).

6.2.1 Approach Category Flaws

While the IEEE categorization of testing approaches described in Table 2.1 is useful, it is not without its faults. One issue, which is not inherent to the categorization itself, is the fact that it is not used consistently (see Table 4.1). The most blatant example of this is that ISO/IEC and IEEE (2017, p. 286) describe mutation testing as a methodology, even though this is not one of the categories *they* created! Additionally, the boundaries between approaches within a category may be unclear: “although each technique is defined independently of all others, in practice [sic] some can be used in combination with other techniques” (ISO/IEC and IEEE, 2021, p. 8). For example, “the test coverage items derived by applying equivalence partitioning can be used to identify the input parameters of test cases derived for scenario testing” (p. 8). Even the categories themselves are not consistently defined, and some approaches are categorized differently by different sources; these differences are tracked so they can be analyzed more systematically.

See #21

OG Beizer

1. Hamburg and Mogyorodi (2024) classify ML model testing as a test level, which they define as “a specific instantiation of a test process”: a vague definition that does not match the one in Table 2.1.
2. Kam (2008, p. 46) says “negative testing is related to the testers’ attitude rather than a specific test approach or test design technique”; while ISO/IEC and IEEE (2021) seem to support this idea of negative testing being at a “higher” level than other approaches, they also imply that it is a test technique (pp. 10, 14).
3. Since keyword-driven testing can be used for automated *or* manual testing (ISO/IEC and IEEE, 2016, pp. 4, 6), the claim that “test cases can be either manual test cases or keyword test cases” (p. 6) is incorrect.
4. ISO/IEC and IEEE (2022, p. 36) say “A/B testing is not a test case generation technique as test inputs are not generated”, where “test case generation technique” may be a synonym of “test design technique”. However, the inclusion of A/B testing under the heading “Test design and execution” in the same document implies that it may be considered a test technique.⁵
5. Retesting and regression testing seem to be separated from the rest of the testing approaches (ISO/IEC and IEEE, 2022, p. 23), but it is not clearly detailed why; Barbosa et al. (2006, p. 3) consider regression testing to be a test level, but since it is not given as an example of a test level throughout the sources that describe them (see Section 2.1.1), it is likely that it is at best not universal and at worst incorrect.
6. Although ad hoc testing is classified as a “technique” (Washizaki, 2024, p. 5-14), it is one in which “no recognized test design technique is used” (Kam, 2008, p. 42).

⁵For simplicity, this implied categorization as “technique” is omitted from Table 6.4.

Some category flaws can be detected automatically, such as test approaches with more than one category. These are given in Table 6.4 and include experience-based testing, which is of particular note. ISO/IEC and IEEE categorize experience-based testing as both a test design technique and a test practice on the same page—twice (2022, Fig. 2, p. 34)! These authors say “experience-based testing practices like exploratory testing ... are not ... techniques for designing test cases”, although they “can use ... test techniques” (2021, p. viii), which they support in (2022, p. 33) along with scripted testing. This implies that “experience-based test design techniques” are used *by* the practice of experience-based testing which is not *itself* a test technique (and similarly with scripted testing). If this is the case, it blurs the line between “practice” and “technique”, which may explain why experience-based testing is categorized inconsistently in the literature.

However, this might mean that a practice such as experience-based testing can be viewed as a “class of test case design techniques” (ISO/IEC and IEEE, 2022, p. 4). If *this* is the case, then test approaches that are a collection of specific subtechniques may be considered practices. The following test approaches are each described as a “class”, “family”, or “collection” of techniques by the sources given, which seems to support this:

- Combinatorial testing (ISO/IEC and IEEE, 2022, p. 3; 2021, p. 2; Washizaki, 2024, p. 5-11)
- Data flow testing (2021, p. 3; implied by Washizaki, 2024, p. 5-13)
- Performance(-related) testing (ISO/IEC and IEEE, 2021, p. 38; Moghadam, 2019, p. 118⁶)
- Security testing (implied by ISO/IEC and IEEE, 2021, p. 40)
- Fault tolerance testing (implied by Washizaki, 2024, p. 4-11)

Of the above, security testing is an outlier, since it is consistently categorized as a test type (ISO/IEC and IEEE, 2022, pp. 9, 22, 26–27; 2021, pp. 7, 40, Tab. A.1; 2017, p. 405; implied by its quality (ISO/IEC, 2023a; Washizaki, 2024, p. 13-4); Firesmith, 2015, p. 53) despite consisting of “a number of techniques” (ISO/IEC and IEEE, 2021, p. 40), although this may be distinct from the notion of “test technique” described in Table 2.1. In addition, specification-based testing and structure-based testing may also be considered “families” since they are quite broad with many subtechniques and are described as “complementary” alongside experience-based testing (ISO/IEC and IEEE, 2021, p. 8, Fig. 2).

Subapproaches of experience-based testing, such as error guessing and exploratory testing, are also categorized ambiguously, causing confusion on how categories and parent-child relations (see Section 2.1.3) interact. ISO/IEC and IEEE (2022, p. 34, emphasis added) say that a previous standard (2021) “describes the experience-based test *design technique* of error guessing. Other experience-based test *practices*

⁶The original source describes “performance testing ... as a family of performance-related testing techniques”, but it makes more sense to consider “performance-related testing” as the “family” with “performance testing” being one of the variabilities (see Section 7.4).

include (but are not limited to) exploratory testing ..., tours, attacks, and checklist-based testing”. This seems to imply that error guessing is both a technique *and* a practice, which does not make sense if these categories are orthogonal. Similarly, the conflicting categorizations of beta testing in Table 6.4 may propagate to its children closed beta testing and open beta testing. Since we infer this (see Section 2.3), we omit it from that table and instead include it in Table 6.6. These kinds of inconsistencies between parent and child test approach categorizations may indicate that categories are not transitive or that more thought must be given to them.

Table 6.4: Test approaches with more than one category.

Approach	Category 1	Category 2
Capacity Testing	Technique (ISO/IEC and IEEE, 2021, p. 38)	Type (ISO/IEC and IEEE, 2022, p. 22; 2013, p. 2; implied by its quality (ISO/IEC, 2023a; ISO/IEC and IEEE, 2021, Tab. A.1); Firesmith, 2015, p. 53)
Checklist-based Testing	Practice (ISO/IEC and IEEE, 2022, p. 34)	Technique (Hamburg and Mogyorodi, 2024)
Data-driven Testing	Practice (ISO/IEC and IEEE, 2022, p. 22)	Technique (Kam, 2008, p. 43; OG Fewster and Graham)
End-to-end Testing	Type (Hamburg and Mogyorodi, 2024)	Technique (Firesmith, 2015, p. 47; Sharma et al., 2021, pp. 601, 603, 605–606)
Endurance Testing	Technique (ISO/IEC and IEEE, 2021, p. 38)	Type (ISO/IEC and IEEE, 2013, p. 2; implied by Firesmith, 2015, p. 55)
Experience-based Testing	Practice (ISO/IEC and IEEE, 2022, pp. 22, 34; 2021, p. viii)	Technique (ISO/IEC and IEEE, 2022, pp. 4, 22; 2021, p. 4; Washizaki, 2024, p. 5-13; Hamburg and Mogyorodi, 2024; Firesmith, 2015, pp. 46, 50; implied by ISO/IEC and IEEE, 2022, p. 34)
Exploratory Testing	Practice (ISO/IEC and IEEE, 2022, pp. 20, 22, 34; 2021, p. viii)	Technique (ISO/IEC and IEEE, 2022, p. 34; Washizaki, 2024, p. 5-14; Firesmith, 2015, p. 50; implied by ISO/IEC and IEEE, 2022, p. 33; inferred from experience-based testing)
Load Testing	Technique (ISO/IEC and IEEE, 2021, p. 38)	Type (ISO/IEC and IEEE, 2022, pp. 5, 20, 22; 2017, p. 253; OG IEEE 2013; Hamburg and Mogyorodi, 2024; implied by Firesmith, 2015, p. 54)
Model-based Testing	Practice (ISO/IEC and IEEE, 2022, p. 22; 2021, p. viii)	Technique (Engström and Petersen, 2015, pp. 1–2; Kam, 2008, p. 4; implied by ISO/IEC and IEEE, 2022, p. 32; 2021, p. 7; 2017, p. 469)

Continued on next page

Table 6.4: Test approaches with more than one category. (Continued)

Approach	Category 1	Category 2
Performance Testing	Technique (ISO/IEC and IEEE, 2021, p. 38)	Type (ISO/IEC and IEEE, 2022, pp. 7, 22, 26–27; 2021, p. 7; implied by Firesmith, 2015, p. 53)
Stress Testing	Technique (ISO/IEC and IEEE, 2021, p. 38)	Type (ISO/IEC and IEEE, 2022, pp. 9, 22; 2017, p. 442; implied by Firesmith, 2015, p. 54)
Alpha Testing	Type (implied by Firesmith, 2015, p. 58)	Level (ISO/IEC and IEEE, 2022, p. 22; inferred from acceptance testing)
Attacks	Practice (ISO/IEC and IEEE, 2022, p. 34)	Technique (implied by Hamburg and Mogyorodi, 2024)
Automated Testing	Practice (ISO/IEC and IEEE, 2022, pp. 20, 22)	Technique (implied by ISO/IEC and IEEE, 2022, p. 35)
Back-to-Back Testing	Practice (ISO/IEC and IEEE, 2022, p. 22)	Technique (implied by ISO/IEC and IEEE, 2022, p. 35)
Beta Testing	Type (implied by Firesmith, 2015, p. 58)	Level (ISO/IEC and IEEE, 2022, p. 22; inferred from acceptance testing)
Continuous Testing	Practice (Washizaki, 2024, p. 6-13)	Technique (implied by ISO/IEC and IEEE, 2022, p. 35)
Error Guessing	Practice (implied by ISO/IEC and IEEE, 2022, p. 34)	Technique (ISO/IEC and IEEE, 2022, pp. 4, 22, 34; 2021, pp. 4, 11, 29, Fig. 2; 2013, p. 3; Washizaki, 2024, p. 5-13; Firesmith, 2015, p. 50)

Continued on next page

Table 6.4: Test approaches with more than one category. (Continued)

Approach	Category 1	Category 2
Integration Testing	Technique (implied by Sharma et al., 2021, pp. 601, 603, 605–606)	Level (ISO/IEC and IEEE, 2022, pp. 12, 20–22, 26–27; 2021, p. 6; Washizaki, 2024, p. 5-7; Hamburg and Mogyorodi, 2024; Sakamoto et al., 2013, p. 343; Peters and Pedrycz, 2000, Tab. 12.3; van Vliet, 2000, p. 438; implied by Barbosa et al., 2006, p. 3)
Manual Testing	Practice (ISO/IEC and IEEE, 2022, p. 22)	Technique (implied by ISO/IEC and IEEE, 2022, p. 35)
Mathematical-based Testing	Practice (ISO/IEC and IEEE, 2022, pp. 22, 36)	Technique (implied by ISO/IEC and IEEE, 2022, p. 36)
Penetration Testing	Technique (ISO/IEC and IEEE, 2021, p. 40; Hamburg and Mogyorodi, 2024)	Type (implied by Firesmith, 2015, p. 57; inferred from security testing)
Procedure Testing	Technique (implied by Firesmith, 2015, p. 47)	Type (ISO/IEC and IEEE, 2022, pp. 7, 22; 2021, p. 39, Tab. A.1; 2017, p. 337; OG IEEE, 2013)
Survivability Testing	Technique (Ghosh and Voas, 1999, p. 39)	Type (implied by its quality (ISO/IEC, 2011); inferred from robustness testing and security testing)
Unit Testing	Technique (implied by Engström and Petersen, 2015, pp. 1–2)	Level (ISO/IEC and IEEE, 2022, pp. 12, 20–22, 26–27; 2021, p. 6; 2017, p. 467; 2016, p. 4; Washizaki, 2024, p. 5-6; Hamburg and Mogyorodi, 2024; Sakamoto et al., 2013, p. 343; Peters and Pedrycz, 2000, Tab. 12.3; van Vliet, 2000, p. 438; implied by Barbosa et al., 2006, p. 3)
Volume Testing	Technique (ISO/IEC and IEEE, 2021, p. 38)	Type (implied by Firesmith, 2015, p. 54; inferred from performance-related testing)

Continued on next page

Table 6.4: Test approaches with more than one category. (Continued)

Approach	Category 1	Category 2
Infrastructure Testing	Type (implied by Firesmith, 2015, p. 57)	Level (implied by Gerrard, 2000a, p. 13; see Table 2.1)
Regression Testing	Technique (implied by ISO/IEC and IEEE, 2022, p. 35)	Level (implied by Barbosa et al., 2006, p. 3)

6.2.2 Synonym Relation Flaws

As mentioned in Section 2.1.2, synonyms do not inherently signify a flaw. Unfortunately, there are many instances of incorrect or ambiguous synonyms, such as the following:

1. A component is an “entity with discrete structure ... within a system considered at a particular level of analysis” (ISO/IEC, 2023b) and “the terms module, component, and unit [sic] are often used interchangeably or defined to be subelements of one another in different ways depending upon the context” with no standardized relationship (ISO/IEC and IEEE, 2017, p. 82). For example, Hamburg and Mogyorodi (2024) define them as synonyms while Baresi and Pezzè (2006, p. 107) say “components differ from classical modules for being re-used in different contexts independently of their development”. Additionally, since components are structurally, functionally, or logically discrete (ISO/IEC and IEEE, 2017, p. 419) and “can be tested in isolation” (Hamburg and Mogyorodi, 2024), “unit/component/module testing” could refer to the testing of both a module *and* a specific function in a module, introducing a further level of ambiguity.
2. ISO/IEC and IEEE (2017, pp. 469, 470; 2013, p. 9) say that “test level” and “test phase” are synonyms, both meaning a “specific instantiation of [a] test sub-process”, but they have other definitions as well. “Test level” can also refer to the scope of a test process; for example, “across the whole organization” or only “to specific projects” (2022, p. 24) and “test phase” can also refer to the “period of time in the software life cycle” when testing occurs (2017, p. 470), usually after the implementation phase (pp. 420, 509; Perry, 2006, p. 56).
3. “Use case testing” is given as a synonym of “scenario testing” by Hamburg and Mogyorodi (2024) but listed separately by ISO/IEC and IEEE (2022, Fig. 2) and described as a “common form of scenario testing” in (2021, p. 20). This implies that use case testing may instead be a child of user scenario testing (see Table 6.5).
4. The terms “test level” and “test stage” are given as synonyms (Hamburg and Mogyorodi, 2024; implied by Gerrard, 2000a, p. 9), but Washizaki (2024, p. 5-6) says “[test] levels can be distinguished based on the object of testing, the *target*, or on the purpose or *objective*” and calls the former “test stages”, giving the term a child relation (see Section 2.1.3) to “test level” instead. However, the examples of “test stages” listed—unit testing, integration testing, system testing, and acceptance testing (Washizaki, 2024, pp. 5-6 to 5-7)—are commonly categorized as “test levels” (see Section 2.1.1).
5. “Operational acceptance testing” and “production acceptance testing” are given as synonyms by Hamburg and Mogyorodi (2024) but listed separately by Firesmith (2015, p. 30).

See #14

OG Hass, 2008

6. The differences between the terms “error”, “failure”, “fault”, “defect” are significant and meaningful (ISO/IEC and IEEE, 2010, pp. 128, 139–140; Washizaki, 2024, p. 12-3; van Vliet, 2000, pp. 399–400), but Patton (2006, pp. 13–14) “just call[s] it what it is and get[s] on with it”, abandoning these four terms, “problem”, “incident”, “anomaly”, “variance”, “inconsistency”, “feature” (!), and “a list of unmentionable terms” in favour of “bug”; after all, “there’s no reason to dice words”!
7. Hamburg and Mogyorodi (2024) claim that code inspections are related to peer reviews but Patton (2006, pp. 94–95) makes them quite distinct.
8. “Walkthroughs” and “structured walkthroughs” are given as synonyms by Hamburg and Mogyorodi (2024) but Peters and Pedrycz (2000, p. 484) imply that they are different, saying a more structured walkthrough may have specific roles.
9. The claim that “test cases can be either manual test cases or keyword test cases” (ISO/IEC and IEEE, 2016, p. 6) implies that “keyword-driven testing” could be a synonym of “automated testing” instead of its child, which seems more reasonable (2016, p. 4; 2022, p. 35).
10. Sharma et al. (2021, p. 601) seem to use the terms “grey-box testing” and “(stepwise) code reading” interchangeably, which would incorrectly imply that they are synonyms.
11. Sneed and Göschl (2000, p. 18) give “white-box testing”, “grey-box testing”, and “black-box testing” as synonyms for “module testing”, “integration testing”, and “system testing”, respectively, but this mapping is incorrect; for example, ISO/IEC and IEEE (2017, p. 444) say “structure-based [(or white-box)] testing is not restricted to use at component level and can be used at all levels” and Sakamoto et al. (2013, pp. 345–346) describe “black-box integration testing”.
12. The previous flaw makes the claim that “red-box testing” is a synonym for “acceptance testing” (Sneed and Göschl, 2000, p. 18) lose credibility.
13. Kam (2008, p. 46) gives “mutation testing” as a synonym of “back-to-back testing”; while the two are related (ISO/IEC and IEEE, 2010, p. 30), the variants used in mutation testing are generated or designed to be detected as incorrect by the test suite (Washizaki, 2024, p. 5-15; similar in van Vliet, 2000, pp. 428–429) which is not a requirement of back-to-back testing.
14. “Conformance testing” is implied to be a synonym of “compliance testing” by Kam (2008, p. 43) which only makes sense because of the vague definition of the latter: “testing to determine the compliance of the component or system”.

There are also cases in which a term is given as a synonym to two (or more) terms that are not synonyms themselves. Sometimes, these terms *are* synonyms;

for example, Hamburg and Mogyorodi (2024) say “use case testing”, “user scenario testing”, and “scenario testing” are all synonyms (although there may be a slight distinction; see Table 6.5 and Synonym Flaw 3). However, this does not always make sense. We identify nine such cases through automatic analysis of the generated graphs, listed below (test approaches in *italics* are synonyms with each other, but not with other terms not in italics):

Better way to handle/display this?

1. **Condition Testing:**

- *Branch Condition Combination Testing* (Patton, 2006, p. 120; Sharma et al., 2021, Fig. 1)
- Branch Condition Testing (Hamburg and Mogyorodi, 2024)
- *Decision Testing* (Washizaki, 2024, p. 5-13)

2. **Soak Testing:**

- Endurance Testing (ISO/IEC and IEEE, 2021, p. 39)
- Reliability Testing⁷ (Gerrard, 2000a, Tab. 2; 2000b, Tab. 1, p. 26)

3. **Functional Testing:**

- Behavioural Testing (van Vliet, 2000, p. 399)
- Correctness Testing (Washizaki, 2024, p. 5-7)
- Specification-based Testing (ISO/IEC and IEEE, 2017, p. 196; Kam, 2008, pp. 44–45; van Vliet, 2000, p. 399; implied by ISO/IEC and IEEE, 2021, p. 129; 2017, p. 431)

4. **Link Testing:**

- Branch Testing (implied by ISO/IEC and IEEE, 2021, p. 24)
- Component Integration Testing (Kam, 2008, p. 45)
- Integration Testing (implied by Gerrard, 2000a, p. 13)

5. **Exhaustive Testing:**

- Branch Condition Combination Testing (Peters and Pedrycz, 2000, p. 464)
- Path Testing (van Vliet, 2000, p. 421)

6. **State-based Testing:**

- State Transition Testing (Firesmith, 2015, p. 47; Barbosa et al., 2006, p. 3)
- State-based Web Browser Testing (Doğan et al., 2014, p. 193)

⁷Endurance testing is given as a child of reliability testing by Firesmith (2015, p. 55), although the terms are not synonyms.

7. Static Verification:

- Static Assertion Checking⁸ (Chalin et al., 2006, p. 343)
- Static Testing (implied by Chalin et al., 2006, p. 343)

8. Testing-to-Fail:

- Forcing Exception Testing (Patton, 2006, pp. 66–67, 78)
- Negative Testing (Patton, 2006, pp. 67, 78, 84–87)

9. Invalid Testing:

- Error Tolerance Testing (implied by Kam, 2008, p. 45)
- Negative Testing (Hamburg and Mogyorodi, 2024; implied by ISO/IEC and IEEE, 2021, p. 10)

6.2.3 Parent-Child Relation Flaws

Parent-child relations (defined in Section 2.1.3) are also not immune to flaws, such as the following:

1. Regression testing and retesting are sometimes given as two distinct approaches (ISO/IEC and IEEE, 2022, p. 8; Firesmith, 2015, p. 34), but sometimes regression testing is defined as a form of “selective retesting” (ISO/IEC and IEEE, 2017, p. 372; Washizaki, 2024, pp. 5-8, 6-5, 7-5 to 7-6; Barbosa et al., 2006, p. 3). Moreover, the two possible variations of regression testing given by van Vliet (2000, p. 411) are “retest-all” and “selective retest”, which is possibly the source of the above misconception. This creates a cyclic relation between regression testing and selective retesting.
2. Performance testing and security testing are given as subtypes of reliability testing by ISO/IEC (2023a), but these are all listed separately by Firesmith (2015, p. 53).
3. Similarly, random testing is a subtechnique of specification-based testing (ISO/IEC and IEEE, 2022, pp. 7, 22; 2021, pp. 5, 20, Fig. 2; Washizaki, 2024, p. 5-12; Hamburg and Mogyorodi, 2024) but is listed separately by Firesmith (2015, p. 46).
4. While Patton (2006, p. 120) implies that condition testing is a subtechnique of path testing, van Vliet (2000, Fig. 13.17) says that multiple condition coverage (which seems to be a synonym of condition coverage (p. 422)) does not subsume and is not subsumed by path coverage.

⁸Chalin et al. (2006, p. 343) list Runtime Assertion Checking (RAC) and Software Verification (SV) as “two complementary forms of assertion checking”; based on how the term “static assertion checking” is used by Lahiri et al. (2013, p. 345), it seems like this should be the complement to RAC instead.

5. ISO/IEC and IEEE (2017, p. 228) provide a definition for “inspections and audits”, despite also giving definitions for “inspection” (p. 227) and “audit” (p. 36); while the first term *could* be considered a superset of the latter two, this distinction doesn’t seem useful.
6. Peters and Pedrycz (2000, Fig. 12.31) imply that decision coverage is a child of both c-use coverage *and* p-use coverage; this seems incorrect, since decisions are the result of p-uses and *not* c-uses (ISO/IEC and IEEE, 2021, pp. 5, 27; 2017, p. 332; van Vliet, 2000, p. 424), and only the p-use relation is implied by ISO/IEC and IEEE (2021, Fig. F.1).

OG 2015

OG Reid 1996

Additionally, some self-referential definitions imply that a test approach is a parent of itself. Since these are by nature self-contained within a given source, these are counted *once* as objective flaws within their sources in Tables 6.1 and 6.2. We identify the following three examples through automatic analysis of our generated graphs (see Section 5.2.1):

1. Performance Testing (Gerrard, 2000a, Tab. 2; 2000b, Tab. 1)
2. System Testing (Firesmith, 2015, p. 23)
3. Usability Testing (Gerrard, 2000a, Tab. 2; 2000b, Tab. 1)

Interestingly, performance testing is *not* described as a subapproach of usability testing by (Gerrard, 2000a;b), which would have been more meaningful information to capture.

There are also pairs of synonyms where one is described as a subapproach of the other, abusing the meaning of “synonym” and causing confusion. We identify 17 of these pairs through automatic analysis of the generated graphs, which are given in Table 6.5 (additional pairs where we infer a flaw are given in Section 6.8.2 for completeness). Of particular note is the relation between path testing and exhaustive testing. While van Vliet (2000, p. 421) claims that path testing done completely “is equivalent to exhaustively testing the program”⁹, this overlooks the effects of input data (ISO/IEC and IEEE, 2021, pp. 129–130; Patton, 2006, p. 121; Peters and Pedrycz, 2000, p. 467) and implementation issues (p. 476) on the code’s behaviour. Exhaustive testing requires “all combinations of input values *and* preconditions ... [to be] tested” (ISO/IEC and IEEE, 2022, p. 4, emphasis added; similar in Hamburg and Mogyorodi, 2024; Patton, 2006, p. 121).

⁹The contradictory definitions of path testing given in Definition Flaw 8 add another layer of complexity to this claim.

Table 6.5: Pairs of test approaches with a parent-child *and* synonym relation.

“Child” → “Parent”	Parent-Child Source(s)	Synonym Source(s)
All Transitions Testing → State Transition Testing	(ISO/IEC and IEEE, 2021, p. 19)	(Kam, 2008, p. 15)
Co-existence Testing → Compatibility Testing	(ISO/IEC, 2023a; ISO/IEC and IEEE, 2022, p. 3; 2021, Tab. A.1)	(ISO/IEC and IEEE, 2021, p. 37)
Domain Testing → Specification-based Testing	(Peters and Pedrycz, 2000, Tab. 12.1)	(Washizaki, 2024, p. 5-10)
Fault Tolerance Testing → Robustness Testing ^a	(Firesmith, 2015, p. 56)	(Hamburg and Mogyorodi, 2024)
Functional Testing → Specification-based Testing ^b	(ISO/IEC and IEEE, 2021, p. 38)	(ISO/IEC and IEEE, 2017, p. 196; Kam, 2008, pp. 44–45; van Vliet, 2000, p. 399; implied by ISO/IEC and IEEE, 2021, p. 129; 2017, p. 431)
Orthogonal Array Testing → Pairwise Testing	(Mandl, 1985, p. 1055)	(Washizaki, 2024, p. 5-11; Valcheva, 2013, p. 473)
Path Testing → Exhaustive Testing	(Peters and Pedrycz, 2000, pp. 466–467, 476; implied by Patton, 2006, pp. 120–121)	(van Vliet, 2000, p. 421)
Performance Testing → Performance-related Testing	(ISO/IEC and IEEE, 2022, p. 22; 2021, p. 38)	(Moghadam, 2019, p. 1187)

Continued on next page

Table 6.5: Pairs of test approaches with a parent-child *and* synonym relation. (Continued)

“Child” → “Parent”	Parent-Child Source(s)	Synonym Source(s)
Static Analysis → Static Testing	(ISO/IEC and IEEE, 2022, pp. 9, 17, 25, 28; Hamburg and Mogyorodi, 2024; Gerrard, 2000a, Fig. 4, p. 12; 2000b, p. 3)	(Peters and Pedrycz, 2000, p. 438)
Structure-based Testing → Structural Testing	(Patton, 2006, pp. 105–121)	(ISO/IEC and IEEE, 2022, p. 9; 2017, pp. 443–444; Hamburg and Mogyorodi, 2024; implied by Barbosa et al., 2006, p. 3)
Use Case Testing → Scenario Testing ^d	(ISO/IEC and IEEE, 2021, p. 20; OG Hass, 2008)	(Hamburg and Mogyorodi, 2024)
Branch Condition Combination Testing → Decision Testing	(implied by the caveat of “atomic conditions” in Hamburg and Mogyorodi, 2024)	(Washizaki, 2024, p. 5-13)
Branch Condition Combination Testing → Exhaustive Testing	(implied by Patton, 2006, p. 121)	(Peters and Pedrycz, 2000, p. 464)
Reviews → Structural Analysis	(Patton, 2006, p. 92)	(implied by Patton, 2006, p. 92)
Structural Testing → Structure-based Testing ^c	(Peters and Pedrycz, 2000, p. 447)	(implied by Peters and Pedrycz, 2000, p. 447)

Continued on next page

Table 6.5: Pairs of test approaches with a parent-child *and* synonym relation. (Continued)

“Child” → “Parent”	Parent-Child Source(s)	Synonym Source(s)
Exploratory Testing → Unscripted Testing	(ISO/IEC and IEEE, 2022, p. 33; 2017, p. 174; Firesmith, 2015, p. 45; implied by Washizaki, 2024, p. 5-14)	(implied by Kuřšovs et al., 2013, p. 214)
Beta Testing → User Testing	(implied by Firesmith, 2015, p. 39)	(implied by Firesmith, 2015, p. 39)

^a Fault tolerance testing may also be a subapproach of reliability testing (ISO/IEC and IEEE, 2017, p. 375; Washizaki, 2024, p. 7-10), which is distinct from robustness testing (Firesmith, 2015, p. 53).

^b See Section 6.3.1.

^c Peters and Pedrycz (2000, p. 447) claim that “structural testing subsumes white box testing”, but also say “structure tests are aimed at exercising the internal logic of a software system” and “in white box testing ..., using detailed knowledge of code, one creates a battery of tests in such a way that they exercise all components of the code (say, statements, branches, paths)” on the same page!

^d See Synonym Flaw 3.

6.2.4 Definition Flaws

Perhaps the most interesting category for those seeking to understand how to apply a given test approach, there are many flaws with how test approaches, as well as supporting terms, are defined:

1. ISO/IEC and IEEE (2022, p. 34) give the “landmark tour” as an example of “a tour used for exploratory testing”, but they also use the analogy of “a tour guide lead[ing] a tourist through the landmarks of a big city” to describe tours in general. Is the distinction between them the fact that landmark tours are pre-planned and follow a decided-upon sequence (p. 34)?
2. Integration testing, system testing, and system integration testing are all listed as “common test levels” (ISO/IEC and IEEE, 2022, p. 12; 2021, p. 6), but the latter two are not defined. This makes the relations between these three terms unclear; system integration testing is listed as a child of both integration testing (Hamburg and Mogyorodi, 2024) and system testing (Firesmith, 2015, p. 23).
3. Similarly, component testing, integration testing, and component integration testing are all mentioned by ISO/IEC and IEEE (2017), but the latter is only defined as “testing of groups of related components” (p. 82). As before, the relations between these three terms are unclear; component integration testing is only listed as a child of integration testing (Hamburg and Mogyorodi, 2024).
4. “Software testing” is often defined to exclude static testing (Firesmith, 2015, p. 13; Ammann and Offutt, 2017, p. 222; Peters and Pedrycz, 2000, p. 439), restricting “testing” to mean dynamic validation (Washizaki, 2024, p. 5-1) or verification “in which a system or component is executed” (ISO/IEC and IEEE, 2017, p. 427). However, “terminology is not uniform among different communities, and some use the term ‘testing’ to refer to static techniques¹⁰ as well” (Washizaki, 2024, p. 5-2). This is done by ISO/IEC and IEEE (2022, pp. 16–17) and Gerrard (2000a, pp. 8–9); the former even explicitly *exclude* static testing in another document (2017, p. 440)!
5. ISO/IEC and IEEE (2010, p. 128) define “error” as “a human action that produces an incorrect result”, but also as “an incorrect result” itself. Since faults are inserted when a developer makes an error (pp. 128, 140; Washizaki, 2024, p. 12-3; van Vliet, 2000, pp. 399–400), this means that they are “incorrect results”, making “error” and “fault” synonyms and the distinction between them less useful.
6. Additionally, “error” can also be defined as “the difference between a computed, observed, or measured value or condition and the true, specified, or theoretically correct value or condition” (ISO/IEC and IEEE, 2010, p. 128;

Q #16: I’m counting this as an ambiguous parent relation occurring in ISTQB; is this sufficient?

Q #17: Should we classify this as a Synonym Flaw at all?

¹⁰Not formally defined, but distinct from the notion of “test technique” described in Table 2.1.

similar in Washizaki, 2024, pp. 17-18 to 17-19, 18-7 to 18-8). While this is a widely used definition, particularly in mathematics, it makes some test approaches ambiguous. For example, back-to-back testing is “testing in which two or more variants of a program are executed with the same inputs, the outputs are compared, and errors are analyzed in case of discrepancies” (ISO/IEC and IEEE, 2010, p. 30; similar in Hamburg and Mogyorodi, 2024), which seems to refer to this definition of “error”.

7. The SWEBOK Guide V4 defines “privacy testing” as testing that “assess[es] the security and privacy of users’ personal data to prevent local attacks” (Washizaki, 2024, p. 5-10). This seems to overlap (both in scope and name) with the definition of “security testing” in (ISO/IEC and IEEE, 2022, p. 7): testing “conducted to evaluate the degree to which a test item, and associated data and information, [sic] are protected so that” only “authorized persons or systems” can use them as intended.
8. Path testing “aims to execute all entry-to-exit control flow paths in a SUT’s control flow graph” (Washizaki, 2024, p. 5-13; similar in Patton, 2006, p. 119), but ISO/IEC and IEEE (2017, p. 316) add that it can also be “designed to execute ... selected paths.”
9. The structure of tours can be defined as either quite general (ISO/IEC and IEEE, 2022, p. 34) or “organized around a special focus” (Hamburg and Mogyorodi, 2024).
10. Alpha testing is performed by “users within the organization developing the software” (ISO/IEC and IEEE, 2017, p. 17), “a small, selected group of potential users” (Washizaki, 2024, p. 5-8), or “roles outside the development organization” conducted “in the developer’s test environment” (Hamburg and Mogyorodi, 2024).
11. Hamburg and Mogyorodi (2024) define “Machine Learning (ML) model testing” and “ML functional performance” in terms of “ML functional performance criteria”, which is defined in terms of “ML functional performance metrics”, which is defined as “a set of measures that relate to the functional correctness of an ML system”. The use of “performance” (or “correctness”) in these definitions is at best ambiguous and at worst incorrect.
12. Load testing is performed with loads “between anticipated conditions of low, typical, and peak usage” (ISO/IEC and IEEE, 2022, p. 5) or loads that are as large as possible (Patton, 2006, p. 86).
13. State testing requires that “all states in the state model ... [are] ‘visited’ ” (ISO/IEC and IEEE, 2021, p. 19), but Patton (2006, pp. 82–83) lists this as only one of its possible criteria.
14. System testing is “conducted on a complete, integrated system” (ISO/IEC and IEEE, 2017, p. 456; similar in Peters and Pedrycz, 2000, Tab. 12.3; van

Q #18:
Shouldn’t this be “is”, referring to “test item”?

Vliet, 2000, p. 439), but Patton (2006, p. 109) says it can also be done on “at least a major portion” of the product.

15. While Patton (2006, p. 92, emphasis added) says that reviews are “*the* process[es] under which static white-box testing is performed”, van Vliet (2000, pp. 418–419) gives correctness proofs as another example.

OG Beizer

16. Kam (2008, p. 46) says that the goal of negative testing is “showing that a component or system does not work” which is not true; if robustness is an important quality for the system, then testing the system “in a way for which it was not intended to be used” (Hamburg and Mogyorodi, 2024) (i.e., negative testing) is one way to help test this!

OG 2015

17. While a computation data use is defined as the “use of the value of a variable in *any* type of statement” (ISO/IEC and IEEE, 2021, p. 2; 2017, p. 83, emphasis added), it is often qualified to *not* be a predicate data use (van Vliet, 2000, p. 424; implied by ISO/IEC and IEEE, 2021, p. 27).

OG ISO1984

18. ISO/IEC and IEEE define an “extended entry (decision) table” both as a decision table where the “conditions consist of multiple values rather than simple Booleans” (2021, p. 18) and one where “the conditions and actions are generally described but are incomplete” (2017, p. 175).

19. ISO/IEC and IEEE (2017) use the same definition for “partial correctness” (p. 314) and “total correctness” (p. 480).

OG 2015

20. ISO/IEC and IEEE (2017, p. 83)’s definition of “all-c-uses testing”—testing that aims to execute all data “use[s] of the value of a variable in any type of statement”—is *much* more vague than the definition they give in (2021, p. 27; similar in van Vliet, 2000, p. 425; Peters and Pedrycz, 2000, p. 479)—testing that exercises “control flow sub-paths from each variable definition to each c-use of that definition (with no intervening definitions)”.

OG 2015

21. “Data definition” is defined as a “statement where a variable is assigned a value” (ISO/IEC and IEEE, 2021, p. 3; 2017, p. 115; similar in 2012, p. 27; van Vliet, 2000, p. 424), but for functional programming languages such as Haskell with immutable variables (Wikibooks Contributors, 2023), this could cause confusion and/or be imprecise.

OG [14]

22. While “error” is defined as “a human action that produces an incorrect result” (Washizaki, 2024, pp. 12-3; van Vliet, 2000, p. 399), Washizaki does not use this consistently, sometimes implying that errors can be intrinsic to software itself (2024, pp. 4-9, 6-5, 7-3, 12-4, 12-9, 12-13).

Q #19: I ignore “syntax errors, runtime errors, and logical errors” (Washizaki, 2024, p. 16-13, cf. p. 18-13) since they seem to be in different domains. Does that make sense?

23. While ergonomics testing is out of scope (as it tests hardware, not software; see Appendix A.1), its definition of “testing to determine whether a component or system and its input devices are being used properly with correct posture” (Hamburg and Mogyorodi, 2024) seems to focus on how the system is *used* as opposed to the system *itself*.

24. Similarly, end-to-end testing is defined as testing “in which business processes are tested from start to finish under production-like circumstances” (Hamburg and Mogyorodi, 2024); it is unclear whether this tests the business processes themselves or the system’s role in performing them.
25. Hamburg and Mogyorodi (2024) describe the term “software in the loop” as a kind of testing, while the source they reference seems to describe “Software-in-the-Loop-Simulation” as a “simulation environment” that may support software integration testing (Knüvener Mackert GmbH, 2022, p. 153); is this a testing approach or a tool that supports testing?
26. While correct, ISTQB’s definition of “specification-based testing” is not helpful: “testing based on an analysis of the specification of the component or system” (Hamburg and Mogyorodi, 2024).
27. While model testing is said to test the object under test, it seems to describe testing the models themselves (Firesmith, 2015, p. 20); using the models to test the object under test seems to be called “driver-based testing” (p. 33).
28. Similarly, it is ambiguous whether “tool/environment testing” refers to testing the tools/environment *themselves* or *using* them to test the object under test; the wording of its subtypes (Firesmith, 2015, p. 25) seems to imply the former.
29. The acronym for System of Systems (SoS) (ISO/IEC and IEEE, 2019b) is used but not defined by Firesmith (2015, p. 23).
30. Van Vliet (2000, p. 425) defines many types of data flow coverage, including all-p-uses, all-p-uses/some-c-uses, and all-c-uses/some-p-uses, but excludes all-c-uses, which is implied by these definitions and defined elsewhere (ISO/IEC and IEEE, 2021, p. 27; 2017, p. 83; Peters and Pedrycz, 2000, p. 479).
31. Van Vliet (2000, pp. 424–425) specifies that every successor of a data definition use needs to be executed as part of all-uses testing, but this condition is not included elsewhere (ISO/IEC and IEEE, 2021, pp. 28–29; 2017, p. 120; Peters and Pedrycz, 2000, pp. 478–479).
32. All-du-paths testing is usually defined as exercising all “loop-free control flow sub-paths from each variable definition to every use (both p-use and c-use) of that definition (with no intervening definitions)” (ISO/IEC and IEEE, 2021, p. 29; similar in 2017, p. 125; Washizaki, 2024, p. 5–13; Peters and Pedrycz, 2000, p. 479); however, paths containing simple cycles may also be required (van Vliet, 2000, p. 425).
33. Van Vliet (2000, pp. 432–433) says that all-p-uses testing is only stronger than all-edges (branch) testing if there are infeasible paths, but ISO/IEC and IEEE (2021, Fig. F.1) do not specify this caveat.

OG Reid, 1996

34. Similarly, van Vliet (2000, pp. 432–433) says that all-du-paths testing is only stronger than all-uses testing if there are infeasible paths, but Washizaki (2024, p. 5-13) does not specify this caveat.
35. Acceptance testing is “usually performed by the purchaser ... with the ... vendor” (ISO/IEC and IEEE, 2017, p. 5), “may or may not involve the developers of the system” (Bourque and Fairley, 2014, p. 4-6), and/or “is often performed under supervision of the user organization” (van Vliet, 2000, p. 439); these descriptions of who the testers are contradict each other *and* all introduce some uncertainty (“usually”, “may or may not”, and “often”, respectively).
36. Bas (2024, p. 16) lists “three [backup] location categories: local, offsite and cloud based [sic]” but does not define or discuss “offsite backups” (pp. 16-17).
37. Gerrard (2000a, Tab. 2) makes a distinction between “transaction verification” and “transaction testing” and uses the phrase “transaction flows” (Fig. 5) but doesn’t explain them.
38. Availability testing isn’t assigned to a test priority (Gerrard, 2000a, Tab. 2), despite the claim that “the test types¹¹ have been allocated a slot against the four test priorities” (p. 13); usability testing and/or performance testing would have been good candidates.

Q #20: Does this merit counting this as an Ambiguity as well as a Contradiction?

Also of note: (ISO/IEC and IEEE, 2022; 2021), from the ISO/IEC/IEEE 29119 family of standards, mention the following 23 test approaches without defining them. This means that out of the 114 test approaches they mention, about 20% have no associated definition!

However, the previous version of this standard, (2013), generally explained two, provided references for two, and explicitly defined one of these terms, for a total of five definitions that could (should) have been included in (2022)! These terms have been underlined, *italicized*, and **bolded**, respectively. Additionally, entries marked with an asterisk* were defined (at least partially) in (2017), which would have been available when creating this family of standards. These terms bring the total count of terms that could (should) have been defined to nine; almost 40% of undefined test approaches could have been defined!

- Acceptance Testing*
- Alpha Testing*
- Beta Testing*
- Capture-Replay Driven Testing
- Data-driven Testing

¹¹“Each type of test addresses a different risk area” (Gerrard, 2000a, p. 12), which is distinct from the notion of “test type” described in Table 2.1.

- Error-based Testing
- Factory Acceptance Testing
- Fault Injection Testing
- Functional Suitability Testing (also mentioned but not defined in (ISO/IEC and IEEE, 2017))
- Integration Testing*
- Model Verification
- Operational Acceptance Testing
- Orthogonal Array Testing
- Production Verification Testing
- Recovery Testing* (Failover/Recovery Testing, Back-up/Recovery Testing, **Backup and Recovery Testing***, Recovery*; see Section 6.5)
- Response-Time Testing
- *Reviews* (ISO/IEC 20246) (Code Reviews*)
- Scalability Testing (defined as a synonym of “capacity testing”; see Section 6.6)
- Statistical Testing
- System Integration Testing (System Integration*)
- System Testing* (also mentioned but not defined in (ISO/IEC and IEEE, 2013))
- *Unit Testing** (IEEE Std 1008-1987, IEEE Standard for Software Unit Testing implicitly listed in the bibliography!)
- User Acceptance Testing

6.2.5 Label Flaws

While some flaws exist because the definition of a term is wrong, others exist because a term’s *name* or *label* is wrong! Terms can be thought of as definition-label pairs, but there is a meaningful distinction between definition flaws and label flaws, which we define in Section 2.2.2. We observe the following label flaws:

1. Since errors are distinct from defects/faults (ISO/IEC and IEEE, 2010, pp. 128, 140; Washizaki, 2024, p. 12-3; van Vliet, 2000, pp. 399–400), error guessing should instead be called “defect guessing” if it is based on a “checklist of potential defects” (ISO/IEC and IEEE, 2021, p. 29) or “fault guessing” if it is a “fault-based technique” (Bourque and Fairley, 2014, p. 4-9) that “anticipate[s] the most plausible faults in each SUT” (Washizaki, 2024, p. 5-13). One (or both) of these proposed terms may be useful in tandem with “error guessing”, which would focus on errors as traditionally defined; this would be a subapproach of error-based testing (implied by van Vliet, 2000, p. 399).
2. Similarly, “fault seeding” is not a synonym of “error seeding” as claimed by ISO/IEC and IEEE (2017, p. 165) and van Vliet (2000, p. 427). The term “error seeding”, also used by Firesmith (2015, p. 34), should be abandoned in favour of “fault seeding”, as it is defined as the “process of intentionally adding known faults to those already in a computer program ... [to] estimat[e] the number of faults remaining” (ISO/IEC and IEEE, 2017, p. 165) based on the ratio between the number of new faults and the number of introduced faults that were discovered (van Vliet, 2000, p. 427).
3. The distinctions between development testing (ISO/IEC and IEEE, 2017, p. 136), developmental testing (Firesmith, 2015, p. 30), and developer testing (p. 39; Gerrard, 2000a, p. 11) are unclear and seem miniscule.
4. “Orthogonal array testing” (Washizaki, 2024, pp. 5-1, 5-11; implied by Valcheva, 2013, pp. 467, 473; Yu et al., 2011, pp. 1573–1577, 1580) and “operational acceptance testing” (Firesmith, 2015, p. 30) have the same acronym (“OAT”).
5. “Installability testing” is given as a test type (ISO/IEC and IEEE, 2022, p. 22; 2021, p. 38; 2017, p. 228), while “installation testing” is given as a test level (van Vliet, 2000, p. 439; implied by Washizaki, 2024, p. 5-8). Since “installation testing” is not given as an example of a test level throughout the sources that describe them (see Section 2.1.1), it is likely that the term “installability testing” with all its related information should be used instead.
6. A typo in (ISO/IEC and IEEE, 2021, Fig. 2) means that “specification-based techniques” is listed twice, when the latter should be “structure-based techniques”.
7. Washizaki (2024, p. 1-1) defines “defect” as “an observable difference between what the software is intended to do and what it does”, but this seems to match the definition of “failure”: the inability of a system “to perform a required function or ... within previously specified limits” (ISO/IEC and IEEE, 2019a, p. 7; 2010, p. 139; similar in van Vliet, 2000, p. 400) that is “externally visible” (ISO/IEC and IEEE, 2019a, p. 7; similar in van Vliet, 2000, p. 400).

Is this a def
flaw?

OG ISO/IEC,
2005

8. Similarly, Washizaki (2024, p. 4-11) says “fault tolerance is a collection of techniques that increase software reliability by detecting errors and then recovering from them or containing their effects if recovery is not possible.” This should either be called “*error* tolerance” or be described as “detecting *faults* and then recovering from them”, since “error” and “fault” have distinct meanings (Washizaki, 2024, p. 5-3; van Vliet, 2000, pp. 399–400). The intent of this term-definition pair is unclear, as the strategies given—“backing up and retrying, using auxiliary code and voting algorithms, and replacing an erroneous value with a phony value that will have a benign effect” (Washizaki, 2024, p. 4-11)—could be used for errors or faults.
9. “Ethical hacking testing” is given as a synonym of penetration testing by Washizaki (2024, p. 13-4), which seems redundant; Gerrard (2000b, p. 28) uses the term “ethical hacking” which is clearer.
10. “Customer acceptance testing” and “contract(ual) acceptance testing” have the same acronym (“CAT”) (Firesmith, 2015, p. 30).
11. The same is true for “hardware-in-the-loop testing” and “human-in-the-loop testing” (“HIL”) (Firesmith, 2015, p. 23), although Preuße et al. (2012, p. 2) use “HiL” for the former.
12. Different capitalizations of the abbreviations of “computation data use” and “predicate data use” are used: the lowercase “c-use” and “p-use” (ISO/IEC and IEEE, 2021, pp. 3, 27–29, 35–36, 114–155, 117–118, 129; 2017, p. 124; Peters and Pedrycz, 2000, p. 477, Tab. 12.6) and the uppercase “C-use” and “P-use” (van Vliet, 2000, pp. 424–425).
13. Similarly for “definition-use” (such as in “definition-use path”), both the lowercase “du” (ISO/IEC and IEEE, 2021, pp. 3, 27, 29, 35, 119–121, 129; Peters and Pedrycz, 2000, pp. 478–479) and the uppercase “DU” (van Vliet, 2000, p. 425) are used.
14. The phrase “continuous automated testing” (Gerrard, 2000a, p. 11) is redundant since Continuous Testing (CT) is already a subapproach of automated testing (ISO/IEC and IEEE, 2022, p. 35; Hamburg and Mogyorodi, 2024).
15. Kam (2008) misspells “state-based” as “state-base” (pp. 13, 15) and “stated-base” (Tab. 1).

6.2.6 Scope Flaws

Some information is given that is technically correct, but it is either given too narrowly, too broadly, or when it should not have been (i.e., it is irrelevant)! We observe the following scope flaws as defined in Section 2.2.2:

1. Washizaki (2024, p. 5-4) says that quality improvement, along with quality assurance, is an aspect of testing that involves “defining methods, tools,

skills, and practices to achieve the specific quality level and objective”; while testing that a system possesses certain qualities is in scope, actively improving the system in response is *not* part of testing.

2. The terms “acceleration tolerance testing” and “acoustic tolerance testing” seem to only refer to software testing in (Firesmith, 2015, p. 56); elsewhere, they seem to refer to testing the acoustic tolerance of rats (Holley et al., 1996) or the acceleration tolerance of astronauts (Morgun et al., 1999, p. 11), aviators (Howe and Johnson, 1995, pp. 27, 42), or catalysts (Liu et al., 2023, p. 1463), which don’t exactly seem relevant...
3. The definition of “math testing” given by Hamburg and Mogyorodi (2024) is too specific to be useful, likely taken from an example instead of a general definition: “testing to determine the correctness of the pay table implementation, the random number generator results, and the return to player computations”.
4. A similar issue exists with multiplayer testing, where its definition specifies “the casino game world” (Hamburg and Mogyorodi, 2024).
5. “Par sheet testing” from (Hamburg and Mogyorodi, 2024) seems to refer to the specific example mentioned in Mistake 12 and does not seem more widely applicable, since a “PAR sheet” is “a list of all the symbols on each reel of a slot machine” (Bluejay, 2024).

6.2.7 Traceability Flaws

The following are examples of traceability flaws, as defined in Section 2.2.2:

1. Washizaki (2024, p. 5-9)’s definition of “elasticity testing” only cites a single source **that doesn’t contain the words “elasticity” or “elastic”!**
2. Kam (2008, p. 42) says “See *boundary value analysis*,” for the glossary entry of “boundary value testing” but does not include “boundary value analysis” in the glossary.
3. The source that Hamburg and Mogyorodi (2024) cite for the definition of “test type” does not seem to actually provide a definition.
4. The same is true for “visual testing” (Hamburg and Mogyorodi, 2024).
5. The same is true for “security attack” (Hamburg and Mogyorodi, 2024).
6. Doğan et al. (2014, p. 184) claim that Sakamoto et al. (2013) define “prime path coverage”, but they do not.

6.3 Functional Testing

“Functional testing” is described alongside many other, likely related, terms. This leads to confusion about what distinguishes these terms, as shown by the following five:

6.3.1 Specification-based Testing

This is defined as “testing in which the principal test basis is the external inputs and outputs of the test item” (ISO/IEC and IEEE, 2022, p. 9). This agrees with a definition of “functional testing”: “testing that ... focuses solely on the outputs generated in response to selected inputs and execution conditions” (ISO/IEC and IEEE, 2017, p. 196). Notably, ISO/IEC and IEEE (2017) lists both as synonyms of “black-box testing” (pp. 431, 196, respectively), despite them sometimes being defined separately. For example, the International Software Testing Qualifications Board (ISTQB) defines “specification-based testing” as “testing based on an analysis of the specification of the component or system” (and gives “black-box testing” as a synonym) and “functional testing” as “testing performed to evaluate if a component or system satisfies functional requirements” (specifying no synonyms) (Hamburg and Mogyorodi, 2024); the latter references ISO/IEC and IEEE (2017, p. 196) (“testing conducted to evaluate the compliance of a system or component with specified functional requirements”) which *has* “black-box testing” as a synonym, and mirrors ISO/IEC and IEEE (2022, p. 21) (testing “used to check the implementation of functional requirements”). Overall, specification-based testing (ISO/IEC and IEEE, 2022, pp. 2-4, 6-9, 22) and black-box testing (Washizaki, 2024, p. 5-10; Souza et al., 2017, p. 3) are test design techniques used to “derive corresponding test cases” (ISO/IEC and IEEE, 2022, p. 11) from “selected inputs and execution conditions” (ISO/IEC and IEEE, 2017, p. 196).

6.3.2 Correctness Testing

Washizaki (2024, p. 5-7) says “test cases can be designed to check that the functional specifications are correctly implemented, which is variously referred to in the literature as conformance testing, correctness testing or functional testing”; this mirrors previous definitions of “functional testing” (ISO/IEC and IEEE, 2022, p. 21; 2017, p. 196) but groups it with “correctness testing”. Since “correctness” is a software quality (ISO/IEC and IEEE, 2017, p. 104; Washizaki, 2024, p. 3-13) which is what defines a “test type” (ISO/IEC and IEEE, 2022, p. 15) (see Section 3.3), it seems consistent to label “functional testing” as a “test type” (ISO/IEC and IEEE, 2022, pp. 15, 20, 22; 2021, pp. 7, 38, Tab. A.1; 2016, p. 4). However, this conflicts with its categorization as a “technique” if considered a synonym of specification-based testing (see Section 6.3.1). Additionally, “correctness testing” is listed separately from “functionality testing” by Firesmith (2015, p. 53).

van Vliet (2000, p. 399) may list these as synonyms; investigate

6.3.3 Conformance Testing

Testing that ensures “that the functional specifications are correctly implemented”, and can be called “conformance testing” or “functional testing” (Washizaki, 2024, p. 5-7). “Conformance testing” is later defined as testing used “to verify that the SUT conforms to standards, rules, specifications, requirements, design, processes, or practices” (Washizaki, 2024, p. 5-7). This definition seems to be a superset of testing methods mentioned earlier as the latter includes “standards, rules, requirements, design, processes, ... [and]” practices in *addition* to specifications!

A complicating factor is that “compliance testing” is also (plausibly) given as a synonym of “conformance testing” (Kam, 2008, p. 43). However, “conformance testing” can also be defined as testing that evaluates the degree to which “results ... fall within the limits that define acceptable variation for a quality requirement” (ISO/IEC and IEEE, 2017, p. 93), which seems to describe something different.

OG PMBOK 5th ed.

6.3.4 Functional Suitability Testing

Procedure testing is called a “type of functional suitability testing” (ISO/IEC and IEEE, 2022, p. 7) but no definition of that term is given. “Functional suitability” is the “capability of a product to provide functions that meet stated and implied needs of intended users when it is used under specified conditions”, including meeting “the functional specification” (ISO/IEC, 2023a). This seems to align with the definition of “functional testing” as related to “black-box/specification-based testing”. “Functional suitability” has three child terms: “functional completeness” (the “capability of a product to provide a set of functions that covers all the specified tasks and intended users’ objectives”), “functional correctness” (the “capability of a product to provide accurate results when used by intended users”), and “functional appropriateness” (the “capability of a product to provide functions that facilitate the accomplishment of specified tasks and objectives”) (ISO/IEC, 2023a). Notably, “functional correctness”, which includes precision and accuracy (ISO/IEC, 2023a; Hamburg and Mogyorodi, 2024), seems to align with the quality/ies that would be tested by “correctness” testing.

6.3.5 Functionality Testing

“Functionality” is defined as the “capabilities of the various ... features provided by a product” (ISO/IEC and IEEE, 2017, p. 196) and is said to be a synonym of “functional suitability” (Hamburg and Mogyorodi, 2024), although it seems like it should really be a synonym of “functional completeness” based on (ISO/IEC, 2023a), which would make “functional suitability” a subapproach. Its associated test type is implied to be a subapproach of build verification testing (Hamburg and Mogyorodi, 2024) and made distinct from “functional testing”; interestingly, security is described as a subapproach of both non-functional and functionality testing (Gerrard, 2000a, Tab. 2). “Functionality testing” is listed separately from “correctness testing” by Firesmith (2015, p. 53).

6.4 Operational (Acceptance) Testing

(CONTRA, LABELS)

There are two names that the literature gives to this test approach:

- *Operational Acceptance Testing (OAT)* (ISO/IEC and IEEE, 2022, p. 22; Hamburg and Mogorodi, 2024) and
- *Operational Testing (OT)* (ISO/IEC, 2018; ISO/IEC and IEEE, 2017, p. 303; Washizaki, 2024, p. 6-9, in the context of software engineering operations; Bourque and Fairley, 2014, pp. 4-6, 4-9).

(CONTRA, SYNS)

Firesmith (2015, p. 30) lists the above terms separately, but they are considered synonyms elsewhere (LambdaTest, 2024; Bocchino and Hamilton, 1996); since Firesmith does not define these terms, it is hard to evaluate his distinction.

6.5 Recovery Testing

“Recovery testing” is “testing ... aimed at verifying software restart capabilities after a system crash or other disaster” (Washizaki, 2024, p. 5-9) including “recover[ing] the data directly affected and re-establish[ing] the desired state of the system” (ISO/IEC, 2023a; similar in Washizaki, 2024, p. 7-10) so that the system “can perform required functions” (ISO/IEC and IEEE, 2017, p. 370). However, the literature also describes similar test approaches with vague or non-existent distinctions between them. We describe these approaches and their flaws here and present the relations between them in Figure 7.1a.

- *Recoverability testing* evaluates “how well a system or software can recover data during an interruption or failure” (Washizaki, 2024, p. 7-10; similar in ISO/IEC, 2023a) and “re-establish the desired state of the system” (2023a). Kam (2008, p. 47) gives this as a synonym for “recovery testing”.
- *Disaster/recovery testing* evaluates if a system can “return to normal operation after a hardware or software failure” (ISO/IEC and IEEE, 2017, p. 140) or if “operation of the test item can be transferred to a different operating site and ... be transferred back again once the failure has been resolved” (2021, p. 37).
 - (OVER, DEFS) These two definitions seem to describe different aspects of the system, where the first is intrinsic to the hardware/software and the second might not be, making this term nonatomic.
- *Backup and recovery testing* “measures the degree to which system state can be restored from backup within specified parameters of time, cost, completeness, and accuracy in the event of failure” (ISO/IEC and IEEE, 2013,

p. 2). This may be what is meant by “recovery testing” in the context of performance-related testing (2022, Fig. 2).

- *Backup/recovery testing* determines the ability of a system “to restor[e] from back-up memory in the event of failure, without transfer[ing] to a different operating site or back-up system” (ISO/IEC and IEEE, 2021, p. 37).
 - (CONTRA, PARS) This given as a subtype of “disaster/recovery testing” which tests if “operation of the test item can be transferred to a different operating site” (2021, p. 37), even though this is *explicitly* excluded from its definition on the same page!
 - (OVER, LABELS) Its name is also quite similar to “backup and recovery testing”, adding further confusion.
- *Failover/recovery testing* determines the ability “to mov[e] to a back-up system in the event of failure, without transfer[ing] to a different operating site” (ISO/IEC and IEEE, 2021, p. 37).
 - (CONTRA, PARS) This is also given as a subtype of “disaster/recovery testing” which tests if “operation of the test item can be transferred to a different operating site” (p. 37), even though this is *explicitly* excluded from its definition on the same page!
 - (AMBI, PARS) While not explicitly related to recovery, *failover testing* “validates the SUT’s ability to manage heavy loads or unexpected failure to continue typical operations” (Washizaki, 2024, p. 5-9) by entering a “backup operational mode in which [these responsibilities] ... are assumed by a secondary system” (Hamburg and Mogyorodi, 2024). Its name implies that it is a child of “failover/recovery testing” but its definition makes it more broad (as it includes handling “heavy loads” where failover/recovery testing does not) which may reverse the direction of this relation.
 - (AMBI, SYNS) Firesmith (2015, p. 56) uses the term “failover and recovery testing” which may be a synonym of “failover/recovery testing”.
- *Restart & recovery (testing)* is listed as a test approach by Gerrard (2000a, Fig. 5) but is not defined (MISS, DEFS) and may simply be a synonym to “recovery testing” (AMBI, SYNS).

6.6 Scalability Testing

(CONTRA, SYNS)

ISO/IEC and IEEE (2021, p. 39) give “scalability testing” as a synonym of “capacity testing” while other sources differentiate between the two (Firesmith, 2015, p. 53; Bas, 2024, pp. 22–23).

(CONTRA, DEFS)

ISO/IEC and IEEE (2021, p. 39) also include the external modification of the system as part of “scalability” but ISO/IEC (2023a) describe it as testing the “capability of a product to handle growing or shrinking workloads or to adapt its capacity to handle variability”, implying that this is done by the system itself.

(WRONG, LABELS)

Washizaki (2024, p. 5-9) says “scalability testing evaluates the capability to use and learn the system and the user documentation” and “focuses on the system’s effectiveness in supporting user tasks and the ability to recover from user errors”. This seems to define “usability testing” with elements of functional and recovery testing.

6.7 Compatibility Testing

(OVER, DEFS)

“Compatibility testing” is defined as “testing that measures the degree to which a test item can function satisfactorily alongside other independent products in a shared environment (co-existence), and where necessary, exchanges information with other systems or components (interoperability)” (ISO/IEC and IEEE, 2022, p. 3). This definition is nonatomic as it combines the ideas of “co-existence” and “interoperability”.

(WRONG, SYNS)

The “interoperability” element of “compatibility testing” is explicitly excluded by ISO/IEC and IEEE (2021, p. 37), (incorrectly) implying that “compatibility testing” and “co-existence testing” are synonyms.

(AMBI, SYNS)

Furthermore, the definition of “compatibility testing” in Kam (2008, p. 43) unhelpfully says “see *interoperability testing*”, adding another layer of confusion to the direction of their relationship.

(WRONG, LABELS)

ISO/IEC and IEEE (2022, pp. 22, 43) say “interoperability testing helps confirm that applications can work on multiple operating systems and devices”, but this seems to instead describe “portability testing”, which evaluates the “capability of a product to be adapted to changes in its requirements, contexts of use, or system environment” (ISO/IEC, 2023a; similar in ISO/IEC and IEEE, 2022, p. 7; 2017, pp. 184, 329; Hamburg and Mogyorodi, 2024), such as being “transferred from one hardware ... environment to another” (ISO/IEC and IEEE, 2021, p. 39).

6.8 Inferred Flaws

Throughout our research, we infer many potential flaws as described in Section 2.3. Some of these have a conflicting source while others do not. Since these are more subjective and are based on our own judgement, we exclude them from any counts of the numbers of flaws but give them here for completeness.

6.8.1 Inferred Synonym Flaws

See Section 6.2.2.

1. Production Acceptance Testing:

- Operational (Acceptance) Testing (Hamburg and Mogyorodi, 2024)¹²
- Production Verification Testing¹³

2. Structure-based Testing:

- Implementation-oriented Testing
- Structural Testing (implied by Peters and Pedrycz, 2000, p. 447)

3. Operational (Acceptance) Testing:

- Field Testing
- Qualification Testing

Additionally, Kam (2008, p. 46) gives “program testing” as a synonym of “component testing” but it probably should be a synonym of “system testing” instead.

6.8.2 Inferred Parent Flaws

As discussed in Section 6.2.3, some pairs of synonyms also have a parent-child relation, abusing the meaning of “synonym” and causing confusion. While Table 6.5 gives the cases where both relations are supported by the literature, some are less explicit. For example, while “*dynamic testing* is sometimes called ... dynamic analysis” (Peters and Pedrycz, 2000, p. 438; implied by ISO/IEC and IEEE, 2017, p. 149; Hamburg and Mogyorodi, 2024), it could be inferred from their static counterparts that dynamic analysis is a *child* of dynamic testing (see ISO/IEC and IEEE, 2022, pp. 9, 17, 25, 28). Additionally, the following automatically generated lists contain examples where at least one of these conflicting relations is *not* supported by the literature but may, nonetheless, be correct. The relations in the first two lists are explicitly given in the literature but may be incorrect, while those in the third list are unsubstantiated by the literature and require more thought before a recommendation can be made.

¹²See Synonym Flaw 5.

¹³“Production acceptance testing” (Firesmith, 2015, p. 30) seems to be the same as “production verification testing” (ISO/IEC and IEEE, 2022, p. 22) but neither is defined.

Pairs given a parent-child relation

1. Programmer Testing → Developer Testing (Firesmith, 2015, p. 39)

Pairs given a synonym relation

1. Structured Walkthroughs → Walkthroughs¹⁴ (Hamburg and Mogyorodi, 2024)
2. Functionality Testing → Functional Suitability Testing (implied by Hamburg and Mogyorodi, 2024)

Pairs that could have a parent/child *or* synonym relation

1. Computation Flow Testing → Computation Testing
2. Field Testing → Operational Testing
3. Monkey Testing → Fuzz Testing
4. Organization-based Testing → Role-based Testing¹⁵
5. Scenario-based Evaluations → Scenario-based Testing
6. System Qualification Testing → System Testing

In addition to this type of flaw, (Gerrard, 2000a, Tab. 2) does *not* give “functionality testing” as a parent of “end-to-end functionality testing”. Finally, Patton (2006, p. 119) says that branch testing is “the simplest form of path testing” which is also implied by ISO/IEC and IEEE (2021, Fig. F.1) and van Vliet (2000, p. 433). This is true in the example Patton gives, but is not necessarily generalizable; one could test the behaviour at branches without testing even a *subset* of complete paths, which ISO/IEC and IEEE (2017, p. 316) give as a definition of “path testing” (see Definition Flaw 8)!

6.8.3 Inferred Category Flaws

See Section 6.2.1.

Table 6.6: Test approaches inferred to have more than one category.

Approach	Category 1	Category 2
A/B Testing	Practice (ISO/IEC and IEEE, 2022, Fig. 2)	Type (inferred from usability testing)

Continued on next page

¹⁴See Synonym Flaw 8.

¹⁵The distinction between organization- and role-based testing in (Firesmith, 2015, pp. 17, 37, 39) seems arbitrary, but further investigation may prove it to be meaningful.

Table 6.6: Test approaches inferred to have more than one category. (Continued)

Approach	Category 1	Category 2
Big-Bang Testing	Level (inferred from integration testing)	Technique (Sharma et al., 2021, pp. 601, 603, 605–606)
Bottom-Up Testing	Level (inferred from integration testing)	Technique (Sharma et al., 2021, pp. 601, 603, 605–606)
Fuzz Testing	Practice (inferred from mathematical-based testing)	Technique (Hamburg and Mogyorodi, 2024; Firesmith, 2015, p. 51; implied by ISO/IEC and IEEE, 2022, p. 36)
Memory Management Testing	Technique (ISO/IEC and IEEE, 2021, p. 38)	Type (inferred from performance-related testing)
Privacy Testing	Technique (ISO/IEC and IEEE, 2021, p. 40)	Type (inferred from security testing)
Sandwich Testing	Level (inferred from integration testing)	Technique (Sharma et al., 2021, pp. 601, 603, 605–606)
Security Audits	Technique (ISO/IEC and IEEE, 2021, p. 40)	Type (inferred from security testing)
Vulnerability Scanning	Technique (ISO/IEC and IEEE, 2021, p. 40)	Type (inferred from security testing)
Closed Beta Testing	Level (inferred from beta testing)	Type (implied by Firesmith, 2015, p. 58)
Open Beta Testing	Level (inferred from beta testing)	Type (implied by Firesmith, 2015, p. 58)
Interface Testing	Type? (Kam, 2008, p. 45)	Level (implied by ISO/IEC and IEEE, 2017, p. 235; Sakamoto et al., 2013, p. 343; inferred from integration testing)

6.8.4 Other Inferred Flaws

The following are flaws that, if were more concrete, would also be included alongside the other flaws:

- “Fuzz testing” is “tagged” (?) as “artificial intelligence” (ISO/IEC and IEEE, 2022, p. 5).
- Gerrard’s definition for “security audits” seems too specific, only applying to “the products installed on a site” and “the known vulnerabilities for those products” (2000b, p. 28).

Chapter 7

Recommendations

As we have shown in Chapter 6, “testing is a mess” (Mosser, 2023, priv. comm.)! It will take a lot of time, effort, expertise, and training to organize these terms (and their relations) logically. However, the hardest step is often the first one, so we attempt to give some examples of how this “rationalization” can occur. These changes often arise when we notice an issue with the current state of the terminology and think about what *we* would do to make it better. We do not claim that these are correct, unbiased, or exclusive, just that they can be used as an inspiration for those wanting to pick up where we leave off.

When redefining terms, we seek to make them:

1. Atomic (e.g., disaster/recovery testing seems to have two disjoint definitions)
2. Straightforward (e.g., backup and recovery testing’s definition implies the idea of performance, but its name does not ; failover/recovery testing, failover and recovery testing, and failover testing are all given separately)
3. Consistent (e.g., backup/recovery testing and failover/recovery testing explicitly exclude an aspect included in its parent disaster/recovery testing)

Likewise, we seek to eliminate classes of flaws that can be detected automatically, such as test approaches that are given as synonyms to multiple distinct approaches (Section 6.2.2) or as parents of themselves (Section 6.2.3), or pairs of approaches with both a parent-child *and* synonym relation (Section 6.2.3).

We give recommendations for the areas of operational (acceptance) testing (Section 7.1), recovery testing (Section 7.2), scalability testing (Section 7.3), performance-related testing (Section 7.4), and compatibility testing (Section 7.5). We provide graphical representations (see Section 5.1) of these subsets when helpful in Figures 7.1 to 7.3, in which arrows representing relations between approaches are coloured based on the source tier (see Section 3.1) that defines them. We colour all proposed approaches and relations orange. We also include inferred approaches and relations in grey for completeness, although they are not explicitly given in the literature (see Section 2.3).

Same label to different phantomsections; is that OK?

7.1 Operational (Acceptance) Testing

Since this terminology is not standardized (see Section 6.4), we propose that “Operational Acceptance Testing (OAT)” and “Operational Testing (OT)” are treated as synonyms for a type of acceptance testing (ISO/IEC and IEEE, 2022, p. 22; Hamburg and Mogyorodi, 2024) that focuses on “non-functional” attributes of the system (LambdaTest, 2024). Indeed, this is how we track this approach in our test approach glossary! We define it as “test[ing] to determine the correct installation, configuration and operation of a module and that it operates securely in the operational environment” (ISO/IEC, 2018) or to “evaluate a system or component in its operational environment” (ISO/IEC and IEEE, 2017, p. 303), particularly “to determine if operations and/or systems administration staff can accept [it]” (Hamburg and Mogyorodi, 2024).

find more academic sources

7.2 Recovery Testing

To remedy the flaws we describe in Section 6.5, we recommend that the literature uses these terms more consistently, resulting in the improved graph in Figure 7.1b. The following proposals “recapture” information from the literature more consistently:

1. Prefer the term “recoverability testing” over “recovery testing” to indicate its focus on a system’s *ability* to recover, not its *performance* of recovering (Kam, 2008, p. 47). “Recovery testing” may be an acceptable synonym, since it seems to be more prevalent in the literature.
2. Introduce the term *recovery performance testing* when evaluating performance metrics of a system’s recovery as a subapproach of recoverability testing and performance-related testing¹ (ISO/IEC and IEEE, 2022, Fig. 2; 2013, p. 2).
3. Introduce separate terms for the different methods of recovery which are all subapproaches of recoverability testing:
 - (a) from backup memory (*backup recovery testing*) (ISO/IEC and IEEE, 2021, p. 37; 2013, p. 2),
 - (b) from a back-up system (*failover testing*) (Washizaki, 2024, p. 5-9; Hamburg and Mogyorodi, 2024), or
 - (c) by transferring operations elsewhere (*transfer recovery testing*) (ISO/IEC and IEEE, 2021, p. 37).

¹See Section 7.4.




Figure 7.1: Graphs of relations between terms related to recovery testing.

7.3 Scalability Testing

We describe the issues with scalability testing terminology in Section 6.6 and show the relations between related approaches (as described by the literature) in Figure 7.2a. These flaws are resolved and/or explained by other sources! Taking this extra information into account results in Figure 7.2b and provides a more accurate description of scalability testing.

(CONTRA, SYNS)

ISO/IEC and IEEE (2021, p. 39) define “scalability testing” as the testing of a system’s ability to “perform under conditions that may need to be supported in the future”. This focus on “the future” is supported by Hamburg and Mogyorodi (2024), who define “scalability” as “the degree to which a component or system can be adjusted for changing capacity”; the original source they reference agrees, defining it as “the measure of a system’s ability to be upgraded to accommodate increased loads” (Gerrard and Thompson, 2002, p. 381). In contrast, capacity testing focuses on the system’s present state, evaluating the “capability of a product

to meet requirements for the maximum limits of a product parameter” (ISO/IEC, 2023a). Therefore, these terms should *not* be synonyms, as done by Firesmith (2015, p. 53) and Bas (2024, pp. 22–23).

(CONTRA, DEFS)

There seems to be an underlying reason that sources disagree on whether external modification of the system is part of scalability testing. Washizaki (2024, p. 5-9) claim that one objective of elasticity testing is “to evaluate scalability”, implying that (ISO/IEC, 2023a)’s notion of “scalability” actually refers to “elasticity”! This makes sense in the context of other definitions provided by Washizaki (2024):

- **Scalability:** “the software’s ability to increase and scale up on its nonfunctional requirements, such as load, number of transactions, and volume of data” (p. 5-5). Based on this definition, scalability testing is then a subtype of load testing and volume testing, as well as potentially transaction flow testing.
- **Elasticity Testing²:** testing that “assesses the ability of the SUT ... to rapidly expand or shrink compute, memory, and storage resources without compromising the capacity to meet peak utilization” (p. 5-9). Based on this definition, elasticity testing is then a subtype of memory management testing (with both being a subtype of resource utilization testing) and stress testing.

This distinction is consistent with how the terms are used in industry: Pandey (2023) says that scalability is the ability to “increase ... performance or efficiency as demand increases over time”, while elasticity allows a system to “tackle changes in the workload [that] occur for a short period”.

(WRONG, LABELS)

Washizaki (2024, p. 5-9)’s definition of “scalability testing” is completely separate from the definitions of “scalability”, “capacity”, and “elasticity” discussed above! Therefore, this definition and its related synonym relations should simply be disregarded since they are inconsistent with the rest of the literature.

7.4 Performance(-related) Testing

“Performance testing” is defined as testing “conducted to evaluate the degree to which a test item accomplishes its designated functions” (ISO/IEC and IEEE, 2022, p. 7; 2017, p. 320; similar in 2021, pp. 38-39; Moghadam, 2019, p. 1187). It does this by “measuring the performance metrics” (Moghadam, 2019, p. 1187; similar in Hamburg and Mogyorodi, 2024) (such as the “system’s capacity for growth” (Gerrard, 2000b, p. 23)), “detecting the functional problems appearing under certain execution conditions” (Moghadam, 2019, p. 1187), and “detecting

²This definition seems correct but its source is unclear; see Traceability Flaw 1.

See #35

violations of non-functional requirements under expected and stress conditions” (Moghadam, 2019, p. 1187; similar in Washizaki, 2024, p. 5-9). It is performed either ...

1. “within given constraints of time and other resources” (ISO/IEC and IEEE, 2022, p. 7; 2017, p. 320; similar in Moghadam, 2019, p. 1187), or
2. “under a ‘typical’ load” (ISO/IEC and IEEE, 2021, p. 39).

It is listed as a subset of performance-related testing, which is defined as testing “to determine whether a test item performs as required when it is placed under various types and sizes of ‘load’ ” (2021, p. 38), along with other approaches like load and capacity testing (ISO/IEC and IEEE, 2022, p. 22). Note that “performance, load and stress testing might considerably overlap in many areas” (Moghadam, 2019, p. 1187). In contrast, Washizaki (2024, p. 5-9) gives “capacity and response time” as examples of “performance characteristics” that performance testing would seek to “assess”, which seems to imply that these are subapproaches to performance testing instead. This is consistent with how some sources treat “performance testing” and “performance-related testing” as synonyms (Washizaki, 2024, p. 5-9;




Figure 7.2: Graphs of relations between terms related to scalability testing.

Moghadam, 2019, p. 1187), as noted in Section 6.2.2. This makes sense because of how general the concept of “performance” is; most definitions of “performance testing” seem to treat it as a category of tests.

However, it seems more consistent to infer that the definition of “performance-related testing” is the more general one often assigned to “performance testing” performed “within given constraints of time and other resources” (ISO/IEC and IEEE, 2022, p. 7; 2017, p. 320; similar in Moghadam, 2019, p. 1187), and “performance testing” is a subapproach of this performed “under a ‘typical’ load” (ISO/IEC and IEEE, 2021, p. 39). This has other implications for relations between these types of testing; for example, “load testing” usually occurs “between anticipated conditions of low, typical, and peak usage” (ISO/IEC and IEEE, 2022, p. 5; 2021, p. 39; 2017, p. 253; Hamburg and Mogyorodi, 2024), so it is a child of “performance-related testing” and a parent of “performance testing”.

After these changes, some finishing touches remain. The “self-loops” mentioned in Section 6.2.3 provide no new information and can be removed. Similarly, the term “soak testing” can be removed. Since it is given as a synonym to both “endurance testing” *and* “reliability testing” (see Section 6.2.2), it makes sense to just use these terms instead of one that is potentially ambiguous. These changes (along with those from Sections 7.2 and 7.3 made implicitly) result in the relations shown in Figure 7.3.



Figure 7.3: Proposed relations between rationalized “performance-related testing” terms.

7.5 Compatibility Testing

“Co-existence” and “interoperability” are often defined separately (ISO/IEC and IEEE, 2017, pp. 73, 237; Hamburg and Mogyorodi, 2024), sometimes explicitly as

a decomposition of “compatibility” (ISO/IEC, 2023a)! Following this precedent, “co-existence testing” and “interoperability testing” should be defined as their own test approaches to make their definitions atomic; ISO/IEC and IEEE (2017) define “interoperability testing” (p. 238) but not “co-existence testing”. The term “compatibility testing” may still be a useful test approach to define, but it should be defined independently of its children: “co-existence testing” and “interoperability testing”.

Chapter 8

Development Process

The following is a rough outline of the steps I have gone through this far for this project:

- Start developing system tests (this was pushed for later to focus on unit tests)
- Test inputting default values as `floats` and `ints`
- Check constraints for valid input
- Check constraints for invalid input
- Test the calculations of:
 - `t_flight`
 - `p_land`
 - `d_offset`
 - `s`
- Test the writing of valid output
- Test for projectile going long
- Integrate system tests into existing unit tests
- Test for assumption violation of `g`
 - Code generation could be flawed, so we can't assume assumptions are respected
 - Test cases shouldn't necessarily match what is done by the code; for example, `g = 0` shouldn't really give a `ZeroDivisionError`; it should be a `ValueError`
 - This inspired the potential for The Use of Assertions in Code
- Test that calculations stop on a constraint violation; this is a requirement should be met by the software (see Section 8.3)

- Test behaviour with empty input file
- Start creation of test summary (for `InputParameters` module)
 - It was difficult to judge test case coverage/quality from the code itself
 - This is not really a test plan, as it doesn't capture the testing philosophy
 - Rationale for each test explains why it supports coverage and how Drasil derived (would derive) it
- Start researching testing
- Implement generation of `__init__.py` files (#3516)
- Start the Generating Requirements subproject

8.1 Improvements to Manual Test Code

Even though this code will eventually be generated by Drasil, it is important that it is still human-readable, for the benefit of those reading the code later. This is one of the goals of Drasil (see #3417 for an example of a similar issue). As such, the following improvements were discovered and implement in the manually created testing code:

- use `pytest`'s parameterization
- reuse functions/data for consistency
- improve import structure
- use `conftest` for running code before all tests of a module

8.1.1 Testing with Mocks

When testing code, it is common to first test lower-level modules, then assume that these modules work when testing higher-level modules. An example would be using an input module to set up test cases for a calculation module after testing the input module. This makes sense when writing test cases manually since it reduces the amount of code that needs to be written and still provides a reasonably high assurance in the software; if there is an issue with the input module that affects the calculation module tests, the issue would be revealed when testing the input module.

However, since these test cases will be generated by Drasil, they can be consistently generated with no additional effort. This means that the testing of each module can be done completely independently, increasing the confidence in the tests.

8.2 The Use of Assertions in Code

While assertions are often only used when testing, they can also be used in the code itself to enforce constraints or preconditions; they act like documentation that determines behaviour! For example, they could be used to ensure that assumptions about values (like the value for gravitational acceleration) are respected by the code, which gives a higher degree of confidence in the code. This process is known as “assertion checking” (Lahiri et al., 2013).

investigate OG sources

8.3 Generating Requirements

I structured my manually created test cases around Projectile’s functional requirements, as these are the most objective aspects of the generated code to test automatically. One of these requirements was “Verify-Input-Values”, which said “Check the entered input values to ensure that they do not exceed the data constraints. If any of the input values are out of bounds, an error message is displayed and the calculations stop.” This led me to develop a test case to ensure that if an input constraint was violated, the calculations would stop (Source Code D.1).

However, this test case failed, since the actual implementation of the code did *not* stop upon an input constraint violation. This was because the code choice for what to do on a constraint violation (Source Code D.2) was “disconnected” from the manually written requirement (Source Code D.3), as described in #3523.

Should I include the definition of Constraints?

This problem has been encountered before (#3259) and presented a good opportunity for generation to encourage reusability and consistency. However, since it makes sense to first verify outputs before actually outputting them and inserting generated requirements among manually created ones seemed challenging, it made sense to first generate an output requirement.

While working on Drasil in the summer of 2019, I implemented the generation of an input requirement across most examples (#1844). I had also attempted to generate an output requirement, but due to time constraints, this was not feasible. The main issue with this change was the desire to capture the source of each output for traceability; this source was attached to the `InstanceModel` (or rarely, `DataDefinition`) and not the underlying `Quantity` that was used for a program’s outputs. The way I had attempted to do this was to add the reference as a `Sentence` in a tuple.

Taking another look at this four years later allowed us to see that we should be storing the outputs of a program as their underlying models, allowing us to keep the source information with it. While there is some discussion about how this might change in the future, for now, all outputs of a program should be `InstanceModels`. Since this change required adding the `Referable` constraints to the output field of `SystemInformation`, the outputs of all examples needed to be updated to satisfy this constraint; this meant that generating the output requirement of each example was nearly trivial once the outputs were specified correctly. After modifying `DataDefinitions` in GlassBR that were outputs to be `InstanceModels` (#3569; #3583), reorganizing the requirements of SWHS (#3589; #3607), and clarifying

cite Dr. Smith

add refs to ‘underlying Theory’ comment and ‘not all outputs be IMs’ comment

add constraints

the outputs of SWHS (#3589), SglPend (#3533), DblPend (#3533), GamePhysics (#3609), and SSP (#3630), the output requirement was ready to be generated.

Chapter 9

Research

It was realized early on in the process that it would be beneficial to understand the different kinds of testing (including what they test, what artifacts are needed to perform them, etc.). This section provides some results of this research, as well as some information on why and how it was performed.

A justification for why we decided to do this should be added

9.1 Existing Taxonomies, Ontologies, and the State of Practice

One thing we may want to consider when building a taxonomy/ontology is the semantic difference between related terms. For example, one ontology found that the term “‘IntegrationTest’ is a kind of Context (with semantic of stage, but not a kind of Activity)” while “‘IntegrationTesting’ has semantic of Level-based Testing that is a kind of Testing Activity [or] ... of Test strategy” (Tebes et al., 2019, p. 157).

A note on testing artifacts is that they are “produced and used throughout the testing process” and include test plans, test procedures, test cases, and test results (Souza et al., 2017, p. 3). The role of testing artifacts is not specified in (Barbosa et al., 2006); requirements, drivers, and source code are all treated the same with no distinction (Barbosa et al., 2006, p. 3).

add acronym?

is this punctuation right?

In (Souza et al., 2017), the ontology (ROoST) is made to answer a series of questions, including “What is the test level of a testing activity?” and “What are the artifacts used by a testing activity?” (Souza et al., 2017, pp. 8-9). The question “How do testing artifacts relate to each other?” (Souza et al., 2017, p. 8) is later broken down into multiple questions, such as “What are the test case inputs of a given test case?” and “What are the expected results of a given test case?” (Souza et al., 2017, p. 21). *These questions seem to overlap with the questions we were trying to ask about different testing techniques.* Tebes et al. (2019, pp. 152-153) may provide some sources for software testing terminology and definitions (this seems to include those suggested by Dr. Carette) in addition to a list of ontologies (some of which have been investigated).

One software testing model developed by the Quality Assurance Institute (QAI) includes the test environment (“conditions ...that both enable and constrain how

testing is performed”, including mission, goals, strategy, “management support, resources, work processes, tools, motivation”), test process (testing “standards and procedures”), and tester competency (“skill sets needed to test software in a test environment”) (Perry, 2006, pp. 5-6).

Another source introduced the notion of an “intervention”: “an act performed (e.g. use of a technique¹ or a process change) to adapt testing to a specific context, to solve a test issue, to diagnose testing or to improve testing” (Engström and Petersen, 2015, p. 1) and noted that “academia tend[s] to focus on characteristics of the intervention [while] industrial standards categorize the area from a process perspective” (Engström and Petersen, 2015, p. 2). It provides a structure to “capture both a problem perspective and a solution perspective with respect to software testing” (Engström and Petersen, 2015, pp. 3-4), but this seems to focus more on test interventions and challenges rather than approaches (Engström and Petersen, 2015, Fig. 5).

9.2 Definitions

- Test case: “the specification of all the entities that are essential for the execution, such as input values, execution and timing conditions, testing procedure, and the expected outcomes” (Washizaki, 2024, pp. 5-1 to 5-2)
- Verification: “the process of evaluating a system or component to determine whether the products of a given development phase satisfy the conditions imposed at the start of that phase” (van Vliet, 2000, p. 400)
- Validation: “the process of evaluating a system or component during or at the end of the development process to determine whether it satisfies specified requirements” (van Vliet, 2000, p. 400)
- Test Suite Reduction: the process of reducing the size of a test suite while maintaining the same coverage (Barr et al., 2015, p. 519); can be accomplished through mutation testing
- Test Case Reduction: the process of “removing side-effect free functions” from an individual test case to “reduc[e] test oracle costs” (Barr et al., 2015, p. 519)

9.2.1 Documentation

- Verification and Validation (V&V) Plan: a document for the “planning of test activities” described by IEEE Standard 1012 (van Vliet, 2000, p. 411)
- Test Plan: “a document describing the scope, approach, resources, and schedule of intended test activities” in more detail than the V&V Plan (van Vliet, 2000, pp. 412-413); should also outline entry and exit conditions for the

¹Not formally defined, but distinct from the notion of “test technique” described in Table 2.1.

testing activities as well as any risk sources and levels (Peters and Pedrycz, 2000, p. 445)

- Test Design documentation: “specifies ... the details of the test approach and identifies the associated tests” (van Vliet, 2000, p. 413)
- Test Case documentation: “specifies inputs, predicted outputs and execution conditions for each test item” (van Vliet, 2000, p. 413)
- Test Procedure documentation: “specifies the sequence of actions for the execution of each test” (van Vliet, 2000, p. 413)
- Test Report documentation: “provides information on the results of testing tasks”, addressing software verification and validation reporting (van Vliet, 2000, p. 413)

9.3 General Testing Notes

- The scope of testing is very dependent on what type of software is being tested, as this informs what information/artifacts are available, which approaches are relevant, and which tacit knowledge is present. For example, a method table is a tool for tracking the “test approaches, testing techniques and test types that are required depending ... on the context of the test object” (Hamburg and Mogyorodi, 2024), although this is more specific to the automotive domain

See #54

OG ISO 26262

- If faults exist in programs, they “must be considered faulty, even if we cannot devise test cases that reveal the faults” (van Vliet, 2000, p. 401)
- “There is no established consensus on which techniques ... are the most effective. The only consensus is that the selection will vary as it should be dependent on a number of factors” (ISO/IEC and IEEE, 2021, p. 128; similar in van Vliet, 2000, p. 440), and it is advised to use many techniques when testing (p. 440). This supports the principle of *independence of testing*: the “separation of responsibilities, which encourages the accomplishment of objective testing” (Hamburg and Mogyorodi, 2024)

9.3.1 Test Oracles

A test oracle is a “source of information for determining whether a test has passed or failed” (ISO/IEC and IEEE, 2022, p. 13) or that “the SUT behaved correctly ... and according to the expected outcomes” and can be “human or mechanical” (Washizaki, 2024, p. 5-5). Oracles provide either “a ‘pass’ or ‘fail’ verdict”; otherwise, “the test output is classified as inconclusive” (Washizaki, 2024, p. 5-5). This process can be “deterministic” (returning a Boolean value) or “probabilistic” (returning “a real number in the closed interval $[0, 1]$ ”) (Barr et al., 2015, p. 509). Probabilistic test oracles can be used to reduce the computation cost (since test

oracles are “typically computationally expensive”) (Barr et al., 2015, p. 509) or in “situations where some degree of imprecision can be tolerated” since they “offer a probability that [a given] test case is acceptable” (Barr et al., 2015, p. 510). The SWEBOK Guide V4 lists “unambiguous requirements specifications, behavioral models, and code annotations” as examples (Washizaki, 2024, p. 5-5), and Barr et al. provides four categories (2015, p. 510):

- Specified test oracle: “judge[s] all behavioural aspects of a system with respect to a given formal specification” (Barr et al., 2015, p. 510)
- Derived test oracle: any “artefact[] from which a test oracle may be derived—for instance, a previous version of the system” or “program documentation”; this includes regression testing, metamorphic testing (Barr et al., 2015, p. 510), and invariant detection (either known in advance or “learned from the program”) (Barr et al., 2015, p. 516)
 - This seems to prove “relative correctness” as opposed to “absolute correctness” (Lahiri et al., 2013, p. 345) since this derived oracle may be wrong!
 - “Two versions can be checked for semantic equivalence to ensure the correctness of [a] transformation” in a process that can be done “incrementally” (Lahiri et al., 2013, p. 345)
 - Note that the term “invariant” may be used in different ways (see (Chalin et al., 2006, p. 348))
- Pseudo-oracle: a type of derived test oracle that is “an alternative version of the program produced independently” (by a different team, in a different language, etc.) (Barr et al., 2015, p. 515) . *We could potentially use the programs generated in other languages as pseudo-oracles!*
- Implicit test oracles: detect “‘obvious’ faults such as a program crash” (potentially due to a null pointer, deadlock, memory leak, etc.) (Barr et al., 2015, p. 510)
- “Lack of an automated test oracle”: for example; a human oracle generating sample data that is “realistic” and “valid”, (Barr et al., 2015, pp. 510-511), crowdsourcing (Barr et al., 2015, p. 520), or a “Wideband Delphi”: “an expert-based test estimation technique that ... uses the collective wisdom of the team members” (Hamburg and Mogyorodi, 2024)

see ISO 29119-11

9.3.2 Generating Test Cases

- “Impl[ies] a reduction in human effort and cost, with the potential to impact the test coverage positively”, and a given “policy could be reused in analogous situations which leads to even more efficiency in terms of required efforts” (Moghadam, 2019, p. 1187)

- A selected “test adequacy criterion can be used in the test selection process [as a ‘test case generator’]. If a 100% statement coverage has not been achieved yet, an additional test case is selected that covers one or more statements yet untested. This generative view is used in many test tools” (van Vliet, 2000, p. 402). This may be similar to the “test data generators” mentioned on (p. 410)

Investigate

- “Dynamic test generation consists of running a program while simultaneously executing the program symbolically in order to gather constraints on inputs from conditional statements encountered along the execution (Godefroid and Luchaup, 2011, p. 23)

OG [11, 6]

- “Generating tests to detect [loop inefficiencies]” is difficult due to “virtual call resolution”, reachability conditions, and order-sensitivity (Dhok and Ramanathan, 2016, p. 896)
- Can be facilitated by “testing frameworks such as JUnit [that] automate the testing process by writing test code” (Sakamoto et al., 2013, p. 344)
- Assertion checking requires “auxiliary invariants”, and while “many ... can be synthesized automatically by invariant generation methods, the undecidable nature (or the high practical complexity) of assertion checking precludes complete automation for a general class of user-supplied assertions” (Lahiri et al., 2013, p. 345)
 - Differential Assertion Checking (DAC) can be supported by “automatic invariant generation” (Lahiri et al., 2013, p. 345)

OG Halfond and Orso, 2007

- *Automated interface discovery* can be used “for test-case generation for web applications” (Doğan et al., 2014, p. 184)

OG Artzi et al., 2008

- “Concrete and symbolic execution” can be used in “a dynamic test generation technique ... for PHP applications” (Doğan et al., 2014, p. 192)
- COBRA is a tool that “generates test cases automatically and applies them to the simulated industrial control system in a SiL Test” (Preuße et al., 2012, p. 2)
- Test case generation is useful for instances where one kind of testing is difficult, but can be generated from a different, simpler kind (e.g., asynchronous testing from synchronous testing (Jard et al., 1999))
- Since some values may not always be applicable to a given scenario (e.g., a test case for zero doesn’t make sense if there is a constraint that the value in question cannot be zero), the user should likely be able to select categories of tests to generate instead of Drasil just generating all possible test cases based on the inputs (Smith and Carette, 2023)

Investigate!

- “Test suite augmentation techniques specialise in identifying and generating” new tests based on changes “that add new features”, but they could be extended to also augment “the expected output” and “the existing *oracles*” (Barr et al., 2015, p. 516)
- The Fault Injection Security Tool (FIST) “automates fault injection analysis of software using program inputs, fault injection functions, and assertions in programs written in C and C++” (Ghosh and Voas, 1999, pp. 40–41). “For example, Booleans are corrupted to their opposite value during execution, integers are corrupted using a random function with a uniform distribution centered around their current value, character strings are corrupted using random values” (p. 41). Unfortunately, “we do not have automatic learning systems for protecting software states, though it is the subject of ongoing research” (p. 44).

9.4 Examples of Metamorphic Relations (MRs)

- The distance between two points should be the same regardless of which one is the “start” point (ISO/IEC and IEEE, 2021, p. 22)
- “If a person smokes more cigarettes, then their expected age of death will probably decrease (and not increase)” (ISO/IEC and IEEE, 2021, p. 22)
- “For a function that translates speech into text[,] ... the same speech at different input volume levels ... [should result in] the same text” (ISO/IEC and IEEE, 2021, p. 22)
- The average of a list of numbers should be equal (within floating-point errors) regardless of the list’s order (Kanewala and Yueh Chen, 2019, p. 67)
- For matrices, if $B = B_1 + B_2$, then $A \times B = A \times B_1 + A \times B_2$ (Kanewala and Yueh Chen, 2019, pp. 68-69)
- Symmetry of trigonometric functions; for example, $\sin(x) = \sin(-x)$ and $\sin(x) = \sin(x + 360^\circ)$ (Kanewala and Yueh Chen, 2019, p. 70)
- Modifying input parameters to observe expected changes to a model’s output (e.g., testing epidemiological models calibrated with “data from the 1918 Influenza outbreak”); by “making changes to various model parameters ... authors identified an error in the output method of the agent based epidemiological model” (Kanewala and Yueh Chen, 2019, p. 70)
- Using machine learning to predict likely MRs to identify faults in mutated versions of a program (about 90% in this case) (Kanewala and Yueh Chen, 2019, p. 71)

9.5 Roadblocks to Testing

- Intractability: it is generally impossible to test a program exhaustively (ISO/IEC and IEEE, 2022, p. 4; Washizaki, 2024, p. 5-5; Peters and Pedrycz, 2000, pp. 439, 461; van Vliet, 2000, p. 421)
- Adequacy: to counter the issue of intractability, it is desirable “to reduce the cardinality of the test suites while keeping the same effectiveness in terms of coverage or fault detection rate” (Washizaki, 2024, p. 5-4) which is difficult to do objectively; see also “minimization”, the process of “removing redundant test cases” (Washizaki, 2024, p. 5-4)
- Undecidability (Peters and Pedrycz, 2000, p. 439): it is impossible to know certain properties about a program, such as if it will halt (i.e., the Halting Problem (Gurfinkel, 2017, p. 4)), so “automatic testing can’t be guaranteed to always work” for all properties (Nelson, 1999)

Add paragraph/section number?

9.5.1 Roadblocks to Testing Scientific Software (Kanewala and Yueh Chen, 2019, p. 67)

- “Correct answers are often unknown”: if the results were already known, there would be no need to develop software to model them (Kanewala and Yueh Chen, 2019, p. 67); in other words, complete test oracles don’t exist “in all but the most trivial cases” (Barr et al., 2015, p. 510), and even if they are, the “automation of mechanized oracles can be difficult and expensive” (Washizaki, 2024, p. 5.5)
- “Practically difficult to validate the computed output”: complex calculations and outputs are difficult to verify (Kanewala and Yueh Chen, 2019, p. 67)
- “Inherent uncertainties”: since scientific software models scenarios that occur in a chaotic and imperfect world, not every factor can be accounted for (Kanewala and Yueh Chen, 2019, p. 67)
- “Choosing suitable tolerances”: difficult to decide what tolerance(s) to use when dealing with floating-point numbers (Kanewala and Yueh Chen, 2019, p. 67)
- “Incompatible testing tools”: while scientific software is often written in languages like FORTRAN, testing tools are often written in languages like Java or C++ (Kanewala and Yueh Chen, 2019, p. 67)

Out of this list, only the first two apply. The scenarios modelled by Drasil are idealized and ignore uncertainties like air resistance, wind direction, and gravitational fluctuations. There are not any instances where special consideration for floating-point arithmetic must be taken; the default tolerance used for relevant testing frameworks has been used and is likely sufficient for future testing. On a related note, the scientific software we are trying to test is already generated in languages with widely-used testing frameworks.

Add example

Add source(s)?

Chapter 10

Extras

Writing Directives

- What macros do I want the reader to know about?

10.1 Writing Directives

I enjoy writing directives (mostly questions) to navigate what I should be writing about in each chapter. You can do this using:

Source Code 10.1: Pseudocode: exWD

```
\begin{writingdirectives}
  \item What macros do I want the reader to know about?
\end{writingdirectives}
```

Personally, I put them at the top of chapter files, just after chapter declarations.

10.2 HREFs

For PDFs, we have (at least) 2 ways of viewing them: on our computers, and printed out on paper. If you choose to view through your computer, reading links (as they are linked in this example, inlined everywhere with “clickable” links) is fine. However, if you choose to read it on printed paper, you will find trouble clicking on those same links. To mitigate this issue, I built the “porthref” macro (see `macros.tex` for the definition) to build links that appear as clickable text when “compiling for computer-focused reading,” and adds links to footnotes when “compiling for printing-focused reading.” There is an option (`compilingforprinting`) in the `manifest.tex` file that controls whether PDF builds should be done for

computers or for printers. For example, by default, McMaster is made with clickable functionality, but if you change the `manifest.tex` option as mentioned, then you will see the link in a footnote (try it out!).

Source Code 10.2: Pseudocode: `exPHref`

```
\porthref{McMaster}{https://www.mcmaster.ca/}
```

10.3 Pseudocode Code Snippets

For pseudocode, you can also use the pseudocode environment, such as that used in Source Code D.5.

10.4 TODOs

While writing, I plastered my thesis with notes for future work because, for whatever reason, I just didn't want to, or wasn't able to, do said work at that time. To help me sort out my notes, I used the `todonotes` package with a few extra macros (defined in `macros.tex`). For example,...

Important notes:

Important: "Important" notes.

Generic inlined notes:

Generic inlined notes.

Notes for later:

Some "easy" notes:

Easy: Easier notes.

Tedious work:

Needs time: Tedious notes.

Questions:

Later: TODO notes for later! For finishing touches, etc.

Q #21: Questions I might have?

Bibliography

- Mominul Ahsan, Stoyan Stoyanov, Chris Bailey, and Alhussein Albarbar. Developing Computational Intelligence for Smart Qualification Testing of Electronic Products. *IEEE Access*, 8:16922–16933, January 2020. ISSN 2169-3536. doi: 10.1109/ACCESS.2020.2967858. URL <https://www.webofscience.com/api/gateway?GWVersion=2&SrcAuth=DynamicDOIArticle&SrcApp=WOS&KeyAID=10.1109%2FACCESS.2020.2967858&DestApp=DOI&SrcAppSID=USW2EC0CB9ABcVz5BcZ70BCfllmtJ&SrcJTitle=IEEE+ACCESS&DestDOIRegistrantName=Institute+of+Electrical+and+Electronics+Engineers>. Place: Piscataway.
- Paul Ammann and Jeff Offutt. *Introduction to Software Testing*. Cambridge University Press, Cambridge, United Kingdom, 2nd edition, 2017. ISBN 978-1-107-17201-2. URL <https://eopcw.com/find/downloadFiles/11>.
- Mohammad Bajammal and Ali Mesbah. Web Canvas Testing Through Visual Inference. In *2018 IEEE 11th International Conference on Software Testing, Verification and Validation (ICST)*, pages 193–203, Västerås, Sweden, 2018. IEEE. ISBN 978-1-5386-5012-7. doi: 10.1109/ICST.2018.00028. URL <https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=8367048>.
- Ellen Francine Barbosa, Elisa Yumi Nakagawa, and José Carlos Maldonado. Towards the Establishment of an Ontology of Software Testing. volume 6, pages 522–525, San Francisco, CA, USA, January 2006.
- Luciano Baresi and Mauro Pezzè. An Introduction to Software Testing. *Electronic Notes in Theoretical Computer Science*, 148(1):89–111, February 2006. ISSN 1571-0661. doi: 10.1016/j.entcs.2005.12.014. URL <https://www.sciencedirect.com/science/article/pii/S1571066106000442>.
- Earl T. Barr, Mark Harman, Phil McMinn, Muzammil Shahbaz, and Shin Yoo. The Oracle Problem in Software Testing: A Survey. *IEEE Transactions on Software Engineering*, 41(5):507–525, 2015. doi: 10.1109/TSE.2014.2372785.
- Mykola Bas. *Data Backup and Archiving*. Bachelor thesis, Czech University of Life Sciences Prague, Praha-Suchdol, Czechia, March 2024. URL https://thes.es.cz/id/60licg/zaverecna_prace_Archive.pdf.
- Josh Berdine, Cristiano Calcagno, and Peter W. O’Hearn. Smallfoot: Modular Automatic Assertion Checking with Separation Logic. In Frank S. de Boer,

- Marcello M. Bonsangue, Susanne Graf, and Willem-Paul de Roever, editors, *Formal Methods for Components and Objects*, pages 115–137, Berlin, Heidelberg, 2006. Springer. ISBN 978-3-540-36750-5. doi: 10.1007/11804192_6.
- Michael Bluejay. Slot Machine PAR Sheets, May 2024. URL <https://easy.vegas/games/slots/par-sheets>.
- Chris Bocchino and William Hamilton. Eastern Range Titan IV/Centaur-TDRSS Operational Compatibility Testing. In *International Telemetering Conference Proceedings*, San Diego, CA, USA, October 1996. International Foundation for Telemetering. ISBN 978-0-608-04247-3. URL https://repository.arizona.edu/bitstream/handle/10150/607608/ITC_1996_96-01-4.pdf?sequence=1&isAllowed=y.
- Pierre Bourque and Richard E. Fairley, editors. *Guide to the Software Engineering Body of Knowledge, Version 3.0*. IEEE Computer Society Press, Washington, DC, USA, 2014. ISBN 0-7695-5166-1. URL www.swebok.org.
- Jacques Carette, July 2024. URL <https://github.com/samm82/TestingTesting/issues/64#issuecomment-2217320875>.
- Jacques Carette, Spencer Smith, Jason Balaci, Ting-Yu Wu, Samuel Crawford, Dong Chen, Dan Szymczak, Brooks MacLachlan, Dan Scime, and Maryyam Niazi. Drasil, February 2021. URL <https://github.com/JacquesCarette/Drasil/tree/v0.1-alpha>.
- Patrice Chalin, Joseph R. Kiniry, Gary T. Leavens, and Erik Poll. Beyond Assertions: Advanced Specification and Verification with JML and ESC/Java2. In Frank S. de Boer, Marcello M. Bonsangue, Susanne Graf, and Willem-Paul de Roever, editors, *Formal Methods for Components and Objects*, pages 342–363, Berlin, Heidelberg, 2006. Springer. ISBN 978-3-540-36750-5. doi: 10.1007/11804192_16.
- ChatGPT (GPT-4o). Defect Clustering Testing, November 2024. URL <https://chatgpt.com/share/67463dd1-d0a8-8012-937b-4a3db0824dcf>.
- Shauvik Roy Choudhary, Husayn Versee, and Alessandro Orso. A Cross-browser Web Application Testing Tool. In *2010 IEEE International Conference on Software Maintenance*, pages 1–6, Timisoara, Romania, September 2010. IEEE. ISBN 978-1-4244-8629-8. doi: 10.1109/ICSM.2010.5609728. URL <https://ieeexplore.ieee.org/abstract/document/5609728>. ISSN: 1063-6773.
- Alan Dennis, Barbara Haley Wixom, and Roberta M. Roth. *System Analysis and Design*. John Wiley & Sons, 5th edition, 2012. ISBN 978-1-118-05762-9. URL https://www.uoitc.edu.iq/images/documents/informatics-institute/Competitive_exam/Systemanalysisanddesign.pdf.
- Monika Dhok and Murali Krishna Ramanathan. Directed Test Generation to Detect Loop Inefficiencies. In *Proceedings of the 2016 24th ACM SIGSOFT*

- International Symposium on Foundations of Software Engineering*, FSE 2016, pages 895–907, New York, NY, USA, November 2016. Association for Computing Machinery. ISBN 978-1-4503-4218-6. doi: 10.1145/2950290.2950360. URL <https://dl.acm.org/doi/10.1145/2950290.2950360>.
- M. Dominguez-Pumar, J. M. Olm, L. Kowalski, and V. Jimenez. Open loop testing for optimizing the closed loop operation of chemical systems. *Computers & Chemical Engineering*, 135:106737, 2020. ISSN 0098-1354. doi: <https://doi.org/10.1016/j.compchemeng.2020.106737>. URL <https://www.sciencedirect.com/science/article/pii/S0098135419312736>.
- Serdar Doğan, Aysu Betin-Can, and Vahid Garousi. Web application testing: A systematic literature review. *Journal of Systems and Software*, 91:174–201, 2014. ISSN 0164-1212. doi: <https://doi.org/10.1016/j.jss.2014.01.010>. URL <https://www.sciencedirect.com/science/article/pii/S0164121214000223>.
- Emelie Engström and Kai Petersen. Mapping software testing practice with software testing research — serp-test taxonomy. In *2015 IEEE Eighth International Conference on Software Testing, Verification and Validation Workshops (ICSTW)*, pages 1–4, 2015. doi: 10.1109/ICSTW.2015.7107470.
- Norman E. Fenton and Shari Lawrence Pfleeger. *Software Metrics: A Rigorous & Practical Approach*. PWS Publishing Company, Boston, MA, USA, 2nd edition, 1997. ISBN 0-534-95425-1.
- Donald G. Firesmith. A Taxonomy of Testing Types, 2015. URL <https://apps.dtic.mil/sti/pdfs/AD1147163.pdf>.
- P. Forsyth, T. Maguire, and R. Kuffel. Real Time Digital Simulation for Control and Protection System Testing. In *2004 IEEE 35th Annual Power Electronics Specialists Conference (IEEE Cat. No.04CH37551)*, volume 1, pages 329–335, Aachen, Germany, 2004. IEEE. ISBN 0-7803-8399-0. doi: 10.1109/PESC.2004.1355765.
- Paul Gerrard. Risk-based E-business Testing - Part 1: Risks and Test Strategy. Technical report, Systeme Evolutif, London, UK, 2000a. URL https://www.agileconnection.com/sites/default/files/article/file/2013/XUS129342file1_0.pdf.
- Paul Gerrard. Risk-based E-business Testing - Part 2: Test Techniques and Tools. Technical report, Systeme Evolutif, London, UK, 2000b. URL wenku.uml.com.cn/document/test/EBTestingPart2.pdf.
- Paul Gerrard and Neil Thompson. *Risk-based E-business Testing*. Artech House computing library. Artech House, Norwood, MA, USA, 2002. ISBN 978-1-58053-570-0. URL <https://books.google.ca/books?id=54UKereAdJ4C>.
- Anup K Ghosh and Jeffrey M Voas. Inoculating software for survivability. *Communications of the ACM*, 42(7):38–44, July 1999. URL <https://dl.acm.org/doi/pdf/10.1145/306549.306563>. Publisher: ACM New York, NY, USA.

- Patrice Godefroid and Daniel Luchaup. Automatic Partial Loop Summarization in Dynamic Test Generation. In *Proceedings of the 2011 International Symposium on Software Testing and Analysis*, ISSTA '11, pages 23–33, New York, NY, USA, July 2011. Association for Computing Machinery. ISBN 978-1-4503-0562-4. doi: 10.1145/2001420.2001424. URL <https://dl.acm.org/doi/10.1145/2001420.2001424>.
- W. Goralski. xDSL loop qualification and testing. *IEEE Communications Magazine*, 37(5):79–83, 1999. doi: 10.1109/35.762860.
- Arie Gurfinkel. Testing: Coverage and Structural Coverage, 2017. URL <https://ece.uwaterloo.ca/~agurfink/ece653w17/assets/pdf/W03-Coverage.pdf>.
- Matthias Hamburg and Gary Mogyorodi, editors. ISTQB Glossary, v4.3, 2024. URL https://glossary.istqb.org/en_US/search.
- Daniel C Holley, Gary D Mele, and Sujata Naidu. NASA Rat Acoustic Tolerance Test 1994-1995: 8 kHz, 16 kHz, 32 kHz Experiments. Technical Report NASA-CR-202117, San Jose State University, San Jose, CA, USA, January 1996. URL <https://ntrs.nasa.gov/api/citations/19960047530/downloads/19960047530.pdf>.
- R. Brian Howe and Robert Johnson. Research Protocol for the Evaluation of Medical Waiver Requirements for the Use of Lisinopril in USAF Aircrew. Interim Technical Report AL/AO-TR-1995-0116, Air Force Materiel Command, Brooks Air Force Base, TX, USA, November 1995. URL <https://apps.dtic.mil/sti/tr/pdf/ADA303379.pdf>.
- Anthony Hunt, Peter Michalski, Dong Chen, Jason Balaci, and Spencer Smith. Drasil - Generate All the Things!, 2021. URL <https://jacquescarette.github.io/Drasil/>.
- IEEE. IEEE Standard for System and Software Verification and Validation. *IEEE Std 1012-2012 (Revision of IEEE Std 1012-2004)*, 2012. doi: 10.1109/IEEESTD.2012.6204026.
- Adisak Intana, Monchanok Thongthep, Phatcharee Thepnimit, Phaplak Saethapan, and Tanawat Monpipat. SYNTTest: Prototype of Syntax Test Case Generation Tool. In *5th International Conference on Information Technology (InCIT)*, pages 259–264. IEEE, 2020. ISBN 978-1-72819-321-2. doi: 10.1109/InCIT50588.2020.9310968.
- ISO. ISO 13849-1:2015 - Safety of machinery –Safety-related parts of control systems –Part 1: General principles for design. *ISO 13849-1:2015*, December 2015. URL <https://www.iso.org/obp/ui#iso:std:iso:13849:-1:ed-3:v1:en>.
- ISO. ISO 21384-2:2021 - Unmanned aircraft systems –Part 2: UAS components. *ISO 21384-2:2021*, December 2021. URL <https://www.iso.org/obp/ui#iso:std:iso:21384:-2:ed-1:v1:en>.

ISO. ISO 28881:2022 - Machine tools –Safety –Electrical discharge machines. *ISO 28881:2022*, April 2022. URL <https://www.iso.org/obp/ui/#iso:std:iso:28881:ed-2:v1:en>.

ISO/IEC. ISO/IEC 25010:2011 - Systems and software engineering –Systems and software Quality Requirements and Evaluation (SQuaRE) –System and software quality models. *ISO/IEC 25010:2011*, March 2011. URL <https://www.iso.org/obp/ui/#iso:std:iso-iec:25010:ed-1:v1:en>.

ISO/IEC. ISO/IEC 2382:2015 - Information technology –Vocabulary. *ISO/IEC 2382:2015*, May 2015. URL <https://www.iso.org/obp/ui/en/#iso:std:iso-iec:2382:ed-1:v2:en>.

ISO/IEC. ISO/IEC TS 20540:2018 - Information technology – Security techniques –Testing cryptographic modules in their operational environment. *ISO/IEC TS 20540:2018*, May 2018. URL <https://www.iso.org/obp/ui/#iso:std:iso-iec:ts:20540:ed-1:v1:en>.

ISO/IEC. ISO/IEC 25010:2023 - Systems and software engineering –Systems and software Quality Requirements and Evaluation (SQuaRE) –Product quality model. *ISO/IEC 25010:2023*, November 2023a. URL <https://www.iso.org/obp/ui/#iso:std:iso-iec:25010:ed-2:v1:en>.

ISO/IEC. ISO/IEC 25019:2023 - Systems and software engineering –Systems and software Quality Requirements and Evaluation (SQuaRE) –Quality-in-use model. *ISO/IEC 25019:2023*, November 2023b. URL <https://www.iso.org/obp/ui/en/#iso:std:iso-iec:25019:ed-1:v1:en>.

ISO/IEC and IEEE. ISO/IEC/IEEE International Standard - Systems and software engineering–Vocabulary. *ISO/IEC/IEEE 24765:2010(E)*, December 2010. doi: 10.1109/IEEESTD.2010.5733835.

ISO/IEC and IEEE. ISO/IEC/IEEE International Standard - Systems and software engineering –Software testing –Part 1: General concepts. *ISO/IEC/IEEE 29119-1:2013*, September 2013. doi: 10.1109/IEEESTD.2013.6588537.

ISO/IEC and IEEE. ISO/IEC/IEEE International Standard - Software and systems engineering –Software testing –Part 5: Keyword-Driven Testing. *ISO/IEC/IEEE 29119-5:2016*, November 2016. doi: 10.1109/IEEESTD.2016.7750539.

ISO/IEC and IEEE. ISO/IEC/IEEE International Standard - Systems and software engineering–Vocabulary. *ISO/IEC/IEEE 24765:2017(E)*, September 2017. doi: 10.1109/IEEESTD.2017.8016712.

ISO/IEC and IEEE. ISO/IEC/IEEE International Standard - Systems and software engineering –Systems and software assurance –Part 1: Concepts and vocabulary. *ISO/IEC/IEEE 15026-1:2019*, March 2019a. doi: 10.1109/IEEESTD.2019.8657410.

ISO/IEC and IEEE. ISO/IEC/IEEE International Standard - Systems and software engineering –Guidelines for the utilization of ISO/IEC/IEEE 15288 in the context of system of systems (SoS). *ISO/IEC/IEEE 21840:2019*, December 2019b. doi: 10.1109/IEEESTD.2019.8929110.

ISO/IEC and IEEE. ISO/IEC/IEEE International Standard - Software and systems engineering –Software testing –Part 4: Test techniques. *ISO/IEC/IEEE 29119-4:2021(E)*, October 2021. doi: 10.1109/IEEESTD.2021.9591574.

ISO/IEC and IEEE. ISO/IEC/IEEE International Standard - Systems and software engineering –Software testing –Part 1: General concepts. *ISO/IEC/IEEE 29119-1:2022(E)*, January 2022. doi: 10.1109/IEEESTD.2022.9698145.

Claude Jard, Thierry Jérón, Lénaïck Tanguy, and César Viho. Remote testing can be as powerful as local testing. In Jianping Wu, Samuel T. Chanson, and Qiang Gao, editors, *Formal Methods for Protocol Engineering and Distributed Systems: Forte XII / PSTV XIX'99*, volume 28 of *IFIP Advances in Information and Communication Technology*, pages 25–40, Beijing, China, October 1999. Springer. ISBN 978-0-387-35578-8. doi: 10.1007/978-0-387-35578-8_2. URL https://doi.org/10.1007/978-0-387-35578-8_2.

Timothy P. Johnson. Snowball Sampling: Introduction. In *Wiley StatsRef: Statistics Reference Online*. John Wiley & Sons, Ltd, 2014. ISBN 978-1-118-44511-2. doi: <https://doi.org/10.1002/9781118445112.stat05720>. URL <https://onlinelibrary.wiley.com/doi/abs/10.1002/9781118445112.stat05720>. eprint: <https://onlinelibrary.wiley.com/doi/pdf/10.1002/9781118445112.stat05720>.

Ben Kam. Web Applications Testing. Technical Report 2008-550, Queen’s University, Kingston, ON, Canada, October 2008. URL <https://research.cs.queensu.ca/TechReports/Reports/2008-550.pdf>.

Cem Kaner, James Bach, and Bret Pettichord. *Lessons Learned in Software Testing: A Context-Driven Approach*. John Wiley & Sons, December 2011. ISBN 978-0-471-08112-8. URL <https://www.wiley.com/en-ca/Lessons+Learned+in+Software+Testing%3A+A+Context-Driven+Approach-p-9780471081128>.

Upulee Kanewala and Tsong Yueh Chen. Metamorphic testing: A simple yet effective approach for testing scientific software. *Computing in Science & Engineering*, 21(1):66–72, 2019. doi: 10.1109/MCSE.2018.2875368.

Knüvener Mackert GmbH. *Knüvener Mackert SPICE Guide*. Knüvener Mackert GmbH, Reutlingen, Germany, 7th edition, 2022. ISBN 978-3-00-061926-7. URL <https://knuevenermackert.com/wp-content/uploads/2021/06/SPICE-BOOKLET-2022-05.pdf>.

Ivans Kuļešovs, Vineta Arnīcane, Guntis Arnīcans, and Juris Borzovs. Inventory of Testing Ideas and Structuring of Testing Terms. 1:210–227, January 2013.

- Shuvendu K. Lahiri, Kenneth L. McMillan, Rahul Sharma, and Chris Hawblitzel. Differential Assertion Checking. In *Proceedings of the 2013 9th Joint Meeting on Foundations of Software Engineering*, ESEC/FSE 2013, pages 345–355, New York, NY, USA, August 2013. Association for Computing Machinery. ISBN 978-1-4503-2237-9. doi: 10.1145/2491411.2491452. URL <https://dl.acm.org/doi/10.1145/2491411.2491452>.
- LambdaTest. What is Operational Testing: Quick Guide With Examples, 2024. URL <https://www.lambdatest.com/learning-hub/operational-testing>.
- Danye Liu, Shaonan Tian, Yu Zhang, Chaoquan Hu, Hui Liu, Dong Chen, Lin Xu, and Jun Yang. Ultrafine SnPd nanoalloys promise high-efficiency electrocatalysis for ethanol oxidation and oxygen reduction. *ACS Applied Energy Materials*, 6(3):1459–1466, January 2023. doi: <https://doi.org/10.1021/acsaem.2c03355>. URL https://pubs.acs.org/doi/pdf/10.1021/acsaem.2c03355?casa_token=ItHfKxeQNbsAAAAA:8zEdU5hi2HfHsSony3ku-lbH902jkHpA-JZw8jIeODzUvFtSdQRdbYhmVq47aX22igR52o2S22mnC88Mxw. Publisher: ACS Publications.
- Robert Mandl. Orthogonal Latin squares: an application of experiment design to compiler testing. *Communications of the ACM*, 28(10):1054–1058, October 1985. ISSN 0001-0782. doi: 10.1145/4372.4375. URL <https://doi.org/10.1145/4372.4375>.
- Robert M. McClure. Introduction, July 2001. URL <http://homepages.cs.ncl.ac.uk/brian.randell/NATO/Introduction.html>.
- Mahshid Helali Moghadam. Machine Learning-Assisted Performance Testing. In *Proceedings of the 2019 27th ACM Joint Meeting on European Software Engineering Conference and Symposium on the Foundations of Software Engineering*, ESEC/FSE 2019, pages 1187–1189, New York, NY, USA, 2019. Association for Computing Machinery. ISBN 978-1-4503-5572-8. doi: 10.1145/3338906.3342484. URL <https://doi.org/10.1145/3338906.3342484>.
- V. V. Morgun, L. I. Voronin, R. R. Kaspransky, S. L. Pool, M. R. Barratt, and O. L. Novinkov. The Russian-US Experience with Development Joint Medical Support Procedures for Before and After Long-Duration Space Flights. Technical report, NASA, Houston, TX, USA, 1999. URL <https://ntrs.nasa.gov/api/citations/2000085877/downloads/2000085877.pdf>.
- E. E. Mukhin, V. M. Nelyubov, V. A. Yukish, E. P. Smirnova, V. A. Solovei, N. K. Kalinina, V. G. Nagaitsev, M. F. Valishin, A. R. Belozeroval, S. A. Enin, A. A. Borisov, N. A. Deryabina, V. I. Khripunov, D. V. Portnov, N. A. Babinov, D. V. Dokhtarenko, I. A. Khodunov, V. N. Klimov, A. G. Razdobarin, S. E. Alexandrov, D. I. Elets, A. N. Bazhenov, I. M. Bukreev, An P. Chernakov, A. M. Dmitriev, Y. G. Ibragimova, A. N. Koval, G. S. Kurskiev, A. E. Litvinov, K. O. Nikolaenko, D. S. Samsonov, V. A. Senichenkov, R. S. Smirnov, S. Yu Tolstyakov, I. B. Tereschenko, L. A. Varshavchik, N. S. Zhiltsov, A. N. Mokeev, P. V. Chernakov, P. Andrew, and M. Kempenaars. Radiation tolerance testing

of piezoelectric motors for ITER (first results). *Fusion Engineering and Design*, 176(article 113017), 2022. ISSN 0920-3796. doi: <https://doi.org/10.1016/j.fuseengdes.2022.113017>. URL <https://www.sciencedirect.com/science/article/pii/S0920379622000175>.

Peter Naur and Brian, editors Randell. Software Engineering: Report on a conference sponsored by the NATO Science Committee, Garmisch, Germany, 7th to 11th October 1968. Brussels, Belgium, January 1969. Scientific Affairs Division, NATO. URL <http://homepages.cs.ncl.ac.uk/brian.randell/NATO/nato1968.PDF>.

Randal C. Nelson. Formal Computational Models and Computability, January 1999. URL https://www.cs.rochester.edu/u/nelson/courses/csc_173/computability/undecidable.html.

Pranav Pandey. Scalability vs Elasticity, February 2023. URL <https://www.linkedin.com/pulse/scalability-vs-elasticity-pranav-pandey/>.

Bhupesh A. Parate, K.D. Deodhar, and V.K. Dixit. Qualification Testing, Evaluation and Test Methods of Gas Generator for IEDs Applications. *Defence Science Journal*, 71(4):462–469, July 2021. doi: 10.14429/dsj.71.16601. URL <https://publications.drdo.gov.in/ojs/index.php/dsj/article/view/16601>.

Ron Patton. *Software Testing*. Sams Publishing, Indianapolis, IN, USA, 2nd edition, 2006. ISBN 0-672-32798-8.

William E. Perry. *Effective Methods for Software Testing*. Wiley Publishing, Inc., Indianapolis, IN, USA, 3rd edition, 2006. ISBN 978-0-7645-9837-1.

J.F. Peters and W. Pedrycz. *Software Engineering: An Engineering Approach*. Worldwide series in computer science. John Wiley & Sons, Ltd., 2000. ISBN 978-0-471-18964-0.

Brian J. Pierre, Felipe Wilches-Bernal, David A. Schoenwald, Ryan T. Elliott, Jason C. Neely, Raymond H. Byrne, and Daniel J. Trudnowski. Open-loop testing results for the pacific DC intertie wide area damping controller. In *2017 IEEE Manchester PowerTech*, pages 1–6, 2017. doi: 10.1109/PTC.2017.7980834.

Sebastian Preuße, Hans-Christian Lapp, and Hans-Michael Hanisch. Closed-loop System Modeling, Validation, and Verification. In *Proceedings of 2012 IEEE 17th International Conference on Emerging Technologies & Factory Automation (ETFA 2012)*, pages 1–8, Krakow, Poland, 2012. IEEE. ISBN 978-1-4673-4736-5. doi: 10.1109/ETFA.2012.6489679. URL <https://ieeexplore.ieee.org/abstract/document/6489679>.

Vasile Rus, Sameer Mohammed, and Sajjan G Shiva. Automatic Clustering of Defect Reports. In *Proceedings of the Twentieth International Conference on Software Engineering & Knowledge Engineering (SEKE 2008)*, pages 291–296,

San Francisco, CA, USA, July 2008. Knowledge Systems Institute Graduate School. ISBN 1-891706-22-5. URL <https://core.ac.uk/download/pdf/48606872.pdf>.

Kazunori Sakamoto, Kaizu Tomohiro, Daigo Hamura, Hironori Washizaki, and Yoshiaki Fukazawa. POGen: A Test Code Generator Based on Template Variable Coverage in Gray-Box Integration Testing for Web Applications. In Vittorio Cortellessa and Dániel Varró, editors, *Fundamental Approaches to Software Engineering*, pages 343–358, Berlin, Heidelberg, March 2013. Springer Berlin Heidelberg. ISBN 978-3-642-37057-1. URL https://link.springer.com/chapter/10.1007/978-3-642-37057-1_25.

Raghvinder S. Sangwan and Phillip A. LaPlante. Test-Driven Development in Large Projects. *IT Professional*, 8(5):25–29, October 2006. ISSN 1941-045X. doi: 10.1109/MITP.2006.122. URL <https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=1717338>.

Sheetal Sharma, Kartika Panwar, and Rakesh Garg. Decision Making Approach for Ranking of Software Testing Techniques Using Euclidean Distance Based Approach. *International Journal of Advanced Research in Engineering and Technology*, 12(2):599–608, February 2021. ISSN 0976-6499. doi: 10.34218/IJARET.12.2.2021.059. URL <https://iaeme.com/Home/issue/IJARET?Volume=12&Issue=2>.

Spencer Smith. Potential Projects, June 2024. URL <https://github.com/JacquesCarette/Drasil/wiki/Potential-Projects>.

Spencer Smith and Jacques Carette. Private Communication, July 2023.

Harry Sneed and Siegfried Göschl. A Case Study of Testing a Distributed Internet-System. *Software Focus*, 1:15–22, September 2000. doi: 10.1002/1529-7950(20009)1:13.3.CO;2-#. URL https://www.researchgate.net/publication/220116945_Testing_software_for_Internet_application.

Erica Souza, Ricardo Falbo, and Nandamudi Vijaykumar. ROoST: Reference Ontology on Software Testing. *Applied Ontology*, 12:1–32, March 2017. doi: 10.3233/AO-170177.

Ephraim Suhir, Laurent Bechou, Alain Bensoussan, and Johann Nicolics. Photovoltaic reliability engineering: quantification testing and probabilistic-design-reliability concept. In *Reliability of Photovoltaic Cells, Modules, Components, and Systems VI*, volume 8825, pages 125–138. SPIE, September 2013. doi: 10.1117/12.2030377. URL <https://www.spiedigitallibrary.org/conference-proceedings-of-spie/8825/88250K/Photovoltaic-reliability-engineering--quantification-testing-and-probabilistic-design-reliability/10.1117/12.2030377.full>.

Guido Tebes, Denis Peppino, Pablo Becker, Gerardo Matturro, Martín Solari, and Luis Olsina. A Systematic Review on Software Testing Ontologies. pages

144–160. August 2019. ISBN 978-3-030-29237-9. doi: 10.1007/978-3-030-29238-6_11.

Guido Tebes, Luis Olsina, Denis Peppino, and Pablo Becker. TestTDO: A Top-Domain Software Testing Ontology. pages 364–377, Curitiba, Brazil, May 2020. ISBN 978-1-71381-853-3.

Daniel Trudnowski, Brian Pierre, Felipe Wilches-Bernal, David Schoenwald, Ryan Elliott, Jason Neely, Raymond Byrne, and Dmitry Kosterev. Initial closed-loop testing results for the pacific DC intertie wide area damping controller. In *2017 IEEE Power & Energy Society General Meeting*, pages 1–5, 2017. doi: 10.1109/PESGM.2017.8274724.

Kwok-Leung Tsui. An Overview of Taguchi Method and Newly Developed Statistical Methods for Robust Design. *IIE Transactions*, 24(5):44–57, May 2007. doi: 10.1080/07408179208964244. URL <https://doi.org/10.1080/07408179208964244>. Publisher: Taylor & Francis.

Matheus A. Tunes, Sean M. Drewry, Jose D. Arregui-Mena, Sezer Picak, Graeme Greaves, Luigi B. Cattini, Stefan Pogatscher, James A. Valdez, Saryu Fensin, Osman El-Atwani, Stephen E. Donnelly, Tarik A. Saleh, and Philip D. Edmondson. Accelerated radiation tolerance testing of Ti-based MAX phases. *Materials Today Energy*, 30(article 101186), October 2022. ISSN 2468-6069. doi: <https://doi.org/10.1016/j.mtener.2022.101186>. URL <https://www.sciencedirect.com/science/article/pii/S2468606922002441>.

Michael Unterkalmsteiner, Robert Feldt, and Tony Gorschek. A Taxonomy for Requirements Engineering and Software Test Alignment. *ACM Transactions on Software Engineering and Methodology*, 23(2):1–38, March 2014. ISSN 1049-331X, 1557-7392. doi: 10.1145/2523088. URL <http://arxiv.org/abs/2307.12477>. arXiv:2307.12477 [cs].

Petya Valcheva. Orthogonal Arrays and Software Testing. In Dimitar G. Velez, editor, *3rd International Conference on Application of Information and Communication Technology and Statistics in Economy and Education*, volume 200, pages 467–473, Sofia, Bulgaria, December 2013. University of National and World Economy. ISBN 978-954-644-586-5. URL <https://icaictsee-2013.unwe.bg/proceedings/ICAICTSEE-2013.pdf>.

Hans van Vliet. *Software Engineering: Principles and Practice*. John Wiley & Sons, Ltd., Chichester, England, 2nd edition, 2000. ISBN 0-471-97508-7.

Hironori Washizaki, editor. *Guide to the Software Engineering Body of Knowledge, Version 4.0*. January 2024. URL <https://waseda.app.box.com/v/SWEBOK4-book>.

Hironori Washizaki. Software Engineering Body of Knowledge (SWEBOK), February 2025. URL <https://www.computer.org/education/bodies-of-knowledge/software-engineering/>.

Wikibooks Contributors. *Haskell/Variables and functions*. Wikimedia Foundation, October 2023. URL https://en.wikibooks.org/wiki/Haskell/Variables_and_functions.

Han Yu, C. Y. Chung, and K. P. Wong. Robust Transmission Network Expansion Planning Method With Taguchi’s Orthogonal Array Testing. *IEEE Transactions on Power Systems*, 26(3):1573–1580, August 2011. ISSN 0885-8950. doi: 10.1109/TPWRS.2010.2082576. URL <https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=5620950>.

Kaiqiang Zhang, Chris Hutson, James Knighton, Guido Herrmann, and Tom Scott. Radiation Tolerance Testing Methodology of Robotic Manipulator Prior to Nuclear Waste Handling. *Frontiers in Robotics and AI*, 7(article 6), February 2020. ISSN 2296-9144. doi: 10.3389/frobt.2020.00006. URL <https://www.frontiersin.org/articles/10.3389/frobt.2020.00006>.

Changlin Zhou, Qun Yu, and Litao Wang. Investigation of the Risk of Electromagnetic Security on Computer Systems. *International Journal of Computer and Electrical Engineering*, 4(1):92, February 2012. URL <http://ijcee.org/papers/457-JE504.pdf>. Publisher: IACSIT Press.

Appendix A

Detailed Scope Analysis

As outlined in Chapter 1, the scope of our research is limited to testing applied to code itself. Throughout our research, we identify many approaches that are out of scope based on this criteria.

A.1 Hardware Testing

While testing the software run *on* or in control *of* hardware is in scope, testing performed on the hardware *itself* is out of scope. The following are some examples of hardware testing approaches we exclude from our research:

- Ergonomics testing and proximity-based testing (Hamburg and Mogyorodi, 2024) are out of scope, since they are used for testing hardware.
- EManations SECurity (EMSEC) testing (ISO, 2021; Zhou et al., 2012, p. 95), which deals with the “security risk” of “information leakage via electromagnetic emanation” (Zhou et al., 2012, p. 95), is also out of scope.
- All the examples of domain-specific testing given by Firesmith (2015, p. 26) are focused on hardware, so these examples are out of scope. However, this might not be representative of *all* kinds of domain-specific testing (e.g., Machine Learning (ML) model testing seems domain-specific), so some subset of this approach may be in scope.
- Similarly, the examples of environmental tolerance testing given by Firesmith (2015, p. 56) do not seem to apply to software. For example, radiation tolerance testing seems to focus on hardware, such as motors (Mukhin et al., 2022), robots (Zhang et al., 2020), or “nanolayered carbide and nitride materials” (Tunes et al., 2022, p. 1). Acceleration tolerance testing seems to focus on astronauts (Morgun et al., 1999, p. 11), aviators (Howe and Johnson, 1995, pp. 27, 42), or catalysts (Liu et al., 2023, p. 1463) and acoustic tolerance testing on rats (Holley et al., 1996), which are even less related! Since these all focus on environment-specific factors that would not impact the code, these examples are out of scope. As with domain-specific testing,

a subset of environmental tolerance testing may be in scope, but since no candidates have been found, this approach is out of scope for now.

- Knüvener Mackert GmbH (2022) uses the terms “software qualification testing” and “system qualification testing” in the context of the automotive industry. While these may be in scope, the more general idea of “qualification testing” seems to refer to the process of making a hardware component, such as an electronic component (Ahsan et al., 2020), gas generator (Parate et al., 2021) or photovoltaic device, “into a reliable and marketable product” (Suhir et al., 2013, p. 1). Therefore, it is currently unclear if this is in scope.
- Orthogonal Array Testing (OAT) can be used when testing software (Mandl, 1985) (in scope) but can also be used when testing hardware (Valcheva, 2013, pp. 471–472), such as “processors ... made from pre-built and pre-tested hardware components” (p. 471) (out of scope). A subset of OAT called “Taguchi’s Orthogonal Array Testing (TOAT)” is used for “experimental design problems in manufacturing” (Yu et al., 2011, p. 1573) or “product and manufacturing process design” (Tsui, 2007, p. 44) and is thus also out of scope.
- Since control systems often have a software *and* hardware component (ISO, 2015; Preuße et al., 2012; Forsyth et al., 2004), only the software component is in scope. In some cases, it is unclear whether the “loops”¹ being tested are implemented by software or hardware, such as those in wide-area damping controllers (Pierre et al., 2017; Trudnowski et al., 2017).
 - A related note: “path coverage” or “path testing” seems to be able to refer to either paths through code (as a subset of control-flow testing) (Washizaki, 2024, p. 5-13) or through a model, such as a finite-state machine (as a subset of model-based testing) (Doğan et al., 2014, p. 184).
- Physical testing (inferred from physical requirements (ISO/IEC and IEEE, 2017, p. 322) and requirements-based testing; see Section 3.2) tests “a physical characteristic that a system or system component must possess” (p. 322).

A.2 V&V of Other Artifacts

While many artifacts produced by the software life cycle can be tested, only testing performed on the code *itself* is in scope. Therefore, we exclude the following test approaches either in full or in part:

- Design reviews and documentation reviews are out of scope, as they focus on the V&V of design (ISO/IEC and IEEE, 2017, pp. 132) and documentation (p. 144), respectively.

¹Humorously, the testing of loops in chemical systems (Dominguez-Pumar et al., 2020) and copper loops (Goralski, 1999) are out of scope.

- Security audits can focus on “an organization’s ... processes and infrastructure” (Hamburg and Mogyorodi, 2024) (out of scope) or “aim to ensure that all of the products installed on a site are secure when checked against the known vulnerabilities for those products” (Gerrard, 2000b, p. 28) (in scope).
- Error seeding is the “process of intentionally adding known faults² to those already in a computer program”, done to both “monitor[] the rate of detection and removal”, which is a part of V&V of the V&V itself (out of scope), “and estimat[e] the number of faults remaining” (ISO/IEC and IEEE, 2017, p. 165), which helps verify the actual code (in scope).
- Fault injection testing, where “faults are artificially introduced² into the SUT [System Under Test]”, can be used to evaluate the effectiveness of a test suite (Washizaki, 2024, p. 5-18), which is a part of V&V of the V&V itself (out of scope), or “to test the robustness of the system in the event of internal and external failures” (ISO/IEC and IEEE, 2022, p. 42), which helps verify the actual code (in scope).
- “Mutation [t]esting was originally conceived as a technique to evaluate test suites in which a mutant is a slightly modified version of the SUT” (Washizaki, 2024, p. 5-15), which is in the realm of V&V of the V&V itself (out of scope). However, it “can also be categorized as a structure-based technique” and can be used to assist fuzz and metamorphic testing (Washizaki, 2024, p. 5-15) (in scope).
- Nontechnical testing (inferred from nontechnical requirements (ISO/IEC and IEEE, 2017, p. 293) and requirements-based testing; see Section 3.2) tests “product and service acquisition or development that is not a property of the product or service” (ISO/IEC and IEEE, 2017, p. 293).

A.3 Static Testing

Throughout the literature, static testing is more ambiguous than dynamic testing, with more ad hoc processes and inconsistent in/exclusion from the scope of software testing in general (see Definition Flaw 4). Furthermore, it seems less relevant to our original goal (the automatic generation of tests). In particular, many techniques require human intervention, either by design (such as code inspections) or to identify and resolve false positives (such as intentional exceptions to linting rules). Nevertheless, understanding the breadth of testing approaches requires a “complete” picture of how software can be tested and how the various approaches relate to one another. Parts of these static approaches may even be generated in

²While error seeding and fault injection testing both introduce faults as part of testing, they do so with different goals: to “estimat[e] the number of faults remaining” (ISO/IEC and IEEE, 2017, p. 165) and “test the robustness of the system” (2022, p. 42), respectively. Therefore, these approaches are not considered synonyms, and the lack of this relation in the literature is not included in Section 6.2.2 as a synonym flaw.

the future! For these reasons, we keep static testing in scope for this stage of our work, even though static testing will likely be removed at a later stage of analysis based on our original motivation.

See #41 and #44

A.4 Vague Terminology

Some terms are so vague that they do not provide any new, meaningful information. For example, the “systematic determination of the extent to which an entity meets its specified criteria” (ISO/IEC and IEEE, 2017, p. 167) is certainly relevant to testing software; while this definition of “evaluation” may be meaningful when defining software testing generally, it does not define a new approach or procedure and applies much more broadly than just to testing. We decided that the following terms are too vague to merit tracking in our glossaries or analyzing further:

See #39, #44, and #28

- **Evaluation:** the “systematic determination of the extent to which an entity meets its specified criteria” (ISO/IEC and IEEE, 2017, p. 167)
- **Product Analysis:** the “process of evaluating a product by manual or automated means to determine if the product has certain characteristics” (ISO/IEC and IEEE, 2017, p. 343)
- **Quality Audit:** “a structured, independent process to determine if project activities comply with organizational and project policies, processes, and procedures” (ISO/IEC and IEEE, 2017, p. 361)
- **Software Product Evaluation:** a “technical operation that consists of producing an assessment of one or more characteristics of a software product according to a specified procedure” (ISO/IEC and IEEE, 2017, p. 424)

OG PMBOK

A.5 Language-specific Approaches

Specific programming languages are sometimes used to define test approaches. If the reliance on a specific programming language is intentional, then this really implies an underlying test approach that may be generalized to other languages. These are therefore considered out-of-scope, including the following examples:

See #63

- “An approach ... for JavaScript testing (referred to as Randomized)” (Doğan et al., 2014, p. 192) is really just random testing used within JavaScript.
- “SQL statement coverage” (Doğan et al., 2014, Tab. 13) is really just statement coverage used specifically for SQL statements.
- Testing for “faults specific to PHP” is just a subcategory of fault-based testing, since “execution failures ... caused by missing an included file, wrong MySQL quer[ies] and uncaught exceptions” (Doğan et al., 2014, Tab. 27) are not exclusive to PHP.

OG Alalfi et al., 2010

OG Artzi et al., 2008

- While “HTML testing” is listed or implied by Gerrard (2000a, Tab. 2; 2000b, Tab. 1, p. 3) and Patton (2006, p. 220), it seems to be a combination of syntax testing, functionality testing, hyperlink testing/link checking, cross-browser compatibility testing, performance testing, content checking (Gerrard, 2000b, p. 3), and grey-box testing (Patton, 2006, pp. 218–220).

A.6 Orthogonally Derived Approaches

Some test approaches appear to be combinations of other (seemingly orthogonal) approaches. While the use of a combination term can sometimes make sense, such as when writing a paper or performing testing that focuses on the intersection between two test approaches, they are sometimes given the same “weight” as their atomic counterparts. For example, Hamburg and Mogyorodi (2024) include “formal reviews” and “informal reviews” in their glossary as separate terms, despite their definitions essentially boiling down to “reviews that follow (or do not follow) a formal process”, which do not provide any new information. These approaches are simply the combinations of “reviews” with “formal” and “informal testing”, respectively. If a source describes an orthogonally derived approach in more detail, such as security audits, we record it as a distinct approach in our test approach glossary with its related information. Otherwise, we consider it out of scope since its details are captured by its in-scope subapproaches. The following are examples of these orthogonally derived approaches, most of which are out of scope:

1. Black box conformance testing (Jard et al., 1999, p. 25)
2. Black-box integration testing (Sakamoto et al., 2013, pp. 345–346)
3. Checklist-based reviews (Hamburg and Mogyorodi, 2024)
4. Closed-loop HiL verification (Preuße et al., 2012, p. 6)
5. Closed-loop protection system testing (Forsyth et al., 2004, p. 331)
6. Endurance stability testing (Firesmith, 2015, p. 55)
7. End-to-end functionality testing (ISO/IEC and IEEE, 2021, p. 20; Gerrard, 2000a, Tab. 2)
8. Formal reviews (Hamburg and Mogyorodi, 2024; Washizaki, 2024, p. 12-14)
9. Grey-box integration testing (Sakamoto et al., 2013, p. 344)
10. Incremental integration testing (Sharma et al., 2021, pp. 601, 603, 605–606)
11. Informal reviews (Hamburg and Mogyorodi, 2024; Washizaki, 2024, p. 12-14)
12. Infrastructure compatibility testing (Firesmith, 2015, p. 53)
13. Invariant-based automatic testing (Doğan et al., 2014, pp. 184–185, Tab. 21), including for “AJAX user interfaces” (p. 191)

<p>Q #22: Is this orthogonal? Investigate legacy testing</p>	<ul style="list-style-type: none"> 14. <u>Legacy system integration (testing)</u> (Gerrard, 2000a, Tab. 2) 15. Manual procedure testing (Firesmith, 2015, p. 47) 16. Manual security audits (Gerrard, 2000b, p. 28) 17. Model-based GUI testing (Doğan et al., 2014, Tab. 1; implied by Sakamoto et al., 2013, p. 356) 18. Model-based web application testing (implied by Sakamoto et al., 2013, p. 356) 19. Non-functional search-based testing (Doğan et al., 2014, Tab. 1) 20. Offline MBT (Hamburg and Mogyorodi, 2024) 21. Online MBT (Hamburg and Mogyorodi, 2024)
<p>Q #23: Is this orthogonal? Investigate role-based testing</p>	<ul style="list-style-type: none"> 22. <u>Role-based reviews</u> (Hamburg and Mogyorodi, 2024) 23. <u>Scenario walkthroughs</u> (Gerrard, 2000a, Fig. 4) 24. <u>Scenario-based</u> reviews (Hamburg and Mogyorodi, 2024)
<p>Q #24: Is this orthogonal? Investigate scenario-based testing</p>	<ul style="list-style-type: none"> 25. Security attacks (Hamburg and Mogyorodi, 2024) 26. Security audits (ISO/IEC and IEEE, 2021, p. 40; Gerrard, 2000b, p. 28) 27. Statistical web testing (Doğan et al., 2014, p. 185) 28. Usability test script(ing) (Hamburg and Mogyorodi, 2024)
<p>Q #25: Is this orthogonal? Investigate scenario-based testing</p>	<ul style="list-style-type: none"> 29. Web application regression testing (Doğan et al., 2014, Tab. 21) 30. White-box unit testing (Sakamoto et al., 2013, pp. 345–346)

There are some cases where the subapproaches of the “compound” approaches listed previously are *not* from separate categories. However, these cases can be explained by insufficient data or by edge cases that require special care. While we assume that the categories given in Table 2.1 are orthogonal, further analysis may disprove this. For now, all of these special cases are affected by at least one of the following conditions:

1. **At least one subapproach is categorized inconsistently.** When a subapproach has more than one category (see Section 6.2.1), it is unclear which one should be used to assess orthogonality.
2. **At least one subapproach’s category is inferred.** When the category of a test approach is not given by the literature but is inferred from related context (see Section 2.3), it is unclear if it can be used to assess orthogonality.

3. **At least one subapproach is only categorized as an approach.** Since “approach” is a catch-all categorization, it does not need to be orthogonal to its subcategories.

4. **A subapproach is explicitly based on another in the same category.** An example of this is stability testing, which tests a “property that an object has with respect to a given failure mode if it cannot exhibit that failure mode” (ISO/IEC and IEEE, 2017, p. 434). This notion of “property” is similar to that of “quality” that the test type category is built on, so it is acceptable that is implied to be a test type by its quality (ISO/IEC and IEEE, 2017, p. 434) and by Firesmith (2015, p. 55).

OG ISO/IEC,
2009

OG ISO/IEC,
2009

Appendix B

Sources by Tier

The following lists of sources comprise the corresponding source tier as defined in Section 3.1.

B.1 Established Standards

ISO/IEC and IEEE, 2022; 2021; 2019a;b; 2017; 2016; 2013; 2010; IEEE, 2012; ISO/IEC, 2023a;b; 2018; 2015; 2011; ISO, 2022; 2015

B.2 Terminology Collections

Hamburg and Mogyorodi, 2024; Doğan et al., 2014; Firesmith, 2015; Washizaki, 2024; Bourque and Fairley, 2014

B.3 Textbooks

Ammann and Offutt, 2017; Dennis et al., 2012; Gerrard and Thompson, 2002; Kaner et al., 2011; Patton, 2006; Perry, 2006; Peters and Pedrycz, 2000; van Vliet, 2000

B.4 Papers and Other Documents

Bajammal and Mesbah, 2018; Barbosa et al., 2006; Baresi and Pezzè, 2006; Barr et al., 2015; Bas, 2024; Berdine et al., 2006; Chalin et al., 2006; ChatGPT (GPT-4o), 2024; Choudhary et al., 2010; Dhok and Ramanathan, 2016; Engström and Petersen, 2015; Forsyth et al., 2004; Gerrard, 2000a;b; Ghosh and Voas, 1999; Godefroid and Luchaup, 2011; Intana et al., 2020; Jard et al., 1999; Kam, 2008; Kanewala and Yueh Chen, 2019; Kuľšovs et al., 2013; Lahiri et al., 2013; LambdaTest, 2024; Mandl, 1985; Moghadam, 2019; Pandey, 2023; Preuße et al., 2012; Rus et al., 2008; Knüvener Mackert GmbH, 2022; Sakamoto et al., 2013; Sangwan and LaPlante, 2006; Sharma et al., 2021; Sneed and Göschl, 2000; Souza et al., 2017; Tsui, 2007; Valcheva, 2013; Yu et al., 2011

Appendix C

Full Lists of Flaws

The following are the full lists of manually identified flaws, first grouped by their manifestation, then by their domain as defined in Section 2.2.

C.1 Full Lists of Flaws by Manifestation

C.1.1 Full List of Mistakes

1. Since errors are distinct from defects/faults (ISO/IEC and IEEE, 2010, pp. 128, 140; Washizaki, 2024, p. 12-3; van Vliet, 2000, pp. 399–400), error guessing should instead be called “defect guessing” if it is based on a “checklist of potential defects” (ISO/IEC and IEEE, 2021, p. 29) or “fault guessing” if it is a “fault-based technique” (Bourque and Fairley, 2014, p. 4-9) that “anticipate[s] the most plausible faults in each SUT” (Washizaki, 2024, p. 5-13). One (or both) of these proposed terms may be useful in tandem with “error guessing”, which would focus on errors as traditionally defined; this would be a subapproach of error-based testing (implied by van Vliet, 2000, p. 399).
2. Similarly, “fault seeding” is not a synonym of “error seeding” as claimed by ISO/IEC and IEEE (2017, p. 165) and van Vliet (2000, p. 427). The term “error seeding”, also used by Firesmith (2015, p. 34), should be abandoned in favour of “fault seeding”, as it is defined as the “process of intentionally adding known faults to those already in a computer program ... [to] estimat[e] the number of faults remaining” (ISO/IEC and IEEE, 2017, p. 165) based on the ratio between the number of new faults and the number of introduced faults that were discovered (van Vliet, 2000, p. 427).
3. Hamburg and Mogyorodi (2024) classify ML model testing as a test level, which they define as “a specific instantiation of a test process”: a vague definition that does not match the one in Table 2.1.
4. Washizaki (2024, p. 5-4) says that quality improvement, along with quality assurance, is an aspect of testing that involves “defining methods, tools, skills, and practices to achieve the specific quality level and objective”; while

testing that a system possesses certain qualities is in scope, actively improving the system in response is *not* part of testing.

5. Washizaki (2024, p. 5-9)’s definition of “elasticity testing” only cites a single source **that doesn’t contain the words “elasticity” or “elastic”!**
6. The terms “acceleration tolerance testing” and “acoustic tolerance testing” seem to only refer to software testing in (Firesmith, 2015, p. 56); elsewhere, they seem to refer to testing the acoustic tolerance of rats (Holley et al., 1996) or the acceleration tolerance of astronauts (Morgun et al., 1999, p. 11), aviators (Howe and Johnson, 1995, pp. 27, 42), or catalysts (Liu et al., 2023, p. 1463), which don’t exactly seem relevant...
7. The differences between the terms “error”, “failure”, “fault”, “defect” are significant and meaningful (ISO/IEC and IEEE, 2010, pp. 128, 139–140; Washizaki, 2024, p. 12-3; van Vliet, 2000, pp. 399–400), but Patton (2006, pp. 13–14) “just call[s] it what it is and get[s] on with it”, abandoning these four terms, “problem”, “incident”, “anomaly”, “variance”, “inconsistency”, “feature” (!), and “a list of unmentionable terms” in favour of “bug”; after all, “there’s no reason to dice words”!
8. Kam (2008, p. 46) says that the goal of negative testing is “showing that a component or system does not work” which is not true; if robustness is an important quality for the system, then testing the system “in a way for which it was not intended to be used” (Hamburg and Mogyorodi, 2024) (i.e., negative testing) is one way to help test this!
9. Kam (2008, p. 42) says “See *boundary value analysis*,” for the glossary entry of “boundary value testing” but does not include “boundary value analysis” in the glossary.
10. A typo in (ISO/IEC and IEEE, 2021, Fig. 2) means that “specification-based techniques” is listed twice, when the latter should be “structure-based techniques”.
11. ISO/IEC and IEEE (2017) use the same definition for “partial correctness” (p. 314) and “total correctness” (p. 480).
12. The definition of “math testing” given by Hamburg and Mogyorodi (2024) is too specific to be useful, likely taken from an example instead of a general definition: “testing to determine the correctness of the pay table implementation, the random number generator results, and the return to player computations”.
13. A similar issue exists with multiplayer testing, where its definition specifies “the casino game world” (Hamburg and Mogyorodi, 2024).
14. “Par sheet testing” from (Hamburg and Mogyorodi, 2024) seems to refer to the specific example mentioned in Mistake 12 and does not seem more widely

OG Beizer

applicable, since a “PAR sheet” is “a list of all the symbols on each reel of a slot machine” (Bluejay, 2024).

15. The source that Hamburg and Mogyorodi (2024) cite for the definition of “test type” does not seem to actually provide a definition.
16. The same is true for “visual testing” (Hamburg and Mogyorodi, 2024).
17. The same is true for “security attack” (Hamburg and Mogyorodi, 2024).
18. Doğan et al. (2014, p. 184) claim that Sakamoto et al. (2013) define “prime path coverage”, but they do not.
19. Peters and Pedrycz (2000, Fig. 12.31) imply that decision coverage is a child of both c-use coverage *and* p-use coverage; this seems incorrect, since decisions are the result of p-uses and *not* c-uses (ISO/IEC and IEEE, 2021, pp. 5, 27; 2017, p. 332; van Vliet, 2000, p. 424), and only the p-use relation is implied by ISO/IEC and IEEE (2021, Fig. F.1).
20. Sharma et al. (2021, p. 601) seem to use the terms “grey-box testing” and “(stepwise) code reading” interchangeably, which would incorrectly imply that they are synonyms.
21. Kam (2008) misspells “state-based” as “state-base” (pp. 13, 15) and “stated-base” (Tab. 1).
22. Sneed and Göschl (2000, p. 18) give “white-box testing”, “grey-box testing”, and “black-box testing” as synonyms for “module testing”, “integration testing”, and “system testing”, respectively, but this mapping is incorrect; for example, ISO/IEC and IEEE (2017, p. 444) say “structure-based [(or white-box)] testing is not restricted to use at component level and can be used at all levels” and Sakamoto et al. (2013, pp. 345–346) describe “black-box integration testing”.
23. The previous flaw makes the claim that “red-box testing” is a synonym for “acceptance testing” (Sneed and Göschl, 2000, p. 18) lose credibility.
24. Kam (2008, p. 46) gives “mutation testing” as a synonym of “back-to-back testing”; while the two are related (ISO/IEC and IEEE, 2010, p. 30), the variants used in mutation testing are generated or designed to be detected as incorrect by the test suite (Washizaki, 2024, p. 5-15; similar in van Vliet, 2000, pp. 428–429) which is not a requirement of back-to-back testing.

OG 2015

OG Reid 1996

FIND SOURCES

OG Hetzel88

more examples?

Appendix D

Code Snippets

Source Code D.1: Tests for main with an invalid input file

```
# from
↳ https://stackoverflow.com/questions/54071312/how-to-pass-command-line-arguments
## \brief Tests main with invalid input file
# \par Types of Testing:
# Dynamic Black-Box (Behavioural) Testing
# Boundary Conditions
# Default, Empty, Blank, Null, Zero, and None
# Invalid, Wrong, Incorrect, and Garbage Data
# Logic Flow Testing
@mark.parametrize("filename", invalid_value_input_files)
@mark.xfail
def test_main_invalid(monkeypatch, filename):
    # from
    ↳ https://stackoverflow.com/questions/10840533/most-pythonic-way-to-delete-a-file
    try:
        remove(output_filename)
    except OSError as e: # this would be "except OSError, e:"
        ↳ before Python 2.6
        if e.errno != ENOENT: # no such file or directory
            raise # re-raise exception if a different error
            ↳ occurred

    assert not path.exists(output_filename)

    with monkeypatch.context() as m:
        m.setattr(sys, 'argv', ['Control.py',
            ↳ str(Path("test/test_input") / f"{filename}.txt")])
```

```
Control.main()
```

```
assert not path.exists(output_filename)
```

Source Code D.2: Projectile’s choice for constraint violation behaviour in code

```
srsConstraints = makeConstraints Warning Warning,
```

Source Code D.3: Projectile’s manually created input verification requirement

```
verifyParamsDesc = foldlSent [S "Check the entered", plural
  ↪ inValue,
  S "to ensure that they do not exceed the" ++. namedRef (datCon
    ↪ [] []) (plural datumConstraint),
  S "If any of the", plural inValue, S "are out of bounds" `sC`
  S "an", phrase errMsg, S "is displayed" `S.andThe` plural
    ↪ calculation, S "stop"]
```

Source Code D.4: “MultiDefinitions” (MultiDefn) Definition

```
-- | 'MultiDefn's are QDefinition factories, used for showing one
  ↪ or more ways
-- we can define a QDefinition.
data MultiDefn e = MultiDefn{
  -- | UID
  _rUId :: UID,
  -- | Underlying quantity it defines.
  _qd :: QuantityDict,
  -- | Explanation of the different ways we can define a quantity.
  _rDesc :: Sentence,
  -- | All possible ways we can define the related quantity.
  _rvs :: NE.NonEmpty (DefiningExpr e)
}
```

Source Code D.5: Pseudocode: Broken QuantityDict Chunk Retriever

```
retrieveQD :: UID -> ChunkDB -> Maybe QuantityDict
retrieveQD u cdb = do
  (Chunk expectedQd) <- lookup u cdb
  pure expectedQd
```
